

Imitation Learning in Continuous Action Spaces: Mitigating Compounding Error without Interaction

Thomas T. Zhang*
UPenn

Daniel Pfrommer†
MIT

Nikolai Matni*
UPenn

Max Simchowitz‡
CMU

Abstract

We study the problem of imitating an expert demonstrator in a continuous state-and-action dynamical system. While imitation learning in discrete settings such as autoregressive language modeling has seen immense success and popularity in recent years, imitation in physical settings such as autonomous driving and robot learning has proven comparably more complex due to the *compounding errors problem*, often requiring elaborate set-ups to perform stably. Recent work has demonstrated that even in benign settings, exponential compounding errors are unavoidable when learning solely from expert-controlled trajectories, suggesting the need for more advanced policy parameterizations or data augmentation. To this end, we present minimal interventions that provably mitigate compounding errors in continuous state-and-action imitation learning. When the system is open-loop stable, we prescribe “action chunking,” i.e., predicting and playing *sequences* of actions in open-loop; when the system is possibly unstable, we prescribe “noise injection,” i.e., adding noise during expert demonstrations. These interventions align with popular choices in modern robot learning, though the benefits we derive are distinct from the effects they were designed to target. Our results draw insights and tools from both control theory and reinforcement learning; however, our analysis reveals novel considerations that do not naturally arise when either literature is considered in isolation.

1 Introduction

Imitation learning (IL) is the problem of learning complex behaviors from data labeled with actions from an expert policy. This methodology encompasses both some of the earliest examples and most recent state-of-the-art in control for autonomous robotic systems [Pomerleau, 1988, Ross and Bagnell, 2010, Bojarski et al., 2016, Teng et al., 2023, Zhao et al., 2023]. Following the rise of large language models (LLMs), IL has also become increasingly prevalent in settings where an agent predicts *discrete tokens*, such as words in a sentence, lines in a proof, or positions on a chessboard [Chen et al., 2021]. Such methods have also seen adoption in the context of both continuous and discrete-action control of continuous state-space dynamical systems in an autoregressive fashion.

However, recent work [Simchowitz et al., 2025] shows that there exist nominally benign systems where *any algorithm* (including e.g. offline RL) for learning suitably regular, autoregressive next-action-prediction policies from offline expert data provably fails: there exist stable (i.e. contractive) systems for which the errors incurred by any learning algorithm compound exponentially-in-horizon when the policy is deployed (i.e. exhibit **compounding errors**), constituting a fundamental challenge to the soundness of imitation learning as a methodology.

*{ttz2, nmatni}@seas.upenn.edu

†dpfrom@mit.edu

‡msimchow@andrew.cmu.edu

As Simchowitz et al. [2025] eliminates the possibility of a simple “fix” to the *learning procedure*, we instead consider how changes to either (1) the policy parameterization or (2) the data-collection process can circumvent this negative result. We thereby elucidate both the design space of “sound” offline learning methodologies and better understand the success of widely-deployed algorithms such as data-augmentation [Laskey et al., 2017, Ke et al., 2021] and action-chunking [Zhao et al., 2023], even in the absence of further interaction with an expert demonstrator [Ross et al., 2011].

Contributions. We provide the first theoretical guarantees in continuous state-action IL for **interventions that provably prevent compounding error strictly from offline demonstrations**. Whereas previous work [Ross et al., 2011, Laskey et al., 2017, Pfrommer et al., 2022] require either iterative interaction with the expert or knowledge of the underlying system, we do so without access to such oracles, using near-“vanilla” behavior cloning. Our results focus on two categories of interventions.

Intervention 1: Action-Chunking. When the environment is benign, we show the algorithmic modification of *action-chunking*, i.e., predicting and playing open-loop sequences of actions, mitigates compounding errors without requiring any modification to the expert data (Theorem 1).

Intervention 2: Expert Noise-Injection. When the environment is less benign, some alteration of the expert data distribution is necessary. We demonstrate *noise-injection*, i.e., adding noise while executing expert actions, is a simple and practical tool for avoiding compounding errors (Theorem 2).

Surprising Takeaways. While Intervention 1 and 2 are reflective of popular practices, our analysis additionally reveals *the intersection of (reinforcement) learning and control leads to new practical and analytical considerations that do not naturally arise when either is considered in isolation*, uncovering phenomena that contrast with the common perspectives of both literatures. In particular:

- Action-chunking has been motivated by both enabling larger policy latency, and encouraging long-horizon planning and stronger policy representations [Chi et al., 2023, Liu et al., 2025]. We illuminate an entirely orthogonal rationale: **action-chunking encourages control-theoretic stability of policies learned**, thereby mitigating possibly exponential compounding errors.
- We show that “coverage” in standard RL [Jiang and Xie, 2024] and “persistent excitation” in control theory [Bai and Sastry, 1985] are overly conservative when describing the role of exploration via noise injection in mitigating compounding errors. Surprisingly, we show that *naive exploration* via white noise suffices, even without ensuring density ratios or when the underlying system is not even controllable [Kailath, 1980]. Rather, we show the **error-directions which are susceptible to compounding error are precisely those along which Intervention 2 provides supervision**.

Our non-trivial analytical insights directly translate to algorithmic modifications with tangible consequences, which we depict in Figure 1.

Related Work. Imitation learning from expert demonstrations has emerged as a dominant technique for learning performant models across applications such as: self-driving vehicles [Hussein et al., 2017, Bojarski et al., 2016, Bansal et al., 2018], visuomotor policies [Finn et al., 2017, Zhang et al., 2018], and large-scale robotic decision-making models [Zitkovich et al., 2023, Black et al., 2024]. As such, the compounding error phenomenon is well-documented, dating back even to the introduction of IL [Pomerleau, 1988].

In discrete state-action settings, the seminal work in Ross and Bagnell [2010], Ross et al. [2011] propose an *interactive* procedure to collect examples of corrective data, seeing widespread adoption [Kelly et al., 2019, Sun et al., 2023]. On the theoretical side, compounding errors appears more benign in discrete settings, with naive behavior cloning (BC) attaining a discrepancy between training and ex-




Figure 1: Visualization of the benefits of action-chunking (Intervention 1) and noise-injection (Intervention 2). **Left:** even on synthetic *globally EISS* (Definition 2.1) dynamics f , frequent feedback can cause exponential compounding error, which action-chunking mitigates. **Right:** on a HalfCheetah-v5 environment, we see *any* amount of noise injection yields significant performance improvement. All experiment details are in Appendix C.

ecution error at most quadratic in the horizon. Recent work by Foster et al. [2024] even demonstrates that modifying the loss may result in performance that has *no* adverse dependence on horizon. However, these works operate in settings ill-suited for continuous control, where the expert policy must be estimated in information-theoretic distances that are not feasible, e.g., even for deterministic policies in continuous action spaces.

Accordingly, prior work which applies IL to continuous control settings has involved more elaborate set-ups to enable stable performance. For example, recent advances in generative policies are typically paired with action-chunked execution (Intervention 1), see e.g., [Chen et al., 2021, Shafuallah et al., 2022, Chi et al., 2023, Zhao and Grover, 2023, Liu et al., 2025]. Other works have considered tools from robust control [Hertneck et al., 2018, Yin et al., 2021] and stability regularization [Sindhwani et al., 2018, Mehta et al., 2025] to promote stability around observed data. Lastly, various works have proposed different forms of data augmentation as a way to promote robustness to distribution shift, including iteratively shaped noise injection during expert demonstrations [Laskey et al., 2017], and noising observed states/actions [Ke et al., 2021, 2024, Block et al., 2024]. Our proposed Intervention 2 can be viewed as a distilled, non-iterative version of DART [Laskey et al., 2017].

A complementary line of work has attempted to understand the theoretical foundations of imitating in continuous settings. Tu et al. [2022] parameterize a scale of “incremental stability” (see Definition 2.1) and study its impact on the statistical generalization of IL. Pfrommer et al. [2022] proposes sufficient conditions for benign compounding errors in a similar setting. However, the resulting algorithms have exceedingly strong requirements, e.g., stability oracles or $\partial \text{input} / \partial \text{state}$ derivative sketching, respectively. Bringing this line of work to a close, Simchowicz et al. [2025] offers definitive evidence that exponential compounding errors cannot be avoided by altering the learning procedure, motivating the interventions we propose. We restate the relevant lower bounds more formally in Theorem A.

2 Preliminaries

We consider a discrete-time, continuous state-action control system with states $\mathbf{x}_t \in \mathcal{X} = \mathbb{R}^{d_x}$ and inputs¹ $\mathbf{u}_t \in \mathcal{U} = \mathbb{R}^{d_u}$, where dynamics deterministically evolve according to $\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{u}_t)$. A deterministic policy π maps histories of states, inputs, and the current time step to a control input $\mathbf{u}_t = \pi(\mathbf{x}_{1:t}, \mathbf{u}_{1:t-1}, t)$. We assume the initial state is drawn $\mathbf{x}_1 \sim D$ for some distribution D fixed throughout. We say π is Markovian and time-invariant if we can simply express $\mathbf{u}_t = \pi(\mathbf{x}_t)$. In this case, we define the closed-loop dynamics $f^\pi(\mathbf{x}, \mathbf{u}) \triangleq f(\mathbf{x}, \pi(\mathbf{x}) + \mathbf{u})$, and $f^\pi(\mathbf{x}) = f^\pi(\mathbf{x}, 0)$.

¹We refer to inputs and actions interchangeably.

We let \mathbb{E}_π (resp. \mathbb{P}_π) denote expectation (resp. law) under $\mathbf{x}_1 \sim D$, the dynamics f , and inputs selected by the policy π . Given two deterministic policies π_1, π_2 , we let $\mathbb{E}_{\pi_1, \pi_2}$ denote the expectation of sequences $(\mathbf{x}_t^{\pi_i}, \mathbf{u}_t^{\pi_i})_{t \geq 1}$ under the dynamics f , coupled so that $\mathbf{x}_1^{\pi_1} = \mathbf{x}_1^{\pi_2} \sim D$. We consider estimation of **deterministic, Markovian** expert policies $\pi^* : \mathcal{X} \rightarrow \mathcal{U}$ given a problem horizon T . Our aim is to learn some policy $\hat{\pi}$ which accumulates low squared-**trajectory error**:

$$\mathbf{J}_{\text{TRAJ}, T}(\hat{\pi}) \triangleq \mathbb{E}_{\hat{\pi}, \pi^*} \left[\sum_{t=1}^T \min \{1, \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 + \|\mathbf{u}_t^{\hat{\pi}} - \mathbf{u}_t^{\pi^*}\|^2\} \right]. \quad (2.1)$$

Above, the practice of taking a minimum with 1 accounts for the possibility of unbounded trajectories on rare events; 1 can be replaced by an arbitrary constant. Upper bounds on $\mathbf{J}_{\text{TRAJ}, T}$ imply upper bounds on the difference in expected Lipschitz costs, see e.g., [Appendix A](#).

Empirical risk minimization. We consider estimates $\hat{\pi}$ of π^* obtained via empirical risk minimization. Given a sample S_n of n trajectories $(\mathbf{x}_t^{(i)}, \mathbf{u}_t^{(i)})_{1 \leq t \leq T, 1 \leq i \leq n}$, the empirical risk is:

$$\mathbf{J}_{\text{EMP}, T}(\hat{\pi}; S_n) \triangleq \sum_{i=1}^n \sum_{t=1}^T \|\hat{\pi}(\mathbf{x}_{1:t}^{(i)}, \mathbf{u}_{1:t-1}^{(i)}, t) - \mathbf{u}_t^{(i)}\|^2. \quad (2.2)$$

In our analysis, we also consider population quantities. Let \mathbb{P}_{demo} denote a probability distribution over demonstrations $(\mathbf{x}_t, \mathbf{u}_t)_{1 \leq t \leq T}$. We let $\mathbb{P}_{\text{demo}} = \mathbb{P}_{\pi^*}$ in [Intervention 1](#), but consider modifications beyond the expert distribution for [Intervention 2](#). In either case we can define a population-level risk:

$$\mathbf{J}_{\text{DEMO}, T}(\hat{\pi}; \mathbb{P}_{\text{demo}}) \triangleq \mathbb{E}_{\mathbb{P}_{\text{demo}}} \left[\sum_{t=1}^T \|\hat{\pi}(\mathbf{x}_{1:t}, \mathbf{u}_{1:t-1}, t) - \mathbf{u}_t\|^2 \right] \quad (2.3)$$

Note that if S_n consists of n i.i.d. trajectories from \mathbb{P}_{demo} , then $\mathbb{E}[\mathbf{J}_{\text{EMP}, T}(\hat{\pi}; S_n)] = \mathbf{J}_{\text{DEMO}, T}(\hat{\pi}; \mathbb{P}_{\text{demo}})$, so minimizing $\mathbf{J}_{\text{EMP}, T}$ is an empirical surrogate for $\mathbf{J}_{\text{DEMO}, T}$. We broadly term $\mathbf{J}_{\text{DEMO}, T}$ the **on-expert error** (and $\mathbf{J}_{\text{EMP}, T}$ its empirical counterpart). Gauging how well various algorithmic interventions (or lack thereof) mitigate compounding errors therefore translates to bounding the *trajectory error* $\mathbf{J}_{\text{TRAJ}, T}$ in terms of the *on-expert error* $\mathbf{J}_{\text{DEMO}, T}$.

The Compounding Errors problem. We formally describe the compounding errors problem. Let alg be a (possibly randomized) mapping from a sample of n trajectories $S_n \stackrel{\text{i.i.d.}}{\sim} \mathbb{P}_{\text{demo}}$ to an imitator policy $\hat{\pi} \sim \text{alg}(S_n)$. The problem instance suffers *exponential compounding errors* if:

$$\mathbb{E}_{\hat{\pi}, S_n} [\mathbf{J}_{\text{TRAJ}, T}(\hat{\pi})] \gtrsim C^T \cdot \mathbb{E}_{\hat{\pi}, S_n} [\mathbf{J}_{\text{DEMO}, T}(\hat{\pi}; \mathbb{P}_{\text{demo}})], \quad (2.4)$$

for some $C > 1$. In other words, imitating via empirical risk minimization on a given demonstration distribution \mathbb{P}_{demo} leads to learned policies $\hat{\pi}$ that suffer exponentially more **trajectory error** rolled out in closed-loop compared to their **on-expert** regression error. Compounding error can be understood through the lens of control-theoretic *stability*, which describes the sensitivity of the dynamics to perturbations of the state or input. We consider a notion of *incremental stability* [[Angeli, 2002](#), [Tran et al., 2016](#)].

Definition 2.1 (EISS). A system $\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{u}_t)$ is (C_{ISS}, ρ) -exponentially incrementally input-to-state stable (EISS) if for all pairs of initial conditions $(\mathbf{x}_1, \mathbf{x}'_1)$ and input sequences $(\{\mathbf{u}_t\}_{t \geq 1}, \{\mathbf{u}'_t\}_{t \geq 1})$, there exist constants $C_{\text{ISS}} \geq 1, \rho \in (0, 1)$ ² such that for any $t \geq 1$:

$$\|\mathbf{x}_t - \mathbf{x}'_t\| \leq C_{\text{ISS}} \rho^{t-1} \|\mathbf{x}_1 - \mathbf{x}'_1\| + C_{\text{ISS}} \sum_{k=1}^{t-1} \rho^{t-1-k} \|\mathbf{u}_k - \mathbf{u}'_k\|, \quad t \geq 1.$$

We say a policy-dynamics pair (π, f) is (C_{ISS}, ρ) -EISS if the induced closed-loop dynamics f^π is (C_{ISS}, ρ) -EISS. We also denote the shorthands $C_{\text{stab}} \triangleq \frac{C_{\text{ISS}}}{1-\rho}$, $c_{\text{stab}} \triangleq C_{\text{stab}}^{-1}$.

In other words, incremental stability ensures bounded input perturbations lead to bounded future state deviations, with their effect decaying in time.³ Our base assumption moving forward is that the *expert-induced* closed-loop system f^{π^*} is incrementally stable. This formalizes a notion of expert robustness by implying the expert can recover from bounded input perturbations. As strong as this may seem, EISS of the expert does not assume away the compounding errors issue (see [Theorem A](#)); for example, if the candidate policy $\hat{\pi}$ destabilizes the system, the resulting “input perturbations” to f^{π^*} are exponentially growing.

Theorem A (Motivating lower bounds, Informal vers. of [Simchowit et al. \[2025, Theorems 1 & 4\]](#)). *There exists families $\mathcal{P}_{\text{stab}} = \{(\pi, g) : (\pi, g) \in \mathcal{P}_{\text{stab}}\}$ and $\mathcal{P}_{\text{unst}} = \{(\pi, g) : (\pi, g) \in \mathcal{P}_{\text{unst}}\}$ of policies and dynamics such that:*

- (i) *For every $(\pi, g) \in \mathcal{P}_{\text{stab}}$, g is open-loop EISS and (π, g) is closed-loop EISS, and π, g are Lipschitz and smooth. However, any learning algorithm which returns smooth, Lipschitz, Markovian policies with state-independent stochasticity must suffer exponential-in- T compounding error (2.4) when learning from n expert trajectories from some $(\pi, g) \in \mathcal{P}_{\text{stab}}$.*
- (ii) *For every $(\pi, g) \in \mathcal{P}_{\text{unst}}$, (π, g) is closed-loop EISS (but g is not open-loop EISS), and π, g are Lipschitz and smooth. However, any learning algorithm, without restriction, suffers exponential-in- T compounding error (2.4) when learning from n expert trajectories on some $(\pi, g) \in \mathcal{P}_{\text{stab}}$.*

In the bound above, $\mathcal{P}_{\text{unst}}, \mathcal{P}_{\text{stab}}$ describe families of problem instances. The bounds ensure that for at least one instance (π, g) in $\mathcal{P}_{\text{stab}}$ (resp. $\mathcal{P}_{\text{unst}}$), the learner suffers exponential-in- T compounding error if that instance (π, g) is the ground truth, and the learner receives n expert demonstrations from that instance. In the case of $\mathcal{P}_{\text{stab}}$, where g is open-loop EISS, the lower bounds only apply to the class of smooth, Lipschitz, Markovian policies with state-independent stochasticity; however, when g are no longer required to be open-loop EISS, the bound holds without restriction. Our results constitute *positive converses*: when g is open-loop EISS, we can use [Intervention 1](#) to construct a smooth, Lipschitz but non-Markovian (i.e., chunked!) policies to bypass [Theorem A.\(i\)](#), and if g is not necessarily EISS, we can still use [Intervention 2](#)—which provides a *different distribution* over expert demonstrations—to circumvent [Theorem A.\(ii\)](#), which otherwise precludes any purely algorithmic changes that do not alter the distribution over training data.

Additional Notation. **Blue** (e.g. π^*) indicates expert-induced quantities, and **red** indicates quantities induced by a learned policy (e.g. $\hat{\pi}$). Positive semi-definite matrices are indicated by $\mathbf{Q} \succeq \mathbf{0}$, and the corresponding partial order $\mathbf{P} \succeq \mathbf{Q} \implies (\mathbf{P} - \mathbf{Q}) \succeq \mathbf{0}$. We use \lesssim, \approx to omit universal constants. In the

²We note traditional definitions of nonlinear stability may track separate β, γ for the transient bound $\beta(\|\mathbf{x} - \mathbf{x}'\|, t)$ and the input gain $\gamma(\|\mathbf{u} - \mathbf{u}'\|)$. For our purposes, it suffices to lump these together under C_{ISS} for clarity.

³The stability definition and ensuing results can be loosened to polynomial decay or local variants with appropriate modifications, though we note the ensuing lower bounds [Theorem A](#) already hold for EISS.

main body, we also use $O_\star(\cdot)$ to omit *polynomial* dependence on instance-dependent constants, but not algorithm-dependent constants or horizon T , e.g. $\frac{T C_{\text{ISS}}}{1-\rho} \sigma_{\mathbf{u}}^2 = O_\star(T \sigma_{\mathbf{u}}^2)$.

3 Action Chunking Suffices in Open-Loop Stable Systems

Action-chunking (Intervention 1) is a popular practice in modern sequential modeling pipelines, where a policy predicts a sequence of actions, of which some number are played *in open-loop* [Chen et al., 2021, Chi et al., 2023, Shafiullah et al., 2022]. There are various intuitions of the practical benefits of action-chunking, ranging from: amenability to multi-modal⁴ prediction, robustness to non-Markovian quirks in the data [Liu et al., 2025], and reducing the effective planning horizon (by dividing the horizon by the chunk length). Yet, we show that even in control settings with *unimodal, Markovian, state-feedback* experts, action-chunking serves a critical role in subverting exponential compounding errors. All proofs and extended details in this section are contained in Appendix A. We may conveniently describe chunking as follows.

Definition 3.1 (Chunking Policy). A deterministic, Markovian chunking policy is specified by a chunk-length ℓ , and mappings $\text{chunk}_i[\pi] : \mathcal{X} \rightarrow \mathcal{U}$, $i \in [\ell]$ such that, for $k \in \mathbb{Z}_{\geq 0}$ and $i \in [\ell]$ and $t = k\ell + i$ for, $\pi(\mathbf{x}_{1:t}, \mathbf{u}_{1:t-1}, t) = \text{chunk}_i[\pi](\mathbf{x}_{k\ell+1}, i)$. We also write $\text{chunk}[\pi](\mathbf{x}) = (\text{chunk}_1(\mathbf{x}), \dots, \text{chunk}_\ell(\mathbf{x}))$. For simplicity, we always assume ℓ divides $T - 1$.

Intervention 1 (Learning over Chunked Policies). We sample S_n as denote n i.i.d. trajectories drawn from the expert distribution \mathbb{P}_{π^\star} . We then minimize $\mathbf{J}_{\text{EMP}, T}(\hat{\pi}; S_n)$ over a class of length- ℓ chunked policies, $\Pi_{\text{chunk}, \ell}$, defined formally in Definition 3.2. We notice that for chunked policies, we have

$$\mathbf{J}_{\text{EMP}, T}(\hat{\pi}; S_n) = \sum_{i=1}^n \sum_{k=1}^{T-1/\ell} \|\mathbf{u}_{1+(k-1)\ell:k\ell}^{(i)} - \text{chunk}[\hat{\pi}](\mathbf{x}_{(k-1)\ell}^{(i)})\|^2. \quad (3.1)$$

Let us formally define the policies induced by chunking with a dynamics model.

Definition 3.2 (Induced Chunking Policy). Let $g : \mathcal{X} \times \mathcal{U} \rightarrow \mathcal{U}$ be a dynamical map (possibly not the true dynamics f), and $\pi : \mathcal{X} \rightarrow \mathcal{U}$ a Markovian, deterministic policy. Given chunk length $\ell \in \mathbb{N}$, we define the induced chunked policy $\tilde{\pi} = \text{chunked}(\pi, g, \ell)$, $\tilde{\pi} : \mathcal{X} \rightarrow (\mathcal{U})^\ell$ as returning

$$\text{chunk}[\tilde{\pi}](\mathbf{x}) = (\pi(\mathbf{x}), \pi(g^\pi(\mathbf{x})), \pi((g^\pi)^2(\mathbf{x})), \dots, \pi((g^\pi)^{\ell-1}(\mathbf{x}))), \quad (3.2)$$

where above $g^\pi(\mathbf{x}) \triangleq g(\mathbf{x}, \pi(\mathbf{x}))$, and $(g^\pi)^i$ is understood as repeated composition.

In other words, $\text{chunked}(\pi, g, \ell)$ returns a policy that, conditioning on the current state, outputs the next ℓ actions given by simulating π on dynamics g in closed-loop. We now lay out the core assumptions moving forward.

Assumption 3.1 (Regularity and Stability). We make the following assumptions:

1. The true dynamics f are (C_{ISS}, ρ) -EISS in open-loop, without loss of generality with $\rho \geq 1/e$.
2. All base policies $\pi \in \Pi \cup \{\pi^\star\}$ in consideration are L_π -Lipschitz: $\|\pi(\mathbf{x}) - \pi(\mathbf{x}')\| \leq L_\pi \|\mathbf{x} - \mathbf{x}'\|$.

In other words, we assume the dynamics f are open-loop stable. All ensuing results regarding imitation learning with chunked policies stem from the following key result.

⁴In the sense of a distribution having multiple modes.

Proposition 3.1. Let [Assumption 3.1](#) hold. Let $(\hat{\pi}, \hat{f})$ be a policy-dynamics pair that is (C_{ISS}, ρ) -EISS, and consider the corresponding chunked policy $\tilde{\pi} = \text{chunked}(\hat{\pi}, \hat{f}, \ell)$. Then the closed-loop system the chunked policy induces on the true dynamics $(\tilde{\pi}, f)$ is $(\tilde{C}, \tilde{\rho})$ -EISS, where $\tilde{C} = \text{poly}(L_{\pi}, C_{\text{ISS}}, \frac{1}{1-\rho})$ and $\tilde{\rho} = \hat{\rho}^{1/2}$, as long as the chunk length is sufficiently long: $\ell > \text{polylog}(L_{\pi}, C_{\text{ISS}}) \cdot \text{polylog}^{-1}(1/\rho)$.

This result states that: as long as a policy “believes” it stabilizes the *simulated* dynamics at hand, then it is guaranteed to be stable on the actual dynamics if it is chunked accordingly. Crucially, without action-chunking, open-loop stability of the nominal dynamics f and closed-loop stability of the expert (π^*, f) need not ensure closed loop stability of $(\hat{\pi}, f)$ for the learned policy in the absence of action chunking $\ell = 1$; see [Theorem A.\(i\)](#).

Contrast this to [Proposition 3.1](#), which depends only on the stability properties of the true system f and the closed-loop simulated system $(\hat{\pi}, \hat{f})$, and requires *no* assumption on the closeness of \hat{f} to f , or $\hat{\pi}$ to any reference policy. This implies the stark benefit of chunking, where relatively short chunk lengths (logarithmic in stability parameters) mark the difference between exponential blow-up and exponential stability. This leads to benign conversion between the on-expert and trajectory error of a chunked imitator policy. We first pose the statistical learning problem over chunked policies.

Assumption 3.2 (Realizability of Chunked Policies). We assume that we have access to a class of possible policy-dynamics pairs $\mathcal{P} \triangleq \{(\pi, g)\}$ such that for all $(\pi, g) \in \mathcal{P}$, (π, g) is EISS with constants (C_{ISS}, ρ) . We assume $(\pi^*, f) \in \mathcal{P}$. From this class, we define given chunk length ℓ the induced policy class: $\Pi_{\text{chunk}, \ell} \triangleq \{\tilde{\pi} = \text{chunked}(\pi, g, \ell) : (\pi, g) \in \mathcal{P}\}$.

We note that if g matches the deployment dynamics f , then $\text{chunked}(\pi, g, \ell)$ returns the same actions as π in closed-loop. Therefore, the expert demonstrations (π^*, f) are realizable as an element of $\Pi_{\text{chunk}, \ell}$ for any ℓ , such that $\inf_{\tilde{\pi} \in \Pi_{\text{chunk}, \ell}} \mathbf{J}_{\text{DEMO}, T}(\tilde{\pi}; \mathbb{P}_{\pi^*}) = 0$. A key consequence of a chunked policy inducing stable closed-loop dynamics is that it accumulates limited compounding error.

Proposition 3.2. Let [Assumption 3.1](#) hold. Let $\tilde{\pi} = \text{chunked}(\hat{\pi}, \hat{f}, \ell) \in \Pi_{\text{chunk}, \ell}$, and assume $(\hat{\pi}, \hat{f})$, $(\tilde{\pi}, f)$ are $(\tilde{C}, \tilde{\rho})$ -EISS. Then, the following bound holds:

$$\mathbf{J}_{\text{TRAJ}, T}(\tilde{\pi}) \leq \text{poly}(\tilde{C}, \frac{1}{1-\tilde{\rho}}) \cdot O_{\star}(\ell) \mathbf{J}_{\text{DEMO}, T}(\tilde{\pi}; \mathbb{P}_{\pi^*}).$$

Therefore, combining [Proposition 3.1](#) and [Proposition 3.2](#) leads to the following compounding error guarantee on any sufficiently chunked policy.

Theorem 1. Let [Assumption 3.1](#) and [Assumption 3.2](#) hold. For sufficiently long chunk-length: $\ell > \text{polylog}(L_{\pi}, C_{\text{ISS}}) \text{polylog}^{-1}(1/\rho)$, let $\tilde{\pi} = \text{chunked}(\hat{\pi}, g, \ell) \in \Pi_{\text{chunk}, \ell}$. The following bound holds on the trajectory error:

$$\mathbf{J}_{\text{TRAJ}, T}(\tilde{\pi}) \leq O_{\star}(\ell) \mathbf{J}_{\text{DEMO}, T}(\tilde{\pi}; \mathbb{P}_{\pi^*}).$$

[Theorem 1](#) implies that when the ambient dynamics f are EISS, then a sufficiently chunked imitator policy will accrue limited compounding errors relative to the on-expert error it sees. In particular, given $\tilde{\pi}$ attaining estimation error $\mathbf{J}_{\text{DEMO}, T}(\tilde{\pi}; \mathbb{P}_{\pi^*}) \leq \inf_{\tilde{\pi} \in \Pi_{\text{chunk}, \ell}} \mathbf{J}_{\text{DEMO}, T}(\tilde{\pi}; \mathbb{P}_{\pi^*}) + \epsilon_{\text{ERM}}^2$, this implies $\mathbf{J}_{\text{TRAJ}, T}(\tilde{\pi}) \leq O_{\star}(\ell) \epsilon_{\text{ERM}}^2$. Notice that ℓ can be chosen to be a problem-dependent constant, and thus does not grow with the horizon T . Hence, the complexity of the chunked class $\Pi_{\text{chunk}, \ell}$, which implicitly determines the the magnitude of ϵ_{ERM} when a statistical learning bound is applied, should also be viewed as T -independent. We further note that $\Pi_{\text{chunk}, \ell}$ extends the dimension of the predicted variable (actions) by a factor of ℓ , but not that of the input variable \mathbf{x} , so the statistical effect of chunking is likely to be benign. We defer further discussion of subtleties that distinguish the chunking policy class $\Pi_{\text{chunk}, \ell}$ compared to an ℓ -times product of Π to [Appendix A](#).

4 Noise Injection Mitigates Compounding Error under Smooth Dynamics

We now consider the difficult setting where the ambient dynamics f may not be open-loop stable. In this case, purely algorithmic interventions like action-chunking are generally insufficient, as erroneous actions can lead to unstable behavior. In fact, we recall [Theorem A](#) states that *no* algorithm, even permitting stochastic and non-Markovian policies, can circumvent exponential compounding errors in the worst-case, provided only data from the expert-induced law \mathbb{P}_{π^*} . This necessitates a modification to the demonstration distribution \mathbb{P}_{demo} beyond the expert’s \mathbb{P}_{π^*} . We consider inducing exploration in the expert dataset via *noise injection*. In the discussion below, we fix a *noise level* $\sigma_u > 0$, which controls the magnitude of the noise added, and a *mixture fraction* $\alpha \in [0, 1]$, that controls the proportion of trajectories collected without noise injection.

Definition 4.1. We define the *expert distribution under noise injection* as the distribution $\mathbb{P}_{\pi^*, \sigma_u}$ over trajectories $(\tilde{\mathbf{x}}_t, \tilde{\mathbf{u}}_t)_{t \geq 1}$ with $\tilde{\mathbf{x}}_1 \sim D$, and $\tilde{\mathbf{u}}_t = \pi^*(\tilde{\mathbf{x}}_t)$, $\tilde{\mathbf{x}}_{t+1} = f(\tilde{\mathbf{x}}_t, \tilde{\mathbf{u}}_t + \sigma_u \mathbf{z}_t)$ for $t \geq 1$, where $\mathbf{z}_t \stackrel{\text{i.i.d.}}{\sim} \text{Unif}(\mathbb{B}^{d_u}(1))$ is drawn uniformly over the unit ball.⁵

In other words, noise injection collects trajectories induced when the expert’s inputs are executed with additive noise $\sigma_u \mathbf{z}_t$. We then consider fitting a policy π by augmenting standard (un-noised) expert trajectories with noise-injected ones.

Intervention 2 (Expert Noise Injection). For the noise-injected distribution $\mathbb{P}_{\pi^*, \sigma_u}$ defined above, provide a sample $S_{n, \sigma, \alpha}$ of $(\mathbf{x}_t^{(i)}, \mathbf{u}_t^{(i)})_{1 \leq t \leq T, 1 \leq i \leq n}$, where for $1 \leq i \leq \lfloor \alpha n \rfloor$ the trajectories are i.i.d. from \mathbb{P}_{π^*} , and the remaining trajectories are drawn i.i.d. from $\mathbb{P}_{\pi^*, \sigma_u}$. Define the corresponding mixture distribution $\mathbb{P}_{\pi^*, \sigma_u, \alpha} \triangleq \alpha \mathbb{P}_{\pi^*} + (1 - \alpha) \mathbb{P}_{\pi^*, \sigma_u}$. We then select $\hat{\pi}$ to minimize $\mathbf{J}_{\text{EMP}, T}(\pi; S_{n, \sigma, \alpha})$, which we observe is an empirical estimate of $\mathbf{J}_{\text{DEMO}, T}(\pi; \mathbb{P}_{\pi^*, \sigma_u, \alpha})$.

In other words, [Intervention 2](#) combines an α fraction of trajectories from the expert demonstrator (the law \mathbb{P}_{π^*}), and augments them with a $1 - \alpha$ fraction executing *noise-injected* expert actions $\pi^*(\mathbf{x}_t) + \sigma_u \mathbf{z}_t$ in the environment (the law $\mathbb{P}_{\pi^*, \sigma_u}$). Note, while the actions under $\mathbb{P}_{\pi^*, \sigma_u}$ are executed noisily, the recorded action labels are *noiseless* $\mathbf{u}_t = \pi^*(\mathbf{x}_t)$, thereby avoiding additional regression error. We now lay out the core assumptions on the expert and dynamics in this section.

Assumption 4.1 (Regularity and Stability). Recall that a function $h : \mathbb{R}^d \rightarrow \mathbb{R}^p$ C -smooth if for all $\mathbf{x}, \mathbf{x}' \in \mathbb{R}^d$, $\|\nabla_{\mathbf{x}} h(\mathbf{x}) - \nabla_{\mathbf{x}} h(\mathbf{x}')\|_2 \leq C \|\mathbf{x} - \mathbf{x}'\|$. We make the following assumptions:

1. The expert policy and true dynamics (π^*, f) are (C_π, C_{reg}) -smooth, respectively.
2. There is some $C_K > 0$ where $\max_t \|\nabla_{\mathbf{x}} \pi^*(\mathbf{x}_t^{\pi^*})\|_{\text{op}} \leq C_K$ holds w.p. 1 over $\mathbf{x}_1^{\pi^*} \sim \mathbb{P}_{\mathbf{x}_1^{\pi^*}}$.
3. The closed-loop system induced by (π^*, f) is (C_{ISS}, ρ) -EISS ([Definition 2.1](#)).

There are two overall philosophies of understanding the merits of data augmentation.

The RL-theoretic perspective. RL-theoretic notions of exploration often take an information-theoretic flavor, where it is captured by notions of “coverage” [[Jin et al., 2021](#), [Zhan et al., 2022](#), [Amortila et al., 2024](#), [Jiang and Xie, 2024](#)]. Coverage analyses rely on density ratios and thus the existence of densities. In continuous state-action spaces, expert (deterministic) policies typically do not have densities, and thus they can be induced by incorporating (possibly shaped) noise to the actions [[Haarnoja et al., 2018](#),

⁵Our results hold for generic bounded noise, but it suffices to consider $\mathbf{z} \sim \text{Unif}(\mathbb{B}^{d_u}(1))$ or $\text{Unif}(\mathbb{S}^{d_u}(1))$.

Schulman et al., 2017]. Crucially, this makes the policy itself noisy—compare this to [Intervention 2](#), where the expert’s recorded action is uncorrupted. When the noise is Gaussian, this practice ensures that the action distribution at a given state admits a Radon-Nikodym derivative with respect to the Lebesgue measure, and maximum-likelihood estimation (MLE) amounts to minimizing square error. Hence, existing analyses of behavior cloning (e.g. via the log-loss [\[Foster et al., 2024\]](#), which reduces to a square loss under Gaussian noising) ensure consistent imitation.

However, this comes at the price of corrupting the demonstrations provided to the learner, which in turn, we show in [Appendix B.4](#), leads to suboptimal rates of estimation. In particular, by reducing imitation learning to MLE over noisy data, the performance of IL is dictated by the capacity of the *stochastic* policy class, as measured by a covering number $\mathcal{N}_{\log}(\Pi, \varepsilon)$ under, e.g., the log-loss. For σ_u -scaled Gaussian noise, this equates to covering under the Euclidean norm at resolution $\approx \sqrt{\sigma_u^2 \varepsilon}$. For non-parametric classes—such as the lower bound constructions leading to [Theorem A](#), this can introduce additional polynomial factors of σ_u^{-1} in the estimation error. These factors of σ_u^{-1} must then be traded off with the error induced by imitating a noisy expert rather than the true expert labels.

In fact, there is a more fundamental problem. Due to the non-linearity of the dynamics, the noisy trajectory distribution suffers a drift from the noiseless one. This drift means policies fitted on the noise-injected trajectories, even with clean actions labeled, i.e., $\mathbb{P}_{\pi^*, \sigma_u}$, necessarily accrue an additive error scaling with σ_u , regardless of the on-expert error, summarized in the following lower bound.

Proposition 4.1 (Drift lower bound, informal). *For any given $\sigma_u > 0$ and $C_\pi > 0$, there exists a pair of two C_π -smooth policies π_1, π_2 such that one trajectory from the rollout distribution under each can distinguish them perfectly, but given trajectories with σ_u^2 -noise injection, any learning algorithm will yield a policy π that incurs $\Omega(C_\pi^2 \sigma_u^4)$ trajectory error, as measured by $\mathbf{J}_{\text{TRAJ}, T}$.*

The formal statement and set-up of [Proposition 4.1](#) is found in [Appendix B.4](#). We notice that this bound scales with C_π , indicating that smoothness is a key quantity in any argument based on noising. As we will see, this lower bound is circumvented by learning from trajectories *both* with and without noise injection, i.e., $\mathbb{P}_{\pi^*, \sigma_u, \alpha}$, as prescribed by [Intervention 2](#).

The control-theoretic perspective. In the control-theoretic literature, *persistence of excitation* (PE) is a well-established sufficient condition for ensuring parameter recovery in system-identification and adaptive control, which in turn yields performant policy synthesis [\[Bai and Sastry, 1985, Narendra and Annaswamy, 1987, Willems et al., 2005, Van Waarde et al., 2020\]](#). A input-sequence is “PE” if it yields a full-rank sequence of states, which guarantees parameter recovery across all modes the system may encounter. Therefore, when an expert policy may output degenerate trajectories in closed-loop,⁶ a natural approach to achieve PE is to inject excitatory noise into the inputs or directly into the system state [\[Annaswamy, 2023\]](#). More modern analyses of both the online linear-quadratic regulator (LQR) problem [\[Dean et al., 2018, Mania et al., 2019, Simchowitz and Foster, 2020\]](#) and of imitation learning [Pfrommer et al. \[2022\]](#) have similarly turned toward PE to ensure desirable learning behavior; the former relying on process noise (i.e., non-degenerate noise entering additively to the state) to excite state variables, the latter assuming the ability to directly perturb states during expert demonstration. By contrast, our setting assumes neither the presence of process noise, nor direct access to the system state. Lastly, we do not even assume the system is *controllable*,⁷ i.e., we also cannot rely on input perturbations

⁶See e.g., cases for linear systems under an optimal LQR controller [\[Polderman, 1986, Lee et al., 2023\]](#).

⁷Informally the ability of a system to be steered from one state to another by applying appropriate control inputs, cf. [\[Kailath, 1980\]](#).




Figure 2: Normalized cumulative reward on the HalfCheetah-v5 environment, with 40 total training trajectories. **Left:** fixing $\sigma_u = 1$, we vary the proportion of expert trajectories $\alpha \in [0, 1]$. We see having *no* noised-injected trajectories is highly suboptimal, but having any number suffices for improved performance, with a minor dip when using *only* noised trajectories (see Proposition 4.1). **Right:** we compare collecting clean action labels as in Intervention 2, versus the noised ones as in an coverage-based approach. We note $\sigma_u = 1$ corresponds to sizable entry-wise input perturbations ≈ 0.4 on an action space of $[-1, 1]^6$. Imitating with noisy labels is therefore catastrophic, yet using clean labels achieves improved performance. Experiment details in Appendix C.

inducing the PE condition. A key object is the controllability Gramian induced by linearizations around the expert trajectory.

Definition 4.2 (Controllability Gramian). Fixing $\mathbf{x}_1^{\pi^*} \sim \mathbb{P}_{\mathbf{x}_1^{\pi^*}}$, and letting $\mathbf{x}_{t+1}^{\pi^*} = f^{\pi^*}(\mathbf{x}_t^{\pi^*}, \mathbf{u}_t^{\pi^*})$, $\mathbf{u}_t^{\pi^*} = \pi^*(\mathbf{x}_t^{\pi^*})$, define the linearizations $\mathbf{A}_t = \nabla_{\mathbf{x}} f(\mathbf{x}_t^{\pi^*}, \mathbf{u}_t^{\pi^*})$, $\mathbf{B}_t = \nabla_{\mathbf{u}} f(\mathbf{x}_t^{\pi^*}, \mathbf{u}_t^{\pi^*})$, $\mathbf{K}_t^{\pi^*} = \nabla_{\mathbf{x}} \pi^*(\mathbf{x}_t^{\pi^*})$, $\mathbf{A}_{s:t}^{\text{cl}} = (\mathbf{A}_{t-1} + \mathbf{B}_{t-1} \mathbf{K}_{t-1}^{\pi^*})(\mathbf{A}_{t-2} + \mathbf{B}_{t-2} \mathbf{K}_{t-2}^{\pi^*}) \cdots (\mathbf{A}_s + \mathbf{B}_s \mathbf{K}_s^{\pi^*})$. The t -step controllability Gramian is defined as: $\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}) \triangleq \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} \mathbf{B}_s \mathbf{B}_s^\top \mathbf{A}_{s+1:t}^{\text{cl}\top}$.

Let us first entertain the implications of attaining PE through Intervention 2 by assuming *one-step controllability*, where $\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}) \succeq \underline{\lambda}_{\mathbf{W}} \mathbf{I}_{d_u}$, $\underline{\lambda}_{\mathbf{W}} > 0$, for all $t \geq 2$, such that under an appropriate input sequence, the (linearized) expert system f^{π^*} can reach any state at any time. Therefore, noise-injection will excite all modes of the linearized system, translating to PE as traditional control theory would suggest. This yields the following (suboptimal) bound when imitating over $\mathbb{P}_{\pi^*, \sigma_u}$.

Suboptimal Proposition 4.2. Let Assumption 4.1 hold, and let $\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}) \succeq \underline{\lambda}_{\mathbf{W}} \mathbf{I}_{d_u}$, $t \geq 2$ w.p. 1 over $\mathbf{x}_1^{\pi^*} \sim D$ for some $\underline{\lambda}_{\mathbf{W}} > 0$. Let $\hat{\pi}$ be a C_π -smooth candidate policy. For σ_u^2 that satisfies $\sigma_u^2 \lesssim O_*(\text{poly}(1/C_\pi, 1/C_{\text{reg}})) \underline{\lambda}_{\mathbf{W}}$, we have:

$$\mathbf{J}_{\text{TRAJ}, T}(\hat{\pi}) \lesssim O_*(T) \underline{\lambda}_{\mathbf{W}}^{-1} \left(\frac{1}{\sigma_u^2} \mathbf{J}_{\text{DEMO}, T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_u}) + C_\pi^2 C_{\text{stab}}^2 \sigma_u^2 \right).$$

The full statement and proof can be found in Appendix B.1. Though this bound avoids exponential-in- T compounding trajectory error, it has several shortcomings. Besides the strictness of one-step controllability (or controllability at all), the bound suffers: 1. a drift term that scales as σ_u^2 , which is even worse than Proposition 4.1 suggests, 2. the requirement on σ_u and resulting bound scaling with $\underline{\lambda}_{\mathbf{W}}$, which is steep for Gramians with fast-decaying spectra. So far, the direct control-theoretic approach provides worse guarantees than the information-theoretic RL one (see Appendix B.4 for details). As such, a combination of algorithmic (e.g., $\mathbb{P}_{\pi^*, \sigma_u} \rightarrow \mathbb{P}_{\pi^*, \sigma_u, \alpha}$) and analytical innovations are required to improve Suboptimal Proposition 4.2.

A new analysis of data augmentation. Though both the RL and control perspectives contain vital takeaways, neither in isolation leads to qualitatively correct assumptions or protocols for imitation over

nonlinear dynamical systems. In light of [Suboptimal Proposition 4.2](#), we make a few key observations. Firstly, compounding errors are not arbitrary state perturbations: they result from policy differences, and thus enter via the input channels. For smooth systems, this implies the trajectory error is primarily contained in the *controllable subspace* $\text{range}(\mathbf{W}_{1:t}^u)$. However, nonlinearity in the dynamics f means error will leak outside of $\text{range}(\mathbf{W}_{1:t}^u)$, which would seem to require PE to detect. Our first key insight is that, as long as we enforce low error on the controllable subspace, the nonlinear error automatically regulates itself.

Proposition 4.3. *Let [Assumption 4.1](#) hold, and assume the candidate policy $\hat{\pi}$ is C_π -smooth. Fix $\mathbf{x}_1^{\hat{\pi}} = \mathbf{x}_1^{\pi^*} = \mathbf{x}_1$, and define $\mathcal{R}_t^{\pi^*} \triangleq \text{range}(\mathbf{W}_{1:t}^u)$. Then for any given $\varepsilon \in [0, 1]$, $T \in \mathbb{N}$, as long as:*

$$\max_{1 \leq t \leq T-1} \sup_{\|\mathbf{v}\| \leq 1, \|\mathbf{w}\| \leq 1, \mathbf{w} \in \mathcal{R}_t^{\pi^*}} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \varepsilon \mathbf{w} + O_*(1) \varepsilon^2 \mathbf{v})\| \leq c_{\text{stab}} \varepsilon,$$

we are guaranteed $\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^}\| \leq \varepsilon$.*

This result shows that if we ensure the “generic” error \mathbf{v} term scales as ε^2 , regularity of $\hat{\pi}, \pi^*$ automatically ensures its contribution to $\hat{\pi} - \pi^*$ is $o(c_{\text{stab}} \varepsilon)$. For smooth systems the nonlinear error is indeed higher-order. However, it remains to control the error term lying in $\mathcal{R}_t^{\pi^*}$, leading to another weakness of [Suboptimal Proposition 4.2](#) in the dependence on the smallest (positive) eigenvalues of $\mathbf{W}_{1:t}^u$. This is unintuitive: small eigendirections of $\mathbf{W}_{1:t}^u$ are those that are hard to excite. In contrast to objectives like parameter recovery, we do not need uniform detection of all directions. In fact, errors should compound slowly on hard-to-excite directions, such that we may safely “ignore” them. Restricting our attention to excitable directions means we only pay for level of excitation we need.

Proposition 4.4. *Let [Assumption 4.1](#) hold. For $\mathbf{x}_1^{\pi^*} \sim D$, let $\{(\lambda_{i,t}, \mathbf{v}_{i,t})\}_{i=1}^{d_x}$ be the eigenvalues and vectors of $\mathbf{W}_{1:t}^u$, $t \geq 2$. Define $\mathcal{R}_t^{\pi^*}(\lambda) \triangleq \text{span}\{\mathbf{v}_{i,t} : \lambda_{i,t} \geq \lambda\}$ and $\mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)}$ the corresponding orthogonal projection. Recall $\mathbb{P}_{\pi^*, \sigma_u, \alpha}$ and set $\alpha = 0.5$. Then, for $\sigma_u \lesssim O_*(\lambda)$, we have:*

$$\mathbb{E}_{\mathbb{P}_{\pi^*}} \|\mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)} \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|_{\text{op}}^2 \lesssim \frac{1}{\sigma_u^2 \lambda} \mathbb{E}_{\mathbb{P}_{\pi^*, \sigma_u, \alpha}} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|^2 + \frac{\sigma_u^2}{\lambda} C_\pi^2 C_{\text{stab}}^4.$$

This is precisely where our *algorithmic* prescription arises: certifying the learned policy $\hat{\pi}$ matches π^* up to first-order on $\mathcal{R}_t^{\pi^*}(\lambda)$ requires data both at $\mathbf{x}_t^{\pi^*}$ and around it (e.g. via noise-injection). Practically, this translates to imitating on the mixture distribution $\mathbb{P}_{\pi^*, \sigma_u, \alpha}$. Therefore, combining [Proposition 4.3](#), which translates imitating π^* well in a neighborhood to low trajectory error, with [Proposition 4.4](#), which guarantees imitating on $\mathbb{P}_{\pi^*, \sigma_u, \alpha}$ matches $\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)$ up to a flexible excitation level, leads to our main guarantee of [Intervention 2](#).

Theorem 2. *Let [Assumption 4.1](#) hold. Let $\hat{\pi}$ be a L_π -Lipschitz, C_π -smooth policy. Then, for $\sigma_u \lesssim O_*(\text{poly}(1/C_\pi, 1/C_{\text{reg}}))$, we have:*

$$\mathbf{J}_{\text{TRAJ}, T}(\hat{\pi}) \lesssim O_*(T) \sigma_u^{-2} \mathbf{J}_{\text{DEMO}, T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_u, 0.5}).$$

Comparisons to the RL and control perspectives. By combining ideas from RL and control, we arrive at conclusions that may be surprising from either perspective. Compared to the RL perspective, 1. we do not have coverage in the usual sense, 2. we avoid accumulating mean-estimation error from imitating noisy action labels, 3. using the mixture distribution $\mathbb{P}_{\pi^*, \sigma_u, \alpha}$ subverts the additive σ_u^4 error in [Proposition 4.1](#). On the control-theoretic side, 1. imitating over $\mathbb{P}_{\pi^*, \sigma_u, \alpha}$ removes the additive σ_u^2 error in [Suboptimal Proposition 4.2](#), 2. we avoid any assumption of controllability *as well as* any dependence

on the small eigendirections of the controllable subspace. In fact, by removing any additive σ_u factor, our bound suggests that we should set the noise-scale σ_u as large as permissible! We emphasize these theoretical predictions are tangible, which we display across [Figure 1](#) and [Figure 2](#), demonstrating: 1. comparable performance gains across many scales of σ_u , 2. the importance of using both clean and noised trajectories, 3. the suboptimality of noising the action *labels*.

5 Discussion and Limitations

Our action-chunking guarantees rely on a structural assumption of $(\hat{r}, \hat{f}) \in \mathcal{P}$ being an EISS pair. We believe either explicitly enforcing this, e.g., via regularization [[Sindhwani et al., 2018](#), [Mehta et al., 2025](#)] or hierarchy [[Matni et al., 2024](#)], or attaining it indirectly via implicit biases [[Chi et al., 2023](#)], are interesting directions of inquiry. We assume smoothness in [Section 4](#), which is not strictly satisfied in some applications, such as in model-predictive control [[Garcia et al., 1989](#)]. We remark our lower bound [Proposition 4.1](#) depends on smoothness in C_π , which implies it is in some sense a fundamental aspect of noise-injection. However, we believe our results should extend to piece-wise notions [[Block et al., 2023](#)], and note ongoing research exploring *smoothing* for learning in dynamical systems [[Suh et al., 2022](#), [Pang et al., 2023](#), [Pfrommer et al., 2024](#)]. In general, we leave a sharp characterization of the role of smoothness and control-theoretic quantities in IL as an open problem. Lastly, we leave investigating iterative interaction (e.g., DAGGER [[Ross et al., 2011](#)]) as future work.

Acknowledgments

TZ gratefully acknowledges a gift from AWS AI to Penn Engineering’s ASSET Center for Trustworthy AI. TZ and NM are supported in part by NSF Award SLES-2331880, NSF CAREER award ECCS-2045834, NSF EECS-2231349, and AFOSR Award FA9550-24-1-0102.

References

- Radosław Adamczak, Rafał Latała, Alexander E Litvak, Krzysztof Oleszkiewicz, Alain Pajor, and Nicole Tomczak-Jaegermann. A short proof of paouris’ inequality. *Canadian Mathematical Bulletin*, 57(1): 3–8, 2014.
- Philip Amortila, Dylan J Foster, Nan Jiang, Ayush Sekhari, and Tengyang Xie. Harnessing density ratios for online reinforcement learning. *arXiv preprint arXiv:2401.09681*, 2024.
- David Angeli. A lyapunov approach to incremental stability properties. *IEEE Transactions on Automatic Control*, 47(3):410–421, 2002.
- Anuradha M Annaswamy. Adaptive control and intersections with reinforcement learning. *Annual Review of Control, Robotics, and Autonomous Systems*, 6(1):65–93, 2023.
- Er-Wei Bai and Sosale Shankara Sastry. Persistency of excitation, sufficient richness and parameter convergence in discrete time adaptive control. *Systems & control letters*, 6(3):153–163, 1985.
- Mayank Bansal, Alex Krizhevsky, and Abhijit Ogale. Chauffeurnet: Learning to drive by imitating the best and synthesizing the worst. *arXiv preprint arXiv:1812.03079*, 2018.

- Kevin Black, Noah Brown, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Lachy Groom, Karol Hausman, Brian Ichter, et al. *pi₀*: A vision-language-action flow model for general robot control. *arXiv preprint arXiv:2410.24164*, 2024.
- Adam Block, Max Simchowitz, and Alexander Rakhlin. Oracle-efficient smoothed online learning for piecewise continuous decision making. In *The Thirty Sixth Annual Conference on Learning Theory*, pages 1618–1665. PMLR, 2023.
- Adam Block, Ali Jadbabaie, Daniel Pfrommer, Max Simchowitz, and Russ Tedrake. Provable guarantees for generative behavior cloning: Bridging low-level stability and high-level behavior. *Advances in Neural Information Processing Systems*, 2024.
- Mariusz Bojarski, Davide Del Testa, Daniel Dworakowski, Bernhard Firner, Beat Flepp, Prasoon Goyal, Lawrence D Jackel, Mathew Monfort, Urs Muller, Jiakai Zhang, et al. End to end learning for self-driving cars. *arXiv preprint arXiv:1604.07316*, 2016.
- Lili Chen, Kevin Lu, Aravind Rajeswaran, Kimin Lee, Aditya Grover, Misha Laskin, Pieter Abbeel, Aravind Srinivas, and Igor Mordatch. Decision transformer: Reinforcement learning via sequence modeling. *Advances in neural information processing systems*, 34:15084–15097, 2021.
- Cheng Chi, Siyuan Feng, Yilun Du, Zhenjia Xu, Eric Cousineau, Benjamin Burchfiel, and Shuran Song. Diffusion policy: Visuomotor policy learning via action diffusion. *arXiv preprint arXiv:2303.04137*, 2023.
- Sarah Dean, Horia Mania, Nikolai Matni, Benjamin Recht, and Stephen Tu. Regret bounds for robust adaptive control of the linear quadratic regulator. *Advances in Neural Information Processing Systems*, 31, 2018.
- Chelsea Finn, Tianhe Yu, Tianhao Zhang, Pieter Abbeel, and Sergey Levine. One-shot visual imitation learning via meta-learning. In *Conference on robot learning*, pages 357–368. PMLR, 2017.
- Dylan J Foster, Adam Block, and Dipendra Misra. Is behavior cloning all you need? understanding horizon in imitation learning. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- Carlos E Garcia, David M Prett, and Manfred Morari. Model predictive control: Theory and practice—a survey. *Automatica*, 25(3):335–348, 1989.
- Xavier Glorot and Yoshua Bengio. Understanding the difficulty of training deep feedforward neural networks. In *Proceedings of the thirteenth international conference on artificial intelligence and statistics*, pages 249–256. JMLR Workshop and Conference Proceedings, 2010.
- Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine. Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. In *International conference on machine learning*, pages 1861–1870. Pmlr, 2018.
- Dan Hendrycks and Kevin Gimpel. Gaussian error linear units (gelus). *arXiv preprint arXiv:1606.08415*, 2016.
- Michael Hertneck, Johannes Köhler, Sebastian Trimpe, and Frank Allgöwer. Learning an approximate model predictive controller with guarantees. *IEEE Control Systems Letters*, 2(3):543–548, 2018.

- Roger A Horn and Charles R Johnson. *Matrix analysis*. Cambridge university press, 2012.
- Ahmed Hussein, Mohamed Medhat Gaber, Eyad Elyan, and Chrisina Jayne. Imitation learning: A survey of learning methods. *ACM Computing Surveys (CSUR)*, 50(2):1–35, 2017.
- Nan Jiang and Tengyang Xie. Offline reinforcement learning in large state spaces: Algorithms and guarantees. 2024.
- Ying Jin, Zhuoran Yang, and Zhaoran Wang. Is pessimism provably efficient for offline rl? In *International Conference on Machine Learning*, pages 5084–5096. PMLR, 2021.
- Thomas Kailath. *Linear systems*, volume 156. Prentice-Hall Englewood Cliffs, NJ, 1980.
- Sham Kakade, Akshay Krishnamurthy, Kendall Lowrey, Motoya Ohnishi, and Wen Sun. Information theoretic regret bounds for online nonlinear control. *Advances in Neural Information Processing Systems*, 33:15312–15325, 2020.
- Liyiming Ke, Jingqiang Wang, Tapomayukh Bhattacharjee, Byron Boots, and Siddhartha Srinivasa. Grasping with chopsticks: Combating covariate shift in model-free imitation learning for fine manipulation. In *2021 IEEE International Conference on Robotics and Automation (ICRA)*, pages 6185–6191. IEEE, 2021.
- Liyiming Ke, Yunchu Zhang, Abhay Deshpande, Siddhartha Srinivasa, and Abhishek Gupta. Ccil: Continuity-based data augmentation for corrective imitation learning. In *The Twelfth International Conference on Learning Representations*, 2024.
- Michael Kelly, Chelsea Sidrane, Katherine Driggs-Campbell, and Mykel J Kochenderfer. Hg-dagger: Interactive imitation learning with human experts. In *2019 International Conference on Robotics and Automation (ICRA)*, pages 8077–8083. IEEE, 2019.
- Michael Laskey, Jonathan Lee, Roy Fox, Anca Dragan, and Ken Goldberg. Dart: Noise injection for robust imitation learning. In *Conference on robot learning*, pages 143–156. PMLR, 2017.
- Bruce D Lee, Ingvar Ziemann, Anastasios Tsiamis, Henrik Sandberg, and Nikolai Matni. The fundamental limitations of learning linear-quadratic regulators. In *2023 62nd IEEE Conference on Decision and Control (CDC)*, pages 4053–4060. IEEE, 2023.
- Yuejiang Liu, Jubayer Ibn Hamid, Annie Xie, Yoonho Lee, Max Du, and Chelsea Finn. Bidirectional decoding: Improving action chunking via closed-loop resampling. In *The Thirteenth International Conference on Learning Representations*, 2025.
- I Loshchilov. Decoupled weight decay regularization. *arXiv preprint arXiv:1711.05101*, 2017.
- Ilya Loshchilov and Frank Hutter. Sgdr: Stochastic gradient descent with warm restarts. *arXiv preprint arXiv:1608.03983*, 2016.
- Horia Mania, Stephen Tu, and Benjamin Recht. Certainty equivalence is efficient for linear quadratic control. *Advances in Neural Information Processing Systems*, 32, 2019.
- Nikolai Matni, Aaron D Ames, and John C Doyle. A quantitative framework for layered multirate control: Toward a theory of control architecture. *IEEE Control Systems Magazine*, 44(3):52–94, 2024.

- Shaunak A Mehta, Yusuf Umut Ciftci, Balamurugan Ramachandran, Somil Bansal, and Dylan P Losey. Stable-bc: Controlling covariate shift with stable behavior cloning. *IEEE Robotics and Automation Letters*, 2025.
- Kumpati S Narendra and Anuradha M Annaswamy. Persistent excitation in adaptive systems. *International Journal of Control*, 45(1):127–160, 1987.
- Tao Pang, HJ Terry Suh, Lujie Yang, and Russ Tedrake. Global planning for contact-rich manipulation via local smoothing of quasi-dynamic contact models. *IEEE Transactions on robotics*, 39(6):4691–4711, 2023.
- Grigoris Paouris. Concentration of mass on convex bodies. *Geometric & Functional Analysis GAFA*, 16(5):1021–1049, 2006.
- Daniel Pfrommer, Thomas Zhang, Stephen Tu, and Nikolai Matni. Tasil: Taylor series imitation learning. *Advances in Neural Information Processing Systems*, 35:20162–20174, 2022.
- Daniel Pfrommer, Swati Padmanabhan, Kwangjun Ahn, Jack Umenberger, Tobia Marcucci, Zakaria Mhammedi, and Ali Jadbabaie. Improved sample complexity of imitation learning for barrier model predictive control. *arXiv preprint arXiv:2410.00859*, 2024.
- Jan Willem Polderman. On the necessity of identifying the true parameter in adaptive lq control. *Systems & control letters*, 8(2):87–91, 1986.
- Dean A Pomerleau. Alvin: An autonomous land vehicle in a neural network. *Advances in neural information processing systems*, 1, 1988.
- Antonin Raffin, Ashley Hill, Adam Gleave, Anssi Kanervisto, Maximilian Ernestus, and Noah Dormann. Stable-baselines3: Reliable reinforcement learning implementations. *Journal of Machine Learning Research*, 22(268):1–8, 2021.
- Stéphane Ross and Drew Bagnell. Efficient reductions for imitation learning. In *Proceedings of the thirteenth international conference on artificial intelligence and statistics*, pages 661–668. JMLR Workshop and Conference Proceedings, 2010.
- Stéphane Ross, Geoffrey Gordon, and Drew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the fourteenth international conference on artificial intelligence and statistics*, pages 627–635. JMLR Workshop and Conference Proceedings, 2011.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Nur Muhammad Shafiullah, Zichen Cui, Ariuntuya Arty Altanzaya, and Lerrel Pinto. Behavior transformers: Cloning k modes with one stone. *Advances in neural information processing systems*, 35: 22955–22968, 2022.
- Max Simchowitz and Dylan Foster. Naive exploration is optimal for online lqr. In *International Conference on Machine Learning*, pages 8937–8948. PMLR, 2020.

- Max Simchowitz, Horia Mania, Stephen Tu, Michael I Jordan, and Benjamin Recht. Learning without mixing: Towards a sharp analysis of linear system identification. In *Conference On Learning Theory*, pages 439–473. PMLR, 2018.
- Max Simchowitz, Daniel Pfrommer, and Ali Jadbabaie. The pitfalls of imitation learning when actions are continuous. *arXiv preprint arXiv:2503.09722*, 2025.
- Vikas Sindhwani, Stephen Tu, and Mohi Khansari. Learning contracting vector fields for stable imitation learning. *arXiv preprint arXiv:1804.04878*, 2018.
- Elias M Stein and Rami Shakarchi. *Functional analysis: introduction to further topics in analysis*, volume 4. Princeton University Press, 2011.
- Hyung Ju Suh, Max Simchowitz, Kaiqing Zhang, and Russ Tedrake. Do differentiable simulators give better policy gradients? In *International Conference on Machine Learning*, pages 20668–20696. PMLR, 2022.
- Xiatao Sun, Shuo Yang, and Rahul Mangharam. Mega-dagger: Imitation learning with multiple imperfect experts. *arXiv preprint arXiv:2303.00638*, 2023.
- Siyu Teng, Xuemin Hu, Peng Deng, Bai Li, Yuchen Li, Yunfeng Ai, Dongsheng Yang, Lingxi Li, Zhe Xuanyuan, Fenghua Zhu, et al. Motion planning for autonomous driving: The state of the art and future perspectives. *IEEE Transactions on Intelligent Vehicles*, 8(6):3692–3711, 2023.
- Mark Towers, Ariel Kwiatkowski, Jordan Terry, John U Balis, Gianluca De Cola, Tristan Deleu, Manuel Goulao, Andreas Kallinteris, Markus Krimmel, Arjun KG, et al. Gymnasium: A standard interface for reinforcement learning environments. *arXiv preprint arXiv:2407.17032*, 2024.
- Duc N Tran, Björn S Rüffer, and Christopher M Kellett. Incremental stability properties for discrete-time systems. In *2016 IEEE 55th Conference on Decision and Control (CDC)*, pages 477–482. IEEE, 2016.
- Stephen Tu, Alexander Robey, Tingnan Zhang, and Nikolai Matni. On the sample complexity of stability constrained imitation learning. In *Learning for Dynamics and Control Conference*, pages 180–191. PMLR, 2022.
- Henk J Van Waarde, Claudio De Persis, M Kanat Camlibel, and Pietro Tesi. Willems’ fundamental lemma for state-space systems and its extension to multiple datasets. *IEEE Control Systems Letters*, 4(3):602–607, 2020.
- Cédric Villani et al. *Optimal transport: old and new*, volume 338. Springer, 2009.
- Jan C Willems, Paolo Rapisarda, Ivan Markovsky, and Bart LM De Moor. A note on persistency of excitation. *Systems & Control Letters*, 54(4):325–329, 2005.
- He Yin, Peter Seiler, Ming Jin, and Murat Arcak. Imitation learning with stability and safety guarantees. *IEEE Control Systems Letters*, 6:409–414, 2021.
- Wenhao Zhan, Baihe Huang, Audrey Huang, Nan Jiang, and Jason Lee. Offline reinforcement learning with realizability and single-policy concentrability. In *Conference on Learning Theory*, pages 2730–2775. PMLR, 2022.

- Tianhao Zhang, Zoe McCarthy, Owen Jow, Dennis Lee, Xi Chen, Ken Goldberg, and Pieter Abbeel. Deep imitation learning for complex manipulation tasks from virtual reality teleoperation. In *2018 IEEE International Conference on Robotics and Automation (ICRA)*, pages 5628–5635. IEEE, 2018.
- Siyan Zhao and Aditya Grover. Decision stacks: Flexible reinforcement learning via modular generative models. *arXiv preprint arXiv:2306.06253*, 2023.
- Tony Z Zhao, Vikash Kumar, Sergey Levine, and Chelsea Finn. Learning fine-grained bimanual manipulation with low-cost hardware. *arXiv preprint arXiv:2304.13705*, 2023.
- Ingvar Ziemann and Stephen Tu. Learning with little mixing. *Advances in Neural Information Processing Systems*, 35:4626–4637, 2022.
- Brianna Zitkovich et al. Rt-2: Vision-language-action models transfer web knowledge to robotic control. In Jie Tan, Marc Toussaint, and Kourosh Darvish, editors, *Proceedings of The 7th Conference on Robot Learning*, volume 229 of *Proceedings of Machine Learning Research*, pages 2165–2183. PMLR, 06–09 Nov 2023.

A Proofs and Additional Details for Section 3

We first introduce some additional definitions:

Definition A.1 (Additional Error Definitions). Given $p \geq 1$, define the p -th power errors:

$$\begin{aligned}\mathbf{J}_{\text{TRAJ},p,T}(\hat{\pi}) &\triangleq \mathbb{E}_{\hat{\pi},\pi^*} \left[\sum_{t=1}^T \min \{1, \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^p + \|\mathbf{u}_t^{\hat{\pi}} - \mathbf{u}_t^{\pi^*}\|^p\} \right] \\ \mathbf{J}_{\text{DEMO},p,T}(\hat{\pi}; \mathbb{P}_{\text{demo}}) &\triangleq \mathbb{E}_{\mathbb{P}_{\text{demo}}} \left[\sum_{t=1}^T \|\hat{\pi}(\mathbf{x}_{1:t}, \mathbf{u}_{1:t-1}, t) - \mathbf{u}_t^{(i)}\|^p \right].\end{aligned}$$

Note that $\mathbf{J}_{\text{TRAJ},2,T} \equiv \mathbf{J}_{\text{TRAJ},T}$, $\mathbf{J}_{\text{DEMO},p,T} \equiv \mathbf{J}_{\text{DEMO},T}$. We further define the trajectory state error:

$$\mathbf{J}_{\text{TRAJ},p,T}^{\mathbf{x}}(\hat{\pi}) \triangleq \mathbb{E}_{\hat{\pi},\pi^*} \left[\sum_{t=1}^T \min \{1, \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^p\} \right].$$

We now state some elementary results.

Lemma A.1. Assume $\hat{\pi}$ is a Markovian, L_{π} -Lipschitz policy. Then:

$$\mathbf{J}_{\text{TRAJ},p,T}(\hat{\pi}) \leq (1 + (2L_{\pi})^p) \mathbf{J}_{\text{TRAJ},p,T}^{\mathbf{x}}(\hat{\pi}) + 2^p \mathbf{J}_{\text{DEMO},p,T}(\hat{\pi}; \mathbb{P}_{\pi^*}).$$

Proof. Following the definition of $\mathbf{J}_{\text{TRAJ},p,T}$, we may add and subtract $\|\mathbf{u}_t^{\hat{\pi}} - \hat{\pi}(\mathbf{x}_t^{\pi^*}) + \hat{\pi}(\mathbf{x}_t^{\pi^*}) - \mathbf{u}_t^{\pi^*}\|^p$ and apply convexity of $\|\cdot\|^p$ to yield:

$$\mathbf{J}_{\text{TRAJ},T}(\hat{\pi}) \leq \mathbb{E}_{\hat{\pi},\pi^*} \left[\sum_{t=1}^T \min \{1, \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^p + 2^p \|\hat{\pi}(\mathbf{x}_t^{\hat{\pi}}) - \hat{\pi}(\mathbf{x}_t^{\pi^*})\|^p + 2^p \|\hat{\pi}(\mathbf{x}_t^{\pi^*}) - \pi^*(\mathbf{x}_t^{\pi^*})\|^p\} \right].$$

Applying Lipschitzness of $\hat{\pi}$ to the second term, and observing the last term is precisely the summand in $\mathbf{J}_{\text{DEMO},p,T}$ completes the proof. \square

Lemma A.2 (Kantorovich-Rubinstein). Define the norm on $\mathbb{R}_x^d \times \mathbb{R}_u^d$: $\|(\mathbf{x}, \mathbf{u})\|_{\mathbf{x}, \mathbf{u}} \triangleq \|\mathbf{x}\| + \|\mathbf{u}\|$. Then, define the class of cost functions $c(\mathbf{x}, \mathbf{u}) \in \mathcal{C}_{\text{Lip}(1)}$ that is 1-Lipschitz in $\|\cdot\|_{\mathbf{x}, \mathbf{u}}$. Then, we have the following:

$$\mathbf{J}_{\text{TRAJ},1,T}(\hat{\pi}) \leq \sup_{c_{1:T} \in \mathcal{C}_{\text{Lip}(1)}} \mathbf{J}_{\text{COST},T}(\hat{\pi}; c_{1:T}).$$

The above is a straightforward application of Kantorovich-Rubinstein strong duality [Villani et al., 2009] by pulling out a conditional expectation over $\mathbf{x}_1^{\pi^*}, \mathbf{x}_1^{\hat{\pi}} = \mathbf{x}_1$ on both sides. The inequality then follows due to the clipping at 1 in the definition of $\mathbf{J}_{\text{TRAJ},1,T}$.

We will often use the following bound on triangular Toeplitz matrices.

Lemma A.3. Given $\rho \in [0, 1)$, define the matrix:

$$\mathbf{A}(\rho) = \begin{bmatrix} 1 & 0 & 0 & \cdots & 0 \\ \rho & 1 & 0 & \cdots & 0 \\ \rho^2 & \rho & 1 & \cdots & 0 \\ \vdots & & & & \\ \rho^{d-1} & \rho^{d-2} & \rho^{d-3} & \cdots & 1 \end{bmatrix} \in \mathbb{R}^{d \times d}.$$

Given $1 \leq p < \infty$, the following bound holds on the induced $\ell^p \rightarrow \ell^p$ operator norm of $\mathbf{A}(\rho)$:

$$\|\mathbf{A}(\rho)\|_{p \rightarrow p} \leq \sum_{s=0}^{d-1} \rho^s \leq \frac{1}{1-\rho}.$$

Proof. We may prove this straightforwardly from an application of the Riesz-Thorin interpolation theorem [Stein and Shakarchi, 2011], which states that fixing $\mathbf{A}(\rho)$, the mapping $1/p \mapsto \|\mathbf{A}(\rho)\|_{p \rightarrow p}$ is log-convex for $p \in [1, \infty)$. In particular, by taking the convex combination $1/p = 1 \cdot (1/p) + 0 \cdot (1-1/p)$, we find:

$$\begin{aligned} \log \|\mathbf{A}(\rho)\|_{p \rightarrow p} &\leq (1/p) \log \|\mathbf{A}(\rho)\|_{1 \rightarrow 1} + (1 - 1/p) \log \|\mathbf{A}(\rho)\|_{\infty \rightarrow \infty} \\ \|\mathbf{A}(\rho)\|_{p \rightarrow p} &\leq \|\mathbf{A}(\rho)\|_{1 \rightarrow 1}^{1/p} \|\mathbf{A}(\rho)\|_{\infty \rightarrow \infty}^{1-1/p}. \end{aligned}$$

We then utilize the basic fact that $\|\mathbf{A}(\rho)\|_{1 \rightarrow 1}$ and $\|\mathbf{A}(\rho)\|_{\infty \rightarrow \infty}$ correspond to the maximum column and row sum of $\mathbf{A}(\rho)$, respectively, which completes the result. \square

We may now state the detailed version of Proposition 3.1.

Proposition A.4 (Full ver. of Proposition 3.1). *Let Assumption 3.1 hold. Let $(\hat{\pi}, \hat{f})$ be a policy-dynamics pair that is (C_{ISS}, ρ) -EISS, and consider the corresponding chunked policy $\tilde{\pi} = \text{chunked}(\hat{\pi}, \hat{f}, \ell)$. Then the closed-loop system the chunked policy induces on the true dynamics $(\tilde{\pi}, f)$ is (\tilde{C}, ρ^{1-a}) -EISS, where $a \in (0, 1)$ and $\tilde{C} = a^{-1} \log(1/\rho)^{-1} \cdot \text{poly}(L_\pi, C_{\text{ISS}})$, as long as the chunk length is sufficiently long: $\ell \gtrsim a^{-1} \log(1/\rho)^{-1} \cdot \log(\text{poly}(L_\pi, \tilde{C}, 1/a))$.*

Proof of Proposition 3.1. Let us define the chunk-indexing shorthand $t_k \triangleq (k-1)\ell + 1$, such that $t_1 = 1$. Toward establishing EISS of the closed-loop chunked system, we want to show for a sequence of input perturbations $\{\mathbf{u}_t\}_{t \geq 1}$ and two trajectories $\{\mathbf{x}_t^{\tilde{\pi}}\}_{t \geq 1}$, $\{\bar{\mathbf{x}}_t^{\tilde{\pi}}\}_{t \geq 1}$ evolving as:

$$\begin{aligned} \mathbf{x}_{t+1}^{\tilde{\pi}} &= f^{\tilde{\pi}}(\mathbf{x}_t^{\tilde{\pi}}, \mathbf{0}), & \mathbf{x}_1^{\tilde{\pi}} &= \mathbf{x}_1 \\ \bar{\mathbf{x}}_{t+1}^{\tilde{\pi}} &= f^{\tilde{\pi}}(\bar{\mathbf{x}}_t^{\tilde{\pi}}, \mathbf{u}_t), & \bar{\mathbf{x}}_1^{\tilde{\pi}} &= \bar{\mathbf{x}}_1, \end{aligned}$$

there exist some constants $C \geq 1$, $\rho \in (0, 1)$ such that:

$$\|\mathbf{x}_T^{\tilde{\pi}} - \bar{\mathbf{x}}_T^{\tilde{\pi}}\| \leq C\rho^{T-1} \|\mathbf{x}_1 - \bar{\mathbf{x}}_1\| + C \sum_{s=1}^{T-1} \rho^{T-1-s} \|\mathbf{u}_s\|.$$

To do so, we prove the following “contractivity” result going between chunks.

Lemma A.5. *Fix some $k \geq 1$. Recall the true dynamics f is (C_{ISS}, ρ) -EISS. Then, the following holds:*

$$\|\mathbf{x}_{t_{k+1}}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_{k+1}}^{\tilde{\pi}}\| \leq \rho^\ell \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}\| + C_{\text{ISS}} \sum_{s=0}^{\ell-1} \rho^{\ell-1-s} \|\mathbf{u}_{t_k+s}\|,$$

where $\rho \triangleq \rho^{1-a}$, as long as $\ell > a^{-1} \text{polylog}(1 + L_\pi, C_{\text{ISS}}, \log(1/\rho), a^{-1})$. As a corollary, setting $\bar{C} = \frac{(1+L_\pi)C_{\text{ISS}}^2}{3a\rho \log(1/\rho)}$, for $1 \leq h \leq \ell$, we have:

$$\|\mathbf{x}_{t_k+h}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k+h}^{\tilde{\pi}}\| \leq \bar{C}\rho^h \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}\| + C_{\text{ISS}} \sum_{s=0}^{h-1} \rho^{h-1-s} \|\mathbf{u}_{t_k+s}\|.$$

Proof of Lemma A.5. Applying (C_{ISS}, ρ) -EISS of the true dynamics f , we have

$$\begin{aligned} \|\mathbf{x}_{t_{k+1}}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_{k+1}}^{\tilde{\pi}}\| &\leq C_{\text{ISS}} \rho^\ell \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}\| + C_{\text{ISS}} \sum_{s=0}^{\ell-1} \rho^{\ell-1-s} \|\mathbf{u}_{t_k+s}^{\tilde{\pi}} - \bar{\mathbf{u}}_{t_k+s}^{\tilde{\pi}} - \mathbf{u}_{t_k+s}\| \\ &\leq C_{\text{ISS}} \rho^\ell \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}\| + C_{\text{ISS}} \sum_{s=0}^{\ell-1} \rho^{\ell-1-s} \|\mathbf{u}_{t_k+s}^{\tilde{\pi}} - \bar{\mathbf{u}}_{t_k+s}^{\tilde{\pi}}\| + C_{\text{ISS}} \sum_{s=0}^{\ell-1} \rho^{\ell-1-s} \|\mathbf{u}_{t_k+s}\|. \end{aligned}$$

where $\mathbf{u}_{t_k+s}^{\tilde{\pi}}$ and $\bar{\mathbf{u}}_{t_k+s}^{\tilde{\pi}}$ are the s -th actions outputted by the chunked policy $\tilde{\pi}$ conditioned on $\mathbf{x}_{t_k}^{\tilde{\pi}}$ and $\bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}$, respectively. We consider the “simulated” dynamical system that generates $\mathbf{u}^{\tilde{\pi}}, \bar{\mathbf{u}}^{\tilde{\pi}}$:

$$\begin{aligned} \mathbf{x}_{s+1} &= \hat{f}^{\hat{\pi}}(\mathbf{x}_s, \mathbf{0}) = \hat{f}(\mathbf{x}_s, \hat{\pi}(\mathbf{x}_s)), \quad s = 0, \dots, \ell-1 \\ \mathbf{u}_{t_k+s}^{\tilde{\pi}} &\triangleq \hat{\pi}(\mathbf{x}_s), \quad \mathbf{x}_0 = \mathbf{x}_{t_k}^{\tilde{\pi}}, \\ \tilde{\mathbf{x}}_{s+1} &= \hat{f}^{\hat{\pi}}(\tilde{\mathbf{x}}_s, \mathbf{0}) = \hat{f}(\tilde{\mathbf{x}}_s, \hat{\pi}(\tilde{\mathbf{x}}_s)), \quad s = 0, \dots, \ell-1 \\ \bar{\mathbf{u}}_{t_k+s}^{\tilde{\pi}} &\triangleq \hat{\pi}(\tilde{\mathbf{x}}_s), \quad \tilde{\mathbf{x}}_0 = \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}. \end{aligned}$$

Crucially, we observe that $(\hat{\pi}, \hat{f})$ is (C_{ISS}, ρ) -EISS, and thus:

$$\|\mathbf{x}_s - \tilde{\mathbf{x}}_s\| \leq C_{\text{ISS}} \rho^s \|\mathbf{x}_0 - \tilde{\mathbf{x}}_0\| = C_{\text{ISS}} \rho^s \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}\|.$$

Therefore, by the L_π -Lipschitzness of $\hat{\pi}$, we have:

$$\|\mathbf{u}_{t_k+s}^{\tilde{\pi}} - \bar{\mathbf{u}}_{t_k+s}^{\tilde{\pi}}\| \leq L_\pi \|\mathbf{x}_s - \tilde{\mathbf{x}}_s\| \leq L_\pi C_{\text{ISS}} \rho^s \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}\|.$$

Plugging this back above, we get:

$$\begin{aligned} \|\mathbf{x}_{t_{k+1}}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_{k+1}}^{\tilde{\pi}}\| &\leq C_{\text{ISS}} \rho^\ell \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}\| + L_\pi C_{\text{ISS}} \cdot C_{\text{ISS}} \sum_{s=0}^{\ell-1} \rho^{\ell-1-s} \rho^s \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}\| \\ &\quad + C_{\text{ISS}} \sum_{s=0}^{\ell-1} \rho^{\ell-1-s} \|\mathbf{u}_{t_k+s}\| \\ &\leq C_{\text{ISS}} \rho^\ell \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}\| + L_\pi C_{\text{ISS}}^2 \rho^{\ell-1} \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}\| + C_{\text{ISS}} \sum_{s=0}^{\ell-1} \rho^{\ell-1-s} \|\mathbf{u}_{t_k+s}\| \\ &\leq (1 + L_\pi) C_{\text{ISS}}^2 \rho^{\ell-1} \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_k}^{\tilde{\pi}}\| + C_{\text{ISS}} \sum_{s=0}^{\ell-1} \rho^{\ell-1-s} \|\mathbf{u}_{t_k+s}\|. \end{aligned}$$

We solve for the requisite chunk-length by solving: $(1 + L_\pi) C_{\text{ISS}}^2 \rho^{\ell-1} \leq \rho^\ell$, where $\rho = \rho^{1-\alpha}$. Rearranging, this amounts to $\ell \in \mathbb{N}$ satisfying:

$$a\ell \geq 1 + \frac{\log((1 + L_\pi) C_{\text{ISS}}^2 \ell)}{\log(1/\rho)}.$$

To remove the ℓ dependence on the right-hand side, we use the following elementary result.

Lemma A.6 (Cf. Simchowitz et al. [2018, Lemma A.4]). *Given $\alpha \geq 1$, for any $\ell \in \mathbb{N}$, $\ell \geq \alpha \log(\ell)$ as soon as $\ell \geq 2\alpha \log(4\alpha)$.*

We observe the above result holds if we add any term that does not depend on ℓ on the right-side of both inequalities. Applying it to the above, since $\log(1/\rho) \geq \log(e) = 1$, setting $\alpha = (a \log(1/\rho))^{-1}$, we have that

$$\ell \geq a^{-1} + \frac{\log((1+L_\pi)C_{\text{ISS}}^2)}{a \log(1/\rho)} + \frac{4 \log(a \log(1/\rho))}{a \log(1/\rho)},$$

implies $a\ell \geq 1 + \frac{\log((1+L_\pi)C_{\text{ISS}}^2)}{\log(1/\rho)}$, which in turn implies $(1+L_\pi)C_{\text{ISS}}^2 \ell \rho^{\ell-1} \leq \rho^\ell$ as required. For the corollary, we observe that the maximum value attained by $r^* \triangleq \max_{h \in \mathbb{N}} (1+L_\pi)C_{\text{ISS}}^2 h \rho^{h-1} / \rho^h = (1+L_\pi)C_{\text{ISS}}^2 h \rho^{a h-1}$ is upper bounded by $\frac{(1+L_\pi)C_{\text{ISS}}^2}{3a\rho \log(1/\rho)}$, completing the result. \square

Toward bounding $\|\mathbf{x}_T^{\tilde{\pi}} - \bar{\mathbf{x}}_T^{\tilde{\pi}}\|$, we define the number of full chunks traversed $K-1 = \lfloor (T-1)/\ell \rfloor$, and the remaining timesteps $h = T - (K-1)\ell - 1$. Further define the shorthands $\mathbf{U}_k = C_{\text{ISS}} \sum_{s=0}^{\ell-1} \rho^{\ell-1-s} \|\mathbf{u}_{t_k+s}\|$ for $k \in [K]$, and $\mathbf{U}_{K+1} = C_{\text{ISS}} \sum_{s=0}^{h-1} \rho^{h-1-s} \|\mathbf{u}_{t_K+s}\|$. Then, for ℓ satisfying Lemma A.5, we use Lemma A.5 to iteratively peel:

$$\begin{aligned} \|\mathbf{x}_T^{\tilde{\pi}} - \bar{\mathbf{x}}_T^{\tilde{\pi}}\| &\leq C_{\text{ISS}} \rho^h \|\mathbf{x}_{t_K}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_K}^{\tilde{\pi}}\| + C_{\text{ISS}} \sum_{s=0}^{h-1} \rho^{h-1-s} \|\mathbf{u}_{t_K+s}^{\tilde{\pi}} - \bar{\mathbf{u}}_{t_K+s}^{\tilde{\pi}} - \mathbf{u}_{t_K+s}\| \\ &\leq \bar{C} \rho^h \|\mathbf{x}_{t_K}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_K}^{\tilde{\pi}}\| + \mathbf{U}_{K+1} \\ &\leq \bar{C} \rho^{\ell+h} \|\mathbf{x}_{t_{K-1}}^{\tilde{\pi}} - \bar{\mathbf{x}}_{t_{K-1}}^{\tilde{\pi}}\| + \bar{C} \rho^\ell \mathbf{U}_K + \mathbf{U}_{K+1} \\ &\vdots \\ &\leq \bar{C} \rho^{T-1} \|\mathbf{x}_1 - \mathbf{x}'_1\| + \bar{C} \sum_{k=1}^{K+1} \rho^{(k-1)\ell} \mathbf{U}_k \\ &\leq \bar{C} \rho^{T-1} \|\mathbf{x}_1 - \mathbf{x}'_1\| + \bar{C} C_{\text{ISS}} \sum_{s=1}^{T-1} \rho^{T-1-s} \|\mathbf{u}_s\|. \end{aligned}$$

This establishes that $(\tilde{\pi}, f)$ is $(\bar{C} C_{\text{ISS}}, \rho^{1-a})$, given the chunk length satisfies $\ell > a^{-1} \log(\text{poly}(1 + L_\pi, C_{\text{ISS}}, \log(1/\rho), a^{-1}))$, and leveraging $\rho \geq 1/e$, we complete the result. \square

Having established that the chunked policy on the true dynamics $(\tilde{\pi}, f)$ is $(\tilde{C}, \hat{\rho}^{1-a})$ stable, we want to show this controls compounding errors when $\tilde{\pi}$ achieves low on-expert error to an expert policy π^* . This is a straightforward application of EISS. In particular, by treating the expert inputs as perturbations to a closed-loop system induced by $(\tilde{\pi}, f)$, we may relate $\mathbf{J}_{\text{TRAJ},1,T}$ to $\mathbf{J}_{\text{DEMO},1,T}$.

Proposition A.7 (Full ver. of Proposition 3.2). *Let Assumption 3.1 hold. Let $\tilde{\pi} = \text{chunked}(\hat{\pi}, \hat{f}, \ell) \in \Pi_{\text{chunk},\ell}$, and assume $(\hat{\pi}, \hat{f})$, $(\tilde{\pi}, f)$ are $(\tilde{C}, \tilde{\rho})$ -EISS. Then, the following bound holds:*

$$\mathbf{J}_{\text{TRAJ},p,T}^{\mathbf{x}}(\tilde{\pi}) \leq \left(\frac{\tilde{C}}{1-\tilde{\rho}} \right)^p \mathbf{J}_{\text{DEMO},p,T}(\tilde{\pi}; \mathbb{P}_{\pi^*}).$$

We have subsequently:

$$\mathbf{J}_{\text{TRAJ},p,T}(\hat{\pi}) \leq \text{poly}^p \left(L_\pi, \tilde{C}, \frac{1}{1-\tilde{\rho}} \right) \mathbf{J}_{\text{DEMO},p,T}(\hat{\pi}; \mathbb{P}_{\pi^*}).$$

Proof. Given $\mathbf{x}_1^{\pi^*} = \mathbf{x}_1^{\tilde{\pi}} \sim D$, we define $\mathbf{x}_t^{\pi^*}, \mathbf{u}_t^{\pi^*}$ and $\mathbf{x}_t^{\tilde{\pi}}, \mathbf{u}_t^{\tilde{\pi}}$ as the states and inputs given by the expert policy π^* and chunked policy $\tilde{\pi}$ in closed-loop. We may then view $(\mathbf{x}_t^{\pi^*}, \mathbf{u}_t^{\pi^*})$ as the resulting trajectory generated by appropriately defined “input perturbations” to the closed-loop chunked system $f^{\tilde{\pi}}: \mathbf{x}_{t+1}^{\pi^*} = f^{\tilde{\pi}}(\mathbf{x}_t^{\pi^*}, \Delta_{\mathbf{u}_t})$, $t \geq 1$, where we define

$$\Delta_{\mathbf{u}_t} \triangleq \mathbf{u}_t^{\pi^*} - \text{chunk}_s[\tilde{\pi}](\mathbf{x}_{t_k}^{\pi^*}),$$

and $t_k = \lfloor \frac{t-1}{\ell} \rfloor$ and $s = t - 1 \bmod \ell$. Therefore, applying the $(\tilde{C}, \tilde{\rho})$ -ISS of $(\tilde{\pi}, f)$, we have:

$$\begin{aligned} \|\mathbf{x}_t^{\tilde{\pi}} - \mathbf{x}_t^{\pi^*}\| &\leq \tilde{C} \sum_{s=1}^{t-1} \tilde{\rho}^{t-1-k} \|\Delta_{\mathbf{u}_t}\| \\ \iff \mathbf{J}_{\text{TRAJ},1,T}^{\mathbf{x}}(\tilde{\pi}) &\leq \frac{\tilde{C}}{1-\tilde{\rho}} \mathbf{J}_{\text{DEMO},1,T}(\tilde{\pi}; \mathbb{P}_{\pi^*}). \end{aligned}$$

The second line follows straightforwardly by summing both sides from $t = 1$ to T and applying an expectation. To extend this bound from $\mathbf{J}_{\text{TRAJ},1,T}^{\mathbf{x}}$ to general $\mathbf{J}_{\text{TRAJ},p,T}^{\mathbf{x}}$, we leverage [Lemma A.3](#). We define the vectors $\mathbf{u}, \mathbf{v} \in \mathbb{R}^{T-1}$:

$$\begin{aligned} \mathbf{u} &= [\|\mathbf{x}_2^{\tilde{\pi}} - \mathbf{x}_2^{\pi^*}\| \quad \cdots \quad \|\mathbf{x}_T^{\tilde{\pi}} - \mathbf{x}_T^{\pi^*}\|]^T \in \mathbb{R}^{T-1}, \\ \mathbf{v} &= [\|\Delta_{\mathbf{u}_1}\| \quad \cdots \quad \|\Delta_{\mathbf{u}_{T-1}}\|]^T \in \mathbb{R}^{T-1}. \end{aligned}$$

We observe that defining $\mathbf{A}(\tilde{\rho})$ as in [Lemma A.3](#), we have $\mathbf{u} = \tilde{C}\mathbf{A}\mathbf{v}$. Taking the p -norm on both sides and applying [Lemma A.3](#) yields: $\|\mathbf{u}\|_p \leq \frac{\tilde{C}}{1-\tilde{\rho}} \|\mathbf{v}\|_p$. Taking the p -th power and applying an expectation over $\mathbf{x}_1 \sim D$ on both sides yields the desired bound on $\mathbf{J}_{\text{TRAJ},p,T}^{\mathbf{x}}$ in terms of $\mathbf{J}_{\text{DEMO},p,T}$.

To extend this a bound on $\mathbf{J}_{\text{TRAJ},p,T}$, we apply a similar bound to [Lemma A.1](#). However, we require some alterations since $\tilde{\pi}$ is not a Markovian policy. We may add and subtract to yield:

$$\|\mathbf{u}_t^{\tilde{\pi}} - \mathbf{u}_t^{\pi^*}\| \leq \|\mathbf{u}_t^{\tilde{\pi}} - \text{chunk}_s[\tilde{\pi}](\mathbf{x}_{t_k}^{\pi^*})\| - \|\mathbf{u}_t^{\pi^*} - \text{chunk}_s[\tilde{\pi}](\mathbf{x}_{t_k}^{\pi^*})\|.$$

Summing up the second term over t yields $\mathbf{J}_{\text{DEMO},T}$. To analyze the first term, we recall that $\mathbf{u}_t^{\tilde{\pi}}$ and $\text{chunk}_s[\tilde{\pi}](\mathbf{x}_{t_k}^{\pi^*})$ result from conditioning on the state every t_k timesteps, then playing the next ℓ actions generated by the simulated closed-loop system $(\hat{\pi}, \hat{f})$. Since by assumption $\hat{f}^{\hat{\pi}}$ is (\tilde{C}, ρ) -ISS, this means that for each t_k and $s = 0, \dots, \ell - 1$,

$$\|\mathbf{x}_{t_k+s} - \tilde{\mathbf{x}}_{t_k+s}\| \leq \tilde{C}\tilde{\rho}^s \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \mathbf{x}_{t_k}^{\pi^*}\|,$$

where $\mathbf{x}_{t_k+s} = (\hat{f}^{\hat{\pi}})^s(\mathbf{x}_{t_k}^{\tilde{\pi}})$, $\tilde{\mathbf{x}}_{t_k+s} = (\hat{f}^{\hat{\pi}})^s(\mathbf{x}_{t_k}^{\pi^*})$. Furthermore, since $\mathbf{u}_{t_k+s}^{\tilde{\pi}} \triangleq \hat{\pi}(\mathbf{x}_{t_k+s})$ and similarly $\text{chunk}_s[\tilde{\pi}](\mathbf{x}_{t_k}^{\pi^*}) \triangleq \hat{\pi}(\tilde{\mathbf{x}}_{t_k+s})$, applying L_{π} -Lipschitzness of $\hat{\pi}$ yields:

$$\begin{aligned} \sum_{k=1}^{T-1/\ell} \sum_{s=0}^{\ell-1} \|\mathbf{u}_{t_k+s}^{\tilde{\pi}} - \text{chunk}_s[\tilde{\pi}](\mathbf{x}_{t_k}^{\pi^*})\|^p &\leq \sum_{k=1}^{T-1/\ell} \sum_{s=0}^{\ell-1} L_{\pi}^p (\tilde{C}\tilde{\rho}^s)^p \|\mathbf{x}_{t_k}^{\tilde{\pi}} - \mathbf{x}_{t_k}^{\pi^*}\|^p \\ &\leq \left(\frac{L_{\pi}\tilde{C}}{1-\tilde{\rho}} \right)^p \sum_{t=1}^T \|\mathbf{x}_t^{\tilde{\pi}} - \mathbf{x}_t^{\pi^*}\|^p. \end{aligned}$$

Putting these pieces together, we get:

$$\begin{aligned}
\mathbf{J}_{\text{TRAJ},p,t}(\tilde{\pi}) &\leq \mathbf{J}_{\text{TRAJ},p,t}^{\mathbf{x}}(\tilde{\pi}) + \sum_{t=1}^T \min\{1, \|\mathbf{u}_t^{\tilde{\pi}} - \mathbf{u}_t^{\pi^*}\|^p\} \\
&\leq \mathbf{J}_{\text{TRAJ},p,t}^{\mathbf{x}}(\tilde{\pi}) + 2^p \mathbf{J}_{\text{DEMO},p,T}(\tilde{\pi}; \mathbb{P}_{\pi^*}) + \left(\frac{2L_{\pi}\tilde{C}}{1-\tilde{\rho}}\right)^p \sum_{t=1}^T \min\{1, \|\mathbf{x}_t^{\tilde{\pi}} - \mathbf{x}_t^{\pi^*}\|^p\} \\
&\leq \left(1 + \left(\frac{2L_{\pi}\tilde{C}}{1-\tilde{\rho}}\right)^p\right) \mathbf{J}_{\text{TRAJ},p,t}^{\mathbf{x}}(\tilde{\pi}) + 2^p \mathbf{J}_{\text{DEMO},p,T}(\tilde{\pi}; \mathbb{P}_{\pi^*}).
\end{aligned}$$

Plugging in the upper bound on $\mathbf{J}_{\text{TRAJ},p,T}^{\mathbf{x}}(\tilde{\pi})$ completes the result. \square

In particular, specifying this result to [Proposition 3.2](#) follows straightforwardly by setting $p = 2$. Therefore, combining [Proposition A.4](#), which says chunking policies induces ISS, with [Proposition A.7](#), which says EISS chunking policies induce low compounding error, yields the final guarantee.

Theorem 3 (Full ver. of [Theorem 1](#)). *Let [Assumption 3.1](#) and [Assumption 3.2](#) hold. Given $a \in (0, 1)$, for sufficiently long chunk-length: $\ell > a^{-1} \log(1/\rho)^{-1} \cdot \log(\text{poly}(L_{\pi}, \hat{C}, 1/a))$, let $\tilde{\pi} = \text{chunked}(\hat{\pi}, g, \ell) \in \Pi_{\text{chunk},\ell}$, such that $(\tilde{\pi}, f)$ is (\tilde{C}, ρ^{1-a}) , with $\tilde{C} = a^{-1} \log(1/\rho)^{-1} \cdot \text{poly}(L_{\pi}, C_{\text{ISS}})$. The following bound holds on the trajectory error induced by $\tilde{\pi}$:*

$$\mathbf{J}_{\text{TRAJ},p,T}(\tilde{\pi}) \lesssim \left(1 + \frac{L_{\pi}\tilde{C}}{1-\rho^{1-a}}\right)^p \mathbf{J}_{\text{DEMO},p,T}(\tilde{\pi}; \mathbb{P}_{\pi^*}).$$

B Proofs and Additional Details for [Section 4](#)

Preliminaries. We first recall the definition of the linearizations around expert trajectories from [Definition 4.2](#).

$$\begin{aligned}
\mathbf{A}_t &= \nabla_{\mathbf{x}} f(\mathbf{x}_t^{\pi^*}, \mathbf{u}_t^{\pi^*}), \quad \mathbf{B}_t = \nabla_{\mathbf{u}} f(\mathbf{x}_t^{\pi^*}, \mathbf{u}_t^{\pi^*}), \quad \mathbf{K}_t^{\pi^*} = \nabla_{\mathbf{x}} \pi^*(\mathbf{x}_t^{\pi^*}), \\
\mathbf{A}_{s:t}^{\text{cl}} &= (\mathbf{A}_{t-1} + \mathbf{B}_{t-1} \mathbf{K}_{t-1}^{\pi^*})(\mathbf{A}_{t-2} + \mathbf{B}_{t-2} \mathbf{K}_{t-2}^{\pi^*}) \cdots (\mathbf{A}_s + \mathbf{B}_s \mathbf{K}_s^{\pi^*}).
\end{aligned} \tag{B.1}$$

We also recall the definition of the controllability Gramian: $\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}) \triangleq \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} \mathbf{B}_s \mathbf{B}_s^{\top} \mathbf{A}_{s+1:t}^{\text{cl} \top}$. For a noising distribution that is zero-mean with covariance $\Sigma_{\mathbf{z}}$, $\mathbf{z}_t \stackrel{\text{i.i.d.}}{\sim} \mathcal{D}(\mathbf{0}, \Sigma_{\mathbf{z}})$, we further define the *noise controllability Gramian*:

$$\mathbf{W}_{1:t}^{\mathbf{z}}(\mathbf{x}_1^{\pi^*}) \triangleq \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} \mathbf{B}_s \Sigma_{\mathbf{z}} \mathbf{B}_s^{\top} \mathbf{A}_{s+1:t}^{\text{cl} \top}.$$

Note that for \mathbf{z}_t sampled from the Euclidean unit ball, we have $\Sigma_{\mathbf{z}} = \frac{1}{(d_u+2)} \mathbf{I}_{d_u} \succeq \frac{1}{3d_u} \mathbf{I}_{d_u}$, and thus:

$$\mathbf{W}_{1:t}^{\mathbf{z}}(\mathbf{x}_1^{\pi^*}) \succeq \frac{1}{3d_u} \mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}).$$

The ensuing results are written for any noising distribution $\mathcal{D}(\mathbf{0}, \Sigma_{\mathbf{z}})$ that are 1-bounded, mean-zero, with covariance $\Sigma_{\mathbf{z}} \succ \mathbf{0}$, unless otherwise stated.

We now establish that the linear (time-varying) system induced by linearizations along expert trajectories inherits (C_{ISS}, ρ) -EISS. We note that though the original dynamics and expert policy are time-invariant, the linearized system is in general not.

Lemma B.1. *Let Assumption 4.1 hold. Given a nominal trajectory generated as*

$$\mathbf{x}_{t+1}^{\pi^*} = f(\mathbf{x}_t^{\pi^*}, \mathbf{u}_t^{\pi^*}), \mathbf{u}_t^{\pi^*} = \pi^*(\mathbf{x}_t^{\pi^*}), t \geq 1, \mathbf{x}_1^{\pi^*} \sim \mathbb{P}_{\mathbf{x}_1^{\pi^*}},$$

and recall the linearizations in Eq. (B.1). Then, the following bounds hold:

$$\|\mathbf{A}_{1:t}^{\text{cl}}\|_{\text{op}} \leq C_{\text{ISS}}\rho^{t-1}, \|\mathbf{A}_{s:t}^{\text{cl}}\|_{\text{op}} \leq C_{\text{ISS}}\rho^{t-s}, \|\mathbf{A}_{s+1:t}^{\text{cl}}\mathbf{B}_s\|_{\text{op}} \leq C_{\text{ISS}}\rho^{t-1-s}, \text{ for all } 1 \leq s \leq t.$$

An equivalent way to view Lemma B.1 is: for an input perturbation sequence $\{\Delta_{\mathbf{u}_t}\}_{t \geq 1}$, the incremental trajectory $\{\Delta_{\mathbf{x}_t}\}_{t \geq 1}$, $\Delta_{\mathbf{x}_t} \triangleq \hat{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}$ induced by linearizations around an expert trajectory $\{\mathbf{x}_t^{\pi^*}\}_{t \geq 1}$ is (C_{ISS}, ρ) -EISS:

$$\Delta_{\mathbf{x}_{t+1}} = (\mathbf{A}_t + \mathbf{B}_t \mathbf{K}_t^{\pi^*}) \Delta_{\mathbf{x}_t} + \mathbf{B}_t \Delta_{\mathbf{u}_t} = \mathbf{A}_{1:t+1}^{\text{cl}} \Delta_{\mathbf{x}_1} + \sum_{s=1}^t \mathbf{A}_{s+1:t+1}^{\text{cl}} \mathbf{B}_s \Delta_{\mathbf{u}_s}.$$

Proof of Lemma B.1. Given the nominal trajectory $\{\mathbf{x}_t^{\pi^*}\}_{t \geq 1}$ generated by $\mathbf{x}_{t+1}^{\pi^*} = f(\mathbf{x}_t^{\pi^*}, \mathbf{u}_t^{\pi^*})$ and the corresponding linearizations $\mathbf{A}_t, \mathbf{B}_t, \mathbf{K}_t^{\pi^*}$ evaluated along the trajectory, consider the trajectory $\{\tilde{\mathbf{x}}_t\}_{t \geq 1}$ generated as $\tilde{\mathbf{x}}_{t+1} = f(\tilde{\mathbf{x}}_t, \pi^*(\tilde{\mathbf{x}}_t) + \mathbf{u}_t, t)$. Expanding the Jacobian linearizations, we have

$$\begin{aligned} \tilde{\mathbf{x}}_{t+1} - \mathbf{x}_{t+1}^{\pi^*} &= f(\tilde{\mathbf{x}}_t, \pi^*(\tilde{\mathbf{x}}_t) + \mathbf{u}_t, t) - f(\mathbf{x}_t^{\pi^*}, \pi^*(\mathbf{x}_t^{\pi^*}), t) \\ &= \mathbf{A}_t(\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}) + \mathbf{B}_t(\pi^*(\tilde{\mathbf{x}}_t) + \mathbf{u}_t - \pi^*(\mathbf{x}_t^{\pi^*})) + O\left(\underbrace{\left\| \begin{bmatrix} \tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*} \\ \pi^*(\tilde{\mathbf{x}}_t) - \pi^*(\mathbf{x}_t^{\pi^*}) + \mathbf{u}_t \end{bmatrix} \right\|^2}_{\triangleq \mathbf{r}_t^{\mathbf{x}}}\right) \\ &= (\mathbf{A}_t + \mathbf{B}_t \mathbf{K}_t^{\pi^*})(\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}) + \mathbf{B}_t(\mathbf{u}_t + \underbrace{O(\|\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}\|^2)}_{\triangleq \mathbf{r}_t^{\mathbf{u}}}) + \mathbf{r}_t^{\mathbf{x}} \\ &= \mathbf{A}_{1:t+1}^{\text{cl}}(\tilde{\mathbf{x}}_1 - \mathbf{x}_1^{\pi^*}) + \sum_{s=1}^t \mathbf{A}_{s+1:t+1}^{\text{cl}}(\mathbf{B}_s(\mathbf{u}_s + \mathbf{r}_s^{\mathbf{u}}) + \mathbf{r}_s^{\mathbf{x}}), \end{aligned} \tag{B.2}$$

We perform a simple sensitivity analysis to isolate $\mathbf{A}_{1:t}^{\text{cl}}$. Defining the displacements $\delta_t^{\mathbf{x}} = \tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}$, and setting $\mathbf{u}_t = \mathbf{0}$, $t \geq 1$, we see that $\frac{\partial}{\partial \delta_1^{\mathbf{x}}} \delta_t^{\mathbf{x}} = \mathbf{A}_{1:t}^{\text{cl}}$, since we observe $\delta_t^{\mathbf{x}} = \mathbf{A}_{1:t}^{\text{cl}} \delta_1^{\mathbf{x}} + \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}}(\mathbf{B}_s \mathbf{r}_s^{\mathbf{u}} + \mathbf{r}_s^{\mathbf{x}})$ is linear in $\delta_1^{\mathbf{x}}$ and the residuals $\mathbf{r}_s^{\mathbf{u}}, \mathbf{r}_s^{\mathbf{x}}$ are higher-order by definition. On the other hand, by the (C_{ISS}, ρ) -EISS of (π^*, f) , we know that $\|\delta_t^{\mathbf{x}}\| \leq C_{\text{ISS}}\rho^{t-1}\|\delta_1^{\mathbf{x}}\|$. By definition of the operator norm, we have $\|\mathbf{A}_{1:t}^{\text{cl}}\|_{\text{op}} = \sup_{\mathbf{v}} \|\mathbf{A}_{1:t}^{\text{cl}} \mathbf{v}\|/\|\mathbf{v}\|$, and thus by a limiting argument $\delta_1^{\mathbf{x}} \rightarrow \mathbf{0}$, we see

$$\|\mathbf{A}_{1:t}^{\text{cl}}\|_{\text{op}} \leq \lim_{\delta_1^{\mathbf{x}} \rightarrow \mathbf{0}} \|\delta_t^{\mathbf{x}}\|/\|\delta_1^{\mathbf{x}}\| \leq C_{\text{ISS}}\rho^{t-1}.$$

To establish a similar bound on $\mathbf{A}_{s:t}^{\text{cl}}$, we observe that $\mathbf{x}_{t+1}^{\pi^*} = f(\mathbf{x}_t^{\pi^*}, \pi^*(\mathbf{x}_t^{\pi^*}))$ is by definition a *time-invariant* closed-loop system, we may apply (C_{ISS}, ρ) -EISS starting from $\delta_s^{\mathbf{x}}$ as the initial displacement such that $\|\delta_t^{\mathbf{x}}\| \leq C_{\text{ISS}}\rho^{t-s}\|\delta_s^{\mathbf{x}}\|$. Applying the same argument yields:

$$\|\mathbf{A}_{s:t}^{\text{cl}}\|_{\text{op}} \leq \lim_{\delta_s^{\mathbf{x}} \rightarrow \mathbf{0}} \|\delta_t^{\mathbf{x}}\|/\|\delta_s^{\mathbf{x}}\| \leq C_{\text{ISS}}\rho^{t-s}.$$

Now, instead setting $\tilde{\mathbf{x}}_1 - \mathbf{x}_1^{\pi^*} = \mathbf{0}$ and an impulse input $\{\mathbf{0}, \dots, \mathbf{0}, \mathbf{u}_k, \mathbf{0}, \dots\}$ for some k , we have $\delta_t^{\mathbf{x}} = \mathbf{A}_{k+1:t}^{\text{cl}} \mathbf{B}_k \mathbf{u}_k + \sum_{s=k}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} (\mathbf{B}_s \mathbf{r}_s^{\mathbf{u}} + \mathbf{r}_s^{\mathbf{x}})$. By the same appeal to EISS of (π^*, f) and limiting argument $\mathbf{u}_k \rightarrow \mathbf{0}$, we have: $\|\mathbf{A}_{k+1:t}^{\text{cl}} \mathbf{B}_k\|_{\text{op}} \leq C_{\text{ISS}} \rho^{t-1-k}$. Notably, this holds for any k and $t \geq k$, completing the proof. \square

Given an expert-induced trajectory $\mathbf{x}_{t+1} = f^{\pi^*}(\mathbf{x}_t)$, $t \in [T-1]$, consider *noise-injected* trajectories $(\tilde{\mathbf{x}}_t, \tilde{\mathbf{u}}_t)_{t \geq 1} \sim \mathbb{P}_{\pi^*, \sigma_{\mathbf{u}}}$ as in Definition 4.1. Our next result demonstrates that the noise-injected trajectories are well-described by the expert linearizations, up to a higher-order term quadratic in the noise-scale $\sigma_{\mathbf{u}}$.

Proposition B.2. *Let Assumption 4.1 hold. Consider noise-injected expert trajectories $\{\tilde{\mathbf{x}}_t, \pi^*(\tilde{\mathbf{x}}_t)\}_{t \geq 1} \sim \mathbb{P}_{\pi^*, \sigma_{\mathbf{u}}}$ for a given initial condition $\tilde{\mathbf{x}}_1 \sim D$: $\tilde{\mathbf{x}}_{t+1} = f(\tilde{\mathbf{x}}_t, \pi^*(\tilde{\mathbf{x}}_t) + \sigma_{\mathbf{u}} \mathbf{z}_t)$, $\mathbf{z}_t \stackrel{\text{i.i.d.}}{\sim} \mathcal{D}(\mathbf{0}, \Sigma_{\mathbf{z}})$. Consider the linearizations along an expert trajectory given in (B.1), setting $\mathbf{x}_1^{\pi^*} = \tilde{\mathbf{x}}_1$. Define the linear and residual components of the noised state $\tilde{\mathbf{x}}_t$:*

$$\tilde{\mathbf{x}}_t^{\text{lin}} \triangleq \mathbf{x}_t^{\pi^*} + \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} \mathbf{B}_s \mathbf{u}_s, \quad \tilde{\mathbf{x}}_t^{\text{res}} \triangleq \tilde{\mathbf{x}}_t - \tilde{\mathbf{x}}_t^{\text{lin}}, \quad t \geq 1. \quad (\text{B.3})$$

Then, as long as $\sigma_{\mathbf{u}} \leq \frac{1}{2} c_{\text{stab}} \frac{\sqrt{1+4C_K^2}}{C_{\pi}}$, and defining $C_r \triangleq C_{\pi} + 4C_{\text{reg}}(1+4C_K^2)$, we have $\|\tilde{\mathbf{x}}_t^{\text{res}}\| \leq C_{\text{stab}}^3 C_r \sigma_{\mathbf{u}}^2$, $t \geq 1$ almost surely over $\tilde{\mathbf{x}}_1 \sim D$ and $\{\mathbf{z}_s\} \stackrel{\text{i.i.d.}}{\sim} \mathcal{D}(\mathbf{0}, \Sigma_{\mathbf{z}})$.

Proof of Proposition B.2. Given the nominal trajectory $\{\mathbf{x}_t^{\pi^*}\}_{t \geq 1}$ generated by $\mathbf{x}_{t+1}^{\pi^*} = f(\mathbf{x}_t^{\pi^*}, \mathbf{u}_t^{\pi^*})$ and the corresponding linearizations $\mathbf{A}_t, \mathbf{B}_t, \mathbf{K}^{\pi^*}$ (B.1) evaluated along the trajectory, consider the trajectory $\{\tilde{\mathbf{x}}_t\}_{t \geq 1}$ generated as $\tilde{\mathbf{x}}_{t+1} = f(\tilde{\mathbf{x}}_t, \pi^*(\tilde{\mathbf{x}}_t) + \mathbf{u}_t, t)$, with $\tilde{\mathbf{x}}_1 = \mathbf{x}_1^{\pi^*}$. Then, following (B.2), we may write:

$$\begin{aligned} \tilde{\mathbf{x}}_{t+1} - \mathbf{x}_{t+1}^{\pi^*} &= f(\tilde{\mathbf{x}}_t, \pi^*(\tilde{\mathbf{x}}_t) + \mathbf{u}_t, t) - f(\mathbf{x}_t^{\pi^*}, \pi^*(\mathbf{x}_t^{\pi^*}), t) \\ &= \mathbf{A}_{1:t+1}^{\text{cl}} (\tilde{\mathbf{x}}_1 - \mathbf{x}_1^{\pi^*}) + \sum_{s=1}^t \mathbf{A}_{s+1:t+1}^{\text{cl}} (\mathbf{B}_s (\mathbf{u}_s + \mathbf{r}_s^{\mathbf{u}}) + \mathbf{r}_s^{\mathbf{x}}) \\ &= \sum_{s=1}^t \mathbf{A}_{s+1:t+1}^{\text{cl}} (\mathbf{B}_s (\mathbf{u}_s + \mathbf{r}_s^{\mathbf{u}}) + \mathbf{r}_s^{\mathbf{x}}). \end{aligned}$$

where we recall $\mathbf{r}_t^{\mathbf{x}}$ and $\mathbf{r}_t^{\mathbf{u}}$ are the second-order remainder terms of the dynamics and policy outputs, respectively. By Assumption 4.1, these are bounded by:

$$\begin{aligned} \|\mathbf{r}_t^{\mathbf{u}}\| &\leq C_{\pi} \|\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}\|^2 \\ \|\mathbf{r}_t^{\mathbf{x}}\| &\leq C_{\text{reg}} (\|\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}\|^2 + \|\pi^*(\tilde{\mathbf{x}}_t) - \pi^*(\mathbf{x}_t^{\pi^*}) + \mathbf{u}_t\|^2) \\ &\leq C_{\text{reg}} (\|\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}\|^2 + 2\|\pi^*(\tilde{\mathbf{x}}_t) - \pi^*(\mathbf{x}_t^{\pi^*})\|^2 + 2\|\mathbf{u}_t\|^2) \\ &\leq C_{\text{reg}} ((1+4C_K^2)\|\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}\|^2 + 4\|\mathbf{r}_t^{\mathbf{u}}\|^2 + 2\|\mathbf{u}_t\|^2) \\ &\leq C_{\text{reg}} ((1+4C_K^2)\|\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}\|^2 + 4C_{\pi}^2 \|\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}\|^4 + 2\|\mathbf{u}_t\|^2). \end{aligned}$$

Defining $\varepsilon_{\tilde{\mathbf{x}}_t} = \|\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}\|$, and $\mathbf{u}_t \stackrel{\text{i.i.d.}}{\sim} \mathcal{D}(\mathbf{0}, \Sigma_{\mathbf{u}}; \sigma_{\mathbf{u}})$ are iid zero-mean, $\Sigma_{\mathbf{u}}$ covariance, $\sigma_{\mathbf{u}}$ -bounded random vectors, we want to bound the mean and covariance of $\tilde{\mathbf{x}}_t$. We note the presence of the quartic term

$4C_\pi^2 \varepsilon_{\mathbf{x}_t}^4$ in our remainder term; we first impose $\varepsilon_{\mathbf{x}_t}^2 \leq \frac{1+4C_K^2}{4C_\pi^2}$ to absorb it into the quadratic term, then show this constraint is obviated for sufficiently small $\|\mathbf{u}_s\|$.

Since (π^*, f) is (C_{ISS}, ρ) -EISS, we have $\varepsilon_{\mathbf{x}_t} \leq C_{\text{ISS}} \sum_{s=1}^{t-1} \rho^{t-s-1} \|\mathbf{u}_s\| \leq \frac{C_{\text{ISS}}}{1-\rho} \max_{s \leq t-1} \|\mathbf{u}_s\|$. Therefore, we have:

$$\begin{aligned} \|\mathbf{r}_t^{\mathbf{u}}\| &\leq C_\pi \varepsilon_{\mathbf{x}_t}^2 \leq C_\pi C_{\text{stab}}^2 \max_{s \leq t-1} \|\mathbf{u}_s\|^2 \\ &\leq C_\pi C_{\text{stab}}^2 \sigma_{\mathbf{u}}^2 \\ \|\mathbf{r}_t^{\mathbf{x}}\| &\leq C_{\text{reg}} \left((1 + 4C_K^2) \varepsilon_{\mathbf{x}_t}^2 + 4C_\pi^2 \varepsilon_{\mathbf{x}_t}^4 + 2\|\mathbf{u}_t\|^2 \right) \\ &\leq C_{\text{reg}} \left(2(1 + 4C_K^2) C_{\text{stab}}^2 \max_{s \leq t-1} \|\mathbf{u}_s\|^2 + 2\|\mathbf{u}_t\|^2 \right) \\ &\leq 4C_{\text{reg}} (1 + 4C_K^2) C_{\text{stab}}^2 \sigma_{\mathbf{u}}^2. \end{aligned}$$

These hold as long as $\sigma_{\mathbf{u}}$ is small enough such that $\varepsilon_{\mathbf{x}_t}^2 \leq C_{\text{stab}}^2 \sigma_{\mathbf{u}}^2 \leq \frac{1+4C_K^2}{4C_\pi^2}$, which holds for $\sigma_{\mathbf{u}} \leq \frac{1}{2} C_{\text{stab}} \frac{\sqrt{1+4C_K^2}}{C_\pi}$. With these perturbation bounds in hand, we now move onto bounding the linear and residual components of $\tilde{\mathbf{x}}_t$. We have immediately:

$$\begin{aligned} \tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*} &= \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} (\mathbf{B}_s (\mathbf{u}_s + \mathbf{r}_s^{\mathbf{u}}) + \mathbf{r}_s^{\mathbf{x}}) \\ \Rightarrow \|\tilde{\mathbf{x}}_t^{\text{res}}\| = \|\tilde{\mathbf{x}}_t^{\text{lin}} - \mathbf{x}_t^{\pi^*}\| &= \left\| \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} (\mathbf{B}_s \mathbf{r}_s^{\mathbf{u}} + \mathbf{r}_s^{\mathbf{x}}) \right\| \\ &\leq \sum_{s=1}^{t-1} \|\mathbf{A}_{s+1:t}^{\text{cl}} \mathbf{B}_s\|_{\text{op}} \|\mathbf{r}_s^{\mathbf{u}}\| + \|\mathbf{A}_{s+1:t}^{\text{cl}}\|_{\text{op}} \|\mathbf{r}_s^{\mathbf{x}}\| \\ &\leq C_{\text{stab}}^3 \underbrace{(C_\pi + 4C_{\text{reg}}(1 + 4C_K^2))}_{\triangleq C_r} \sigma_{\mathbf{u}}^2. \end{aligned}$$

This completes the proof. \square

We now proceed with the *one-step controllable* setting, where $\mathbf{W}_{1:t}^{\mathbf{u}} \succ \underline{\lambda}_{\mathbf{W}} \mathbf{I}$ for all $t \geq 2$, leading up to [Suboptimal Proposition 4.2](#), where we also fit $\hat{\pi}$ purely on noise-injected trajectories, in order to grasp the core ideas and the remaining key deficiencies.

B.1 One-step Controllable Case: Persistency of Excitation

We consider settings where the controllability Gramians induced by linearizations around an expert trajectory are always full-rank.

Assumption B.1 (Linearized one-step controllability). Let $\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}) \succeq \underline{\lambda}_{\mathbf{W}} \mathbf{I}_{d_u}$, $t \geq 2$ w.p. 1 over $\mathbf{x}_1^{\pi^*} \sim D$ for some $\underline{\lambda}_{\mathbf{W}} > 0$. Consider the noise-controllability Gramians $\mathbf{W}_{1:t}^z$ as defined in [Definition 4.2](#). Accordingly, there exists $\underline{\lambda}_z > 0$ such that w.p. 1 over $\mathbf{x}_1^{\pi^*} \sim D$, $\mathbf{W}_{1:t}^z(\mathbf{x}_1^{\pi^*}) \succeq \underline{\lambda}_z \mathbf{I}_{d_x}$, $t \geq 2$.

[Proposition B.2](#) in conjunction with [Assumption B.1](#) implies the noise-injected expert states $\tilde{\mathbf{x}}_t$ form a *full-rank* covariance around $\mathbf{x}_t^{\pi^*}$ for each timestep $t = 2, \dots, T$. This corresponds with the well-known notion of *persistency of excitation* from the control literature [[Annaswamy, 2023](#)]. As a consequence of [Proposition B.2](#), we have the following excitation bound.

Corollary B.1. Let [Assumption 4.1](#) hold and C_r be as defined in [Proposition B.2](#). Recall the noise-controllability Gramian $\mathbf{W}_{1:t}^z$ as in [Assumption B.1](#). As long as:

$$\sigma_u \lesssim \min \left\{ \lambda_{\min}^+ (\mathbf{W}_{1:t}^z (\mathbf{x}_1^{\pi^*})) c_{\text{stab}}^4 C_r^{-1}, c_{\text{stab}} \frac{\sqrt{1 + 4C_K^2}}{C_\pi} \right\},$$

the following holds almost surely over $\tilde{\mathbf{x}}_1 = \mathbf{x}_1^{\pi^*} \sim \mathbb{P}_{\mathbf{x}_1^{\pi^*}}$ and $\{\mathbf{z}_s\} \stackrel{\text{i.i.d.}}{\sim} \mathcal{D}(\mathbf{0}, \Sigma_z)$:

$$\mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} \left[(\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}) (\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})^\top \right] \geq \frac{\sigma_u^2}{2} \mathbf{W}_{1:t}^z (\mathbf{x}_1^{\pi^*}). \quad (\text{B.4})$$

Proof of Corollary B.1. Denoting $\mathbf{C} = \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} \mathbf{B}_s \mathbf{u}_s$ and $\mathbf{E} = \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} (\mathbf{B}_s \mathbf{r}_s^u + \mathbf{r}_s^x)$, we bound the second moment of $\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}$:

$$\begin{aligned} \mathbb{E}_{\mathbf{u}} \left[(\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}) (\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})^\top \right] &= \mathbb{E}_{\mathbf{u}} \left[(\mathbf{C} + \mathbf{E})(\mathbf{C} + \mathbf{E})^\top \right] \\ &= \mathbb{E}_{\mathbf{u}} \left[\mathbf{C}\mathbf{C}^\top \right] + \mathbb{E}_{\mathbf{u}} \left[\mathbf{E}\mathbf{E}^\top \right] + \mathbb{E}_{\mathbf{u}} \left[\mathbf{E}\mathbf{C}^\top + \mathbf{C}\mathbf{E}^\top \right] \\ &\geq \mathbb{E}_{\mathbf{u}} \left[\mathbf{C}\mathbf{C}^\top \right] + \mathbb{E}_{\mathbf{u}} \left[\mathbf{E}\mathbf{C}^\top + \mathbf{C}\mathbf{E}^\top \right] \end{aligned}$$

By Weyl's inequality [[Horn and Johnson, 2012](#)], we have for each $k = 1, \dots, \text{rank}(\mathbb{E}_{\mathbf{u}} [\mathbf{C}\mathbf{C}^\top])$:

$$\begin{aligned} \left| \lambda_k(\mathbb{E}_{\mathbf{u}}[(\mathbf{C} + \mathbf{E})(\mathbf{C} + \mathbf{E})^\top]) - \lambda_k(\mathbb{E}_{\mathbf{u}}[\mathbf{C}\mathbf{C}^\top]) \right| &\leq 2\|\mathbf{E}\mathbf{C}^\top\|_{\text{op}} \\ &\leq 2 \left(\frac{C_{\text{ISS}}}{1 - \rho} \sigma_u \right) (C_{\text{stab}}^3 (C_\pi + 4C_{\text{reg}}(1 + 4C_K^2)) \sigma_u^2) \\ &= 2C_{\text{stab}}^4 C_r \sigma_u^3. \end{aligned}$$

Rearranging the above yields, for each $k = 1, \dots, \text{rank}(\mathbb{E}_{\mathbf{u}} [\mathbf{C}\mathbf{C}^\top])$:

$$\lambda_k(\mathbb{E}_{\mathbf{u}}[(\mathbf{C} + \mathbf{E})(\mathbf{C} + \mathbf{E})^\top]) \geq \lambda_k(\mathbb{E}_{\mathbf{u}}[\mathbf{C}\mathbf{C}^\top]) - 2C_{\text{stab}}^4 (C_\pi + 4C_{\text{reg}}(1 + 4C_K^2)) \sigma_u^3.$$

Therefore, for sufficiently small σ_u such that:

$$\sigma_u \leq \frac{1}{4} \frac{\lambda_{\min}^+(\mathbb{E}_{\mathbf{u}}[\mathbf{C}\mathbf{C}^\top])}{\sigma_u^2} c_{\text{stab}}^4 C_r^{-1},$$

where $\lambda_{\min}^+(\cdot)$ denotes the smallest positive eigenvalue, we have $\lambda_k(\mathbb{E}_{\mathbf{u}}[(\mathbf{C} + \mathbf{E})(\mathbf{C} + \mathbf{E})^\top]) \geq \frac{1}{2} \lambda_k(\mathbb{E}_{\mathbf{u}}[\mathbf{C}\mathbf{C}^\top])$, $k = 1, \dots, \text{rank}(\mathbb{E}_{\mathbf{u}} [\mathbf{C}\mathbf{C}^\top])$ such that

$$\begin{aligned} \mathbb{E}_{\mathbf{u}} \left[(\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}) (\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})^\top \right] &\geq \frac{1}{2} \mathbb{E}_{\mathbf{u}} [\mathbf{C}\mathbf{C}^\top] \\ &= \frac{1}{2} \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} \mathbf{B}_s \Sigma_{\mathbf{u}} \mathbf{B}_s^\top \mathbf{A}_{s+1:t}^{\text{cl}^\top}. \end{aligned}$$

□

Proposition B.2 demonstrates that noise injection yields full-rank exploration around the expert trajectory that is essentially described by the controllability Gramian induced by linearizations around the expert trajectory. In this case, we show that a policy $\hat{\pi}$ attaining low on-expert error does not suffer exponential compounding error. The first ingredient is an adapted result from [Pfrommer et al. \[2022\]](#) that certifies low trajectory error as long as policies are persistently close in a tube around the expert trajectory.

Proposition B.3 (TaSIL [[Pfrommer et al., 2022](#)]). Assume the closed-loop system induced by (π^*, f) is (C_{ISS}, ρ) -EISS. For any (deterministic) policy $\hat{\pi}$ and initial state \mathbf{x}_1 , let $\mathbf{x}_1^{\hat{\pi}} = \mathbf{x}_1^{\pi^*} = \mathbf{x}_1$, and consider the closed-loop trajectories generated by $\hat{\pi}$ and π^* :

$$\mathbf{x}_{t+1}^{\hat{\pi}} = f(\mathbf{x}_t^{\hat{\pi}}, \hat{\pi}(\mathbf{x}_t^{\hat{\pi}})), \quad \mathbf{x}_{t+1}^{\pi^*} = f(\mathbf{x}_t^{\pi^*}, \pi^*(\mathbf{x}_t^{\pi^*})), \quad t \geq 1. \quad (\text{B.5})$$

Then for any given $\varepsilon > 0$, $T \in \mathbb{N}$, as long as:

$$\max_{1 \leq t \leq T-1} \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \varepsilon \mathbf{w})\| \leq c_{\text{stab}} \varepsilon,$$

we are guaranteed $\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\| \leq \varepsilon$.

An elementary proof to [Proposition B.3](#) can be found in e.g., [Simchowitz et al. \[2025, Lemma I.4\]](#). Our next ingredient demonstrates that if noise injection induces full-rank state covariances, closeness in a tube with radius proportional to the noise variance is certified, up to higher-order perturbations from smoothness.

Lemma B.4. Let [Assumption 4.1](#) hold, and let [Assumption B.1](#) hold with $\underline{\lambda}_z > 0$. Let $\{\mathbf{x}_t^{\pi^*}\}_{t=1}^T, \{\tilde{\mathbf{x}}_t\}_{t=1}^T$ be expert and noise-injected states initialized from a given $\mathbf{x}_1^{\pi^*} = \tilde{\mathbf{x}}_1$. Let $\hat{\pi}$ be any C_π -smooth (deterministic) policy. For sufficiently small noise-scale $\sigma_{\mathbf{u}} \lesssim \min\{c_{\text{stab}}^3 C_{\mathbf{r}}^{-1} \sqrt{\underline{\lambda}_z}, c_{\text{stab}} \sqrt{1 + 4C_{\mathbf{K}}^2 C_\pi^{-1}}\}$, the following holds for each $t = 1, \dots, T-1$:

$$\sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \varepsilon \mathbf{w})\|^2 \leq 16 \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2] + 9C_\pi^2 C_{\text{stab}}^4 \sigma_{\mathbf{u}}^4,$$

for any $\varepsilon \leq \sigma_{\mathbf{u}} \sqrt{\lambda_{\min}(\mathbf{W}_{1:t}^z(\mathbf{x}_1^{\pi^*}))/2}$.

Proof. Toward upper-bounding the left-hand side of the desired inequality, we have:

$$\begin{aligned} & \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \varepsilon \mathbf{w})\|^2 \\ & \leq \sup_{\|\mathbf{w}\| \leq 1} 2\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) + \varepsilon \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) \mathbf{w}\|^2 + 8C_\pi^2 \varepsilon^4 \\ & \leq 4\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|^2 + \sup_{\|\mathbf{w}\| \leq 1} 4\|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) \mathbf{w}\|^2 \varepsilon^2 + 8C_\pi^2 \varepsilon^4 \\ & \leq 4\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|^2 + 4\|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|_{\text{op}}^2 \varepsilon^2 + 8C_\pi^2 \varepsilon^4, \end{aligned} \quad (\text{B.6})$$

where we use the fact that $\hat{\pi} - \pi^*$ is at worst $2C_\pi$ -smooth, and repeatedly apply $(a+b)^2 \leq 2a^2 + 2b^2$. We now lower bound $\mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2]$. Recall the linear and residual decomposition of $\tilde{\mathbf{x}}_t = \tilde{\mathbf{x}}_t^{\text{lin}} + \tilde{\mathbf{x}}_t^{\text{res}}$ from [Proposition B.2](#). Applying the C_π -smoothness of $\hat{\pi}$ and π^* , we have:

$$\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t) = (\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) + \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top (\tilde{\mathbf{x}}_t^{\text{lin}} - \mathbf{x}_t^{\pi^*} + \tilde{\mathbf{x}}_t^{\text{res}}) + \mathbf{r}_t^{\pi^*},$$

where $\|\mathbf{r}_t^\pi\| \leq 2C_\pi \|\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}\|^2 \leq 2C_\pi C_{\text{stab}}^2 \sigma_{\mathbf{u}}^2$ by applying C_π -smoothness and (C_{ISS}, ρ) -EISS (Definition 2.1) under $\sigma_{\mathbf{u}}$ -bounded input perturbations. Therefore, we may lower bound:

$$\begin{aligned} \|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2 &\geq \frac{1}{2} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) + \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top (\tilde{\mathbf{x}}_t^{\text{lin}} - \mathbf{x}_t^{\pi^*})\|^2 - \|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \tilde{\mathbf{x}}_t^{\text{res}} + \mathbf{r}_t^\pi\|^2 \\ &\geq \frac{1}{2} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) + \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top (\tilde{\mathbf{x}}_t^{\text{lin}} - \mathbf{x}_t^{\pi^*})\|^2 \\ &\quad - 2\|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \tilde{\mathbf{x}}_t^{\text{res}}\|^2 - 8C_\pi^2 C_{\text{stab}}^4 \sigma_{\mathbf{u}}^4. \end{aligned}$$

Taking the expectation on both sides, we have:

$$\begin{aligned} \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2] &\geq \frac{1}{2} \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) + \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top (\tilde{\mathbf{x}}_t^{\text{lin}} - \mathbf{x}_t^{\pi^*})\|^2] \\ &\quad - 2\mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \tilde{\mathbf{x}}_t^{\text{res}}\|^2] - 8C_\pi^2 C_{\text{stab}}^4 \sigma_{\mathbf{u}}^4. \end{aligned}$$

Notably, $\mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\tilde{\mathbf{x}}_t^{\text{lin}} - \mathbf{x}_t^{\pi^*}] = \mathbf{0}$, and thus the first term on the right-hand side can be expanded to yield:

$$\begin{aligned} &\mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) + \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top (\tilde{\mathbf{x}}_t^{\text{lin}} - \mathbf{x}_t^{\pi^*})\|^2] \\ &= \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|^2] + \text{tr} \left(\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [(\tilde{\mathbf{x}}_t^{\text{lin}} - \mathbf{x}_t^{\pi^*})(\tilde{\mathbf{x}}_t^{\text{lin}} - \mathbf{x}_t^{\pi^*})^\top] \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) \right) \\ &= \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|^2] + \text{tr} \left(\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}) \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) \right), \end{aligned}$$

On the other hand, expanding the second term yields:

$$\begin{aligned} \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \tilde{\mathbf{x}}_t^{\text{res}}\|^2] &= \text{tr} \left(\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\tilde{\mathbf{x}}_t^{\text{res}} \tilde{\mathbf{x}}_t^{\text{res}\top}] \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) \right) \\ &\leq \text{tr} \left(\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) \right) C_{\text{stab}}^6 C_{\mathbf{r}}^2 \sigma_{\mathbf{u}}^4, \end{aligned}$$

where we applied Proposition B.2 for the second line. Therefore, for sufficiently small noise level:

$$\sigma_{\mathbf{u}}^2 \leq \frac{1}{8} C_{\text{stab}}^6 C_{\mathbf{r}}^{-2} \lambda_{\min}(\mathbf{W}_{1:t}^{\mathbf{z}}(\mathbf{x}_1^{\pi^*})),$$

we may combine the first and second terms to yield:

$$\begin{aligned} &\mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2] \\ &\geq \frac{1}{2} \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|^2] + \frac{1}{4} \text{tr} \left(\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}) \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) \right) - 8C_\pi^2 C_{\text{stab}}^4 \sigma_{\mathbf{u}}^4 \\ &\geq \frac{1}{2} \left(\mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|^2] + \frac{1}{2} \lambda_{\min}(\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*})) \|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|_{\text{op}}^2 \right) - 8C_\pi^2 C_{\text{stab}}^4 \sigma_{\mathbf{u}}^4, \end{aligned}$$

where we used the elementary inequalities $\text{tr}(\mathbf{P}\mathbf{Q}) \geq \lambda_{\min}(\mathbf{P})\text{tr}(\mathbf{Q}) \geq \lambda_{\min}(\mathbf{P})\lambda_{\max}(\mathbf{Q})$, for any $\mathbf{P} \succ \mathbf{0}$, $\mathbf{Q} \geq \mathbf{0}$. Notably, the validity of this inequality rests on $\mathbf{P} \triangleq \mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}) \succ \mathbf{0}$ granted by Assumption B.1. Rearranging (B.6) yields:

$$\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|^2 + \|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|_{\text{op}}^2 \varepsilon^2 \geq \frac{1}{4} \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \varepsilon \mathbf{w})\|^2 - 2C_\pi^2 \varepsilon^4.$$

For $\varepsilon^2 \leq \frac{1}{2} \lambda_{\min}(\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*})) = \frac{\sigma_{\mathbf{u}}^2}{2} \lambda_{\min}(\mathbf{W}_{1:t}^{\mathbf{z}}(\mathbf{x}_1^{\pi^*}))$, plugging this into the above sequence of inequalities yields:

$$\begin{aligned} \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2] &\geq \frac{1}{2} \left(\frac{1}{4} \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \varepsilon \mathbf{w})\|^2 - 2C_{\pi}^2 \varepsilon^4 \right) - 8C_{\pi}^2 C_{\text{stab}}^4 \sigma_{\mathbf{u}}^4 \\ &\geq \frac{1}{16} \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \varepsilon \mathbf{w})\|^2 - \frac{1}{2} C_{\pi}^2 \varepsilon^4 - 8C_{\pi}^2 C_{\text{stab}}^4 \sigma_{\mathbf{u}}^4. \end{aligned}$$

We have trivially that $\varepsilon^2 \leq \frac{1}{2} \lambda_{\min}(\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*})) \leq \frac{\sigma_{\mathbf{u}}^2}{2} C_{\text{stab}}^2$, and thus rearranging the inequality yields the desired inequality:

$$\sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \varepsilon \mathbf{w})\|^2 \leq 16 \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2] + 9C_{\pi}^2 C_{\text{stab}}^4 \sigma_{\mathbf{u}}^4.$$

□

Therefore, using [Lemma B.4](#) to certify the tube condition in [Proposition B.3](#) yields the (suboptimal) imitation guarantee.

Suboptimal Proposition 4.2. Let [Assumption 4.1](#) hold, and let $\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}) \succeq \underline{\lambda}_{\mathbf{W}} \mathbf{I}_{d_{\mathbf{u}}}$, $t \geq 2$ w.p. 1 over $\mathbf{x}_1^{\pi^*} \sim D$ for some $\underline{\lambda}_{\mathbf{W}} > 0$. Let $\hat{\pi}$ be a C_{π} -smooth candidate policy. For $\sigma_{\mathbf{u}}^2$ that satisfies $\sigma_{\mathbf{u}}^2 \lesssim O_{\star}(\text{poly}(1/C_{\pi}, 1/C_{\text{reg}})) \underline{\lambda}_{\mathbf{W}}$, we have:

$$\mathbf{J}_{\text{TRAJ}, T}(\hat{\pi}) \lesssim O_{\star}(T) \underline{\lambda}_{\mathbf{W}}^{-1} \left(\frac{1}{\sigma_{\mathbf{u}}^2} \mathbf{J}_{\text{DEMO}, T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_{\mathbf{u}}}) + C_{\pi}^2 C_{\text{stab}}^2 \sigma_{\mathbf{u}}^2 \right).$$

Proof of Suboptimal Proposition 4.2. Using the identity for non-negative random variable Z supported on $[0, 1]$, $\mathbb{E}[Z] = \int_0^1 \mathbb{P}[Z > \varepsilon] d\varepsilon$, we have:

$$\begin{aligned} \mathbb{E}_{\mathbf{x}_1 \sim \mathbb{P}_{\mathbf{x}_1^{\pi^*}}} \left[\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 \wedge 1 \right] &= \int_0^1 \mathbb{P} \left[\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 > \varepsilon \right] d\varepsilon \\ &= \int_0^{\tau} \mathbb{P} \left[\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 > \varepsilon \right] d\varepsilon + \int_{\tau}^1 \mathbb{P} \left[\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 > \varepsilon \right] d\varepsilon \\ &\leq \int_0^{\tau} \mathbb{P} \left[\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 > \varepsilon \right] d\varepsilon + \mathbb{P} \left[\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 > \tau \right] \end{aligned}$$

where we choose a splitting point $\tau \in [0, 1]$ to be determined later. Now, applying [Proposition B.3](#) yields:

$$\begin{aligned} &\int_0^{\tau} \mathbb{P} \left[\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 > \varepsilon \right] d\varepsilon + \mathbb{P} \left[\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 > \tau \right] \\ &\leq \int_0^{\tau} \mathbb{P} \left[\max_{1 \leq t \leq T} \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \sqrt{\varepsilon} \mathbf{w})\|^2 > c_{\text{stab}}^2 \varepsilon \right] d\varepsilon \\ &\quad + \mathbb{P} \left[\max_{1 \leq t \leq T} \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \sqrt{\tau} \mathbf{w})\|^2 > c_{\text{stab}}^2 \tau \right]. \end{aligned}$$

For the first term, we have:

$$\begin{aligned}
& \int_0^\tau \mathbb{P} \left[\max_{1 \leq t \leq T} \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \sqrt{\varepsilon} \mathbf{w})\|^2 > c_{\text{stab}}^2 \varepsilon \right] d\varepsilon \\
& \leq \int_0^\tau \mathbb{P} \left[\max_{1 \leq t \leq T} \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \sqrt{\tau} \mathbf{w})\|^2 > c_{\text{stab}}^2 \varepsilon \right] d\varepsilon \\
& \leq \min \left\{ \tau, C_{\text{stab}}^2 \mathbb{E}_{\mathbf{x}_1} \left[\max_{1 \leq t \leq T} \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \tau \mathbf{w})\|^2 \right] \right\},
\end{aligned}$$

where the last line arises from combining the trivial bound $\int_0^\tau \mathbb{P}[Z > \varepsilon] d\varepsilon \leq \tau$ and by performing the variable substitution $\varepsilon' = c_{\text{stab}}^2 \varepsilon$, then applying the identity $\mathbb{E}[Z] = \int_0^1 \mathbb{P}[Z > \varepsilon'] d\varepsilon'$. Therefore, setting $\tau \leq \tilde{\sigma}_{\mathbf{u}}^2$, we apply [Lemma B.4](#) to get:

$$\begin{aligned}
& \int_0^\tau \mathbb{P} \left[\max_{1 \leq t \leq T} \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \varepsilon \mathbf{w})\|^2 > \frac{1-\rho}{C_{\text{ISS}}} \varepsilon \right] d\varepsilon \\
& \leq \min \left\{ \tau, C_{\text{stab}}^2 \mathbb{E}_{\mathbf{x}_1} \left[\max_{1 \leq t \leq T} 16 \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2] + \bar{C} C_\pi^2 \sigma_{\mathbf{u}}^4 \right] \right\} \\
& \leq \min \left\{ \tau, C_{\text{stab}}^2 \bar{C} C_\pi^2 \sigma_{\mathbf{u}}^4 + 16 C_{\text{stab}}^2 \sum_{t=1}^T \mathbb{E}_{\mathbf{x}_t} [\|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2] \right\}.
\end{aligned}$$

For the second term, we apply Markov's inequality and similarly bound:

$$\begin{aligned}
& \mathbb{P} \left[\max_{1 \leq t \leq T} \sup_{\|\mathbf{w}\| \leq 1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \sqrt{\tau} \mathbf{w})\|^2 > c_{\text{stab}}^2 \tau \right] \\
& \leq C_{\text{stab}}^2 \tau^{-1} \mathbb{E}_{\mathbf{x}_1} \left[\max_{1 \leq t \leq T} 16 \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2] + 9 C_\pi^2 C_{\text{stab}}^4 \sigma_{\mathbf{u}}^4 \right] \\
& \leq C_{\text{stab}}^2 \tau^{-1} \left(9 C_\pi^2 C_{\text{stab}}^4 \sigma_{\mathbf{u}}^4 + 16 \sum_{t=1}^T \mathbb{E}_{\mathbf{x}_t} [\|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2] \right).
\end{aligned}$$

Combining the two bounds and setting $\tau = \tilde{\sigma}_{\mathbf{u}}^2$ yields a bound on $\mathbb{E}_{\mathbf{x}_1 \sim \mathbb{P}_{\mathbf{x}_1^{\pi^*}}} [\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 \wedge 1]$ in terms of $\mathbf{J}_{\text{DEMO},2,T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_{\mathbf{u}}})$ and an additive drift term. By summing over each $1 \leq t \leq T$, we get a bound on $\mathbf{J}_{\text{TRAJ},2,T}^{\mathbf{x}}(\hat{\pi})$, accruing a T factor. Now, by [Lemma A.1](#), we have:

$$\mathbf{J}_{\text{TRAJ},2,T}(\hat{\pi}) \leq (1 + 4L_\pi^2) \mathbf{J}_{\text{TRAJ},2,T}^{\mathbf{x}}(\hat{\pi}) + 4\mathbf{J}_{\text{DEMO},2,T}(\hat{\pi}; \mathbb{P}_{\pi^*}).$$

It remains to relate $\mathbf{J}_{\text{DEMO},2,T}(\hat{\pi}; \mathbb{P}_{\pi^*})$ to $\mathbf{J}_{\text{DEMO},2,T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_{\mathbf{u}}})$. Since the injected noise is by definition $\sigma_{\mathbf{u}}$ -bounded, applying (C_{ISS}, ρ) -EISS of (π^*, f) yields w.p. 1 over any $\mathbf{x}_1^{\pi^*}$ and $\{\mathbf{z}\} \stackrel{\text{i.i.d.}}{\sim} \mathcal{D}(\mathbf{0}, \Sigma_{\mathbf{z}})$:

$$\begin{aligned}
\|\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}\| & \leq C_{\text{ISS}} \sum_{s=1}^{t-1} \rho^{t-1-s} \|\sigma_{\mathbf{u}} \mathbf{z}_s\| \\
& \leq C_{\text{stab}} \sigma_{\mathbf{u}}.
\end{aligned}$$

In other words, for a given $\mathbf{z}_t \sim \mathcal{D}(\mathbf{0}, \Sigma_{\mathbf{z}})$ we always have:

$$\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\| \leq \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \sigma_{\mathbf{u}} \mathbf{z}_t)\| + 2L_\pi C_{\text{stab}} \sigma_{\mathbf{u}}.$$

Squaring both sides and taking an expectation yields the following bound on $\mathbf{J}_{\text{DEMO},2,T}(\hat{\pi}; \mathbb{P}_{\pi^*})$:

$$\mathbf{J}_{\text{DEMO},2,T}(\hat{\pi}; \mathbb{P}_{\pi^*}) \lesssim \mathbf{J}_{\text{DEMO},2,T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_u}) + T L_{\pi}^2 C_{\text{stab}}^2 \sigma_u^2.$$

Putting the pieces together, we have:

$$\begin{aligned} \mathbf{J}_{\text{TRAJ},2,T}(\hat{\pi}) &\leq (1 + 4L_{\pi}^2) \mathbf{J}_{\text{TRAJ},2,T}^{\mathbf{x}}(\hat{\pi}) + 4\mathbf{J}_{\text{DEMO},2,T}(\hat{\pi}; \mathbb{P}_{\pi^*}) \\ &\lesssim (1 + 4L_{\pi}^2) C_{\text{stab}}^2 \underline{\lambda}_{\mathbf{z}}^{-1} T \left(\frac{1}{\sigma_u^2} \mathbf{J}_{\text{DEMO},2,T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_u}) + C_{\pi}^2 C_{\text{stab}}^4 \sigma_u^2 \right) \\ &\quad + \mathbf{J}_{\text{DEMO},2,T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_u}) + T L_{\pi}^2 C_{\text{stab}}^2 \sigma_u^2. \end{aligned}$$

When $\mathcal{D}(\mathbf{0}, \Sigma_{\mathbf{z}})$ is the uniform distribution over the ball, we have $\underline{\lambda}_{\mathbf{z}} \approx \underline{\lambda}_{\mathbf{W}}/d_u$. Lumping terms together, this completes the proof of [Suboptimal Proposition 4.2](#). \square

This result says that if noise injection fully excites the state space, then the trajectory error is bounded by the on-expert error evaluated on the noise-injected law $\mathbb{P}_{\pi^*, \sigma_u}$ plus a higher-order error term from smoothness. Note that simply regressing on the expert trajectories without noise injection, even the smooth one-step controllable case considered here, can suffer from exponential compounding error (see [Simchowitz et al. \[2025, Theorem 4\]](#)). Though this is a marked improvement upon vanilla behavior cloning, this set-up leaves open a couple deficiencies. Firstly, performing behavior cloning on $\mathbb{P}_{\pi^*, \sigma_u}$ yields a drift term $\approx \sigma_u^2$ that persists even when $\mathbf{J}_{\text{DEMO},T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_u})$ is small; this introduces a trade-off on the noise-scale, where larger σ_u benefits the excitation, but exacerbates the drift. We demonstrate in [Appendix B.4](#) that this additive factor is fundamental. Secondly, one-step controllability—and in a similar vein persistency of excitation—is a strong condition (e.g. requires $d_u = d_x$); typically we do not expect inputs to be able to excite every mode in a system, let alone instantaneously.

B.2 Departing from Controllability

We now consider the case where we lack controllability, one-step or otherwise. In other words, the linear controllability Gramians need not be full-rank: $\text{rank}(\mathbf{W}_{1:t}^u(\mathbf{x}_1^{\pi^*})) < d_x$. Furthermore, as promised in the body, we hope to lift the inverse dependence on the smallest positive eigenvalue of controllability Gramian, including when it is rank-deficient. On the technical front, a few barriers are present. Firstly, the state-covariance bound in [Corollary B.1](#) imposes a constraint on σ_u scaling with the smallest positive eigenvalue of $\mathbf{W}_{1:t}^z$. Secondly, [Proposition B.3](#) requires certifying that $\hat{\pi}$ and π^* match on a (full-dimensional) ball around the expert trajectory, and subsequently the “expectation-to-uniform” bound in [Lemma B.4](#) requires a full-rank covariance.

Given these technical difficulties, we introduce the notion of the “reachable subspace” under the linearized system under the expert.

Definition B.1. Fix any $\mathbf{x}_1^{\pi^*} \sim D$. Recall the expert linearizations from [Eq. \(B.1\)](#). Define the *reachable subspace* of the expert closed-loop system at time t :

$$\mathcal{R}_t^{\pi^*} \triangleq \left\{ \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} \mathbf{B}_{ts} \mathbf{u}_s \mid \{\mathbf{u}_s\}_{s=1}^{t-1} \subset \mathbb{R}^{d_u} \right\}.$$

The following facts hold:

- $\mathcal{R}_t^{\pi^*}$ is a linear subspace of \mathbb{R}^{d_x} .
- Given any positive-definite $\Sigma \succ \mathbf{0}$, the associated controllability Gramian satisfies $\text{rank}\left(\sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} \mathbf{B}_s \Sigma \mathbf{B}_s^\top \mathbf{A}_{s+1:t}^{\text{cl}\top}\right) = \dim(\mathcal{R}_t^{\pi^*})$ for each $t \geq 1$.

Let $\{(\lambda_{i,t}, \mathbf{v}_{i,t})\}_{i=1}^{d_x}$ be the eigenvalues and vectors of $\mathbf{W}_{1:t}^{\mathbf{u}}$, $t \geq 2$.⁸ Let us further define the reachable subspace *truncated at λ* :

$$\mathcal{R}_t^{\pi^*}(\lambda) \triangleq \text{span}\{\mathbf{v}_{i,t} : \lambda_{i,t} \geq \lambda\},$$

as well as the corresponding orthogonal projection matrix $\mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)}$. We also abuse notation and denote $\mathcal{R}_t^{\pi^*}(\lambda)^\perp$ as the subspace component of $\mathcal{R}_t^{\pi^*}$ orthogonal to $\mathcal{R}_t^{\pi^*}(\lambda)$.

In line with the body, we will consider $\mathcal{D}(\mathbf{0}, \Sigma_z) = \text{Unif}(\mathbb{B}^{d_u}(1))$, such that $\mathbf{W}_{1:t}^z \succeq \frac{1}{3d_u} \mathbf{W}_{1:t}^{\mathbf{u}}$. As previewed in the body, the main guiding intuition moving forward is as follows: 1. by smoothness of the dynamics, most of the error should be contained in the (linearized) reachable subspace, 2. the small eigendirections of the controllability Gramian are precisely those that are hard-to-excite, and thus should accumulate compounding errors slowly enough to “ignore” them. We start by proving a restricted “Jacobian sketching” result (cf. [Proposition 4.4](#)). We note that though we present [Proposition 4.3](#) first in the body, we will in fact use an extended version of it that relies on the subsequent result.

Proposition B.5 (Full ver. of [Proposition 4.4](#)). *Let [Assumption 4.1](#) hold. For $\mathbf{x}_1^{\pi^*} \sim D$, define $\mathcal{R}_t^{\pi^*}(\lambda) \triangleq \text{span}\{\mathbf{v}_{i,t} : \lambda_{i,t} \geq \lambda\}$ and $\mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)}$ as in [Definition B.1](#), for some $\lambda \geq \lambda_{\min}^+(\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}))$. Then, for $\sigma_{\mathbf{u}}$ satisfying:*

$$\sigma_{\mathbf{u}} \lesssim \min \left\{ \lambda d_u^{-1} c_{\text{stab}}^4 C_{\mathbf{r}}^{-1}, c_{\text{stab}} \frac{\sqrt{1 + 4C_{\mathbf{K}}^2}}{C_{\pi}} \right\} = O_*(\lambda).$$

we have the following bound for each $t \geq 2$:

$$\|\mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)} \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|_{\text{op}}^2 \lesssim \frac{d_u}{\sigma_{\mathbf{u}}^2 \lambda} \left(\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|^2 + \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} \|(\hat{\pi} - \pi^*)(\tilde{\mathbf{x}}_t)\|^2 \right) + \frac{d_u \sigma_{\mathbf{u}}^2}{\lambda} C_{\pi}^2 C_{\text{stab}}^4.$$

We note that [Proposition 4.4](#) is recovered by applying an expectation over $\mathbf{x}_1^{\pi^*}$ on both sides of the inequality.

Proof of [Proposition B.5](#). First, we consider the following adaptation of [Corollary B.1](#)

Corollary B.2. *Let [Assumption 4.1](#) hold and $C_{\mathbf{r}}$ be as defined in [Proposition B.2](#). Fix any $t \geq 2$. For $\lambda \geq \lambda_{\min}^+(\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}))$, set $\mathbf{P} = \mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)}$ as in [Definition B.1](#). As long as:*

$$\sigma_{\mathbf{u}} \lesssim \min \left\{ \lambda d_u^{-1} c_{\text{stab}}^4 C_{\mathbf{r}}^{-1}, c_{\text{stab}} \frac{\sqrt{1 + 4C_{\mathbf{K}}^2}}{C_{\pi}} \right\},$$

the following holds almost surely over $\tilde{\mathbf{x}}_1 = \mathbf{x}_1^{\pi^*} \sim D$ and $\{\mathbf{z}_s\} \stackrel{\text{i.i.d.}}{\sim} \mathcal{D}(\mathbf{0}, \Sigma_z)$:

$$\mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} \left[\mathbf{P} (\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}) (\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})^\top \mathbf{P} \right] \geq \frac{\sigma_{\mathbf{u}}^2}{2} \mathbf{P} \mathbf{W}_{1:t}^z(\mathbf{x}_1^{\pi^*}) \mathbf{P}. \quad (\text{B.7})$$

⁸Though we omit it for clarity, recall all these quantities implicitly condition on $\mathbf{x}_1^{\pi^*}$.

The proof of [Corollary B.2](#) follows from a one-line modification in the proof of [Corollary B.1](#), where instead of requiring Weyl's inequality to hold over all positive eigenvalues $k = 1, \dots, \text{rank}(\mathbf{W}_{1:t}^u)$, we need only to consider up to $k = 1, \dots, p$, $p = \dim(\mathcal{R}_t^{\pi^*}(\lambda))$, for which $\lambda_p(\mathbf{W}_{1:t}^z) \gtrsim d_u^{-1} \lambda_p(\mathbf{W}_{1:t}^u) \geq d_u^{-1} \lambda$.

We proceed by applying the C_π -smoothness of $\hat{\pi}$ and π^* , we have:

$$\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t) = (\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) + \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top (\tilde{\mathbf{x}}_t^{\text{lin}} - \mathbf{x}_t^{\pi^*} + \tilde{\mathbf{x}}_t^{\text{res}}) + \mathbf{r}_t^\pi,$$

where $\|\mathbf{r}_t^\pi\| \leq 2C_\pi \|\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*}\| \leq 2C_\pi C_{\text{stab}}^2 \sigma_u^2$ by applying C_π -smoothness and (C_{ISS}, ρ) -EISS ([Definition 2.1](#)) under σ_u -bounded input perturbations. Therefore, we may lower bound:

$$\begin{aligned} \|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\| &\geq \|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top (\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})\| - \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) + \mathbf{r}_t^\pi\| \\ &\geq \|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top (\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})\| - \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\| - 2C_\pi C_{\text{stab}}^2 \sigma_u^2. \end{aligned}$$

Rearranging the above inequality, squaring both sides, and applying the inequality $(a + b + c)^2 \leq 3(a^2 + b^2 + c^2)$ we have:

$$\|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top (\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})\|^2 \leq 3(\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|^2 + \|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2 + 2C_\pi^2 C_{\text{stab}}^4 \sigma_u^4).$$

Taking an expectation over the noise injection on both sides, we may apply [Corollary B.2](#) on the left-hand side: for σ_u satisfying the requirements therein, we have:

$$\begin{aligned} &\mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [\|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top (\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})\|^2] \\ &= \text{tr} \left(\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [(\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})(\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})^\top] \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) \right) \\ &\geq \text{tr} \left(\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)} \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} [(\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})(\tilde{\mathbf{x}}_t - \mathbf{x}_t^{\pi^*})^\top] \mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)} \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) \right) \\ &\geq \frac{\sigma_u^2}{2} \text{tr} \left(\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)} \mathbf{W}_{1:t}^z(\mathbf{x}_1^{\pi^*}) \mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)} \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*}) \right) \\ &\gtrsim \sigma_u^2 \|\mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)} \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|_{\text{op}}^2 d_u^{-1} \lambda, \end{aligned}$$

where we applied [Corollary B.2](#) on the second-to-last line, and for the last line we used by definition $\lambda_{\min}^+(\mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)} \mathbf{W}_{1:t}^z(\mathbf{x}_1^{\pi^*}) \mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)}) \gtrsim d_u^{-1} \lambda_{\min}^+(\mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)} \mathbf{W}_{1:t}^u(\mathbf{x}_1^{\pi^*}) \mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)}) \gtrsim d_u^{-1} \lambda$. Thus, re-arranging the inequalities, we have

$$\begin{aligned} &\|\mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda)} \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|_{\text{op}}^2 \\ &\lesssim \frac{d_u}{\sigma_u^2 \lambda} \left(\|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|^2 + \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^{\pi^*}} \|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2 \right) + \frac{d_u \sigma_u^2}{\lambda} C_\pi^2 C_{\text{stab}}^4, \end{aligned}$$

which completes the result. □

In light of [Proposition B.5](#), we have demonstrated that small estimation error along both un-noised and noise-injected states implies a first-order closeness of $\hat{\pi}$ and π^* along a subspace of our choosing. However, by choosing the excitation threshold λ that we guarantee closeness above, we do not track: 1. error in the reachable subspace below the λ threshold, 2. error for non-linearity. As stated, [Proposition B.3](#) requires uniform closeness on a ε -scaled unit ball, which [Proposition B.5](#) does not grant. Our next step is to prove the full version of [Proposition 4.3](#).

Proposition B.6. Let [Assumption 4.1](#) hold. For any initial state \mathbf{x}_1 , let $\mathbf{x}_1^{\hat{\pi}} = \mathbf{x}_1^{\pi^*} = \mathbf{x}_1$, and consider the closed-loop trajectories generated by $\hat{\pi}$ and π^* . Define the constant $C_{\text{rem}} \triangleq 2C_{\text{reg}}(3 + 2C_K^2 + 2C_\pi^2)$. Fix any sequence $\{\lambda_t\}_{t=1}^{T-1}$, where each $\lambda_t \in [\lambda_{\min}^+(\mathbf{W}_{1:t}^u(\mathbf{x}_1^{\pi^*})), \lambda_{\max}(\mathbf{W}_{1:t}^u(\mathbf{x}_1^{\pi^*}))]$. Then for any given $\varepsilon \in [0, 1]$, $T \in \mathbb{N}$, as long as:

$$\max_{1 \leq t \leq T-1} \sup_{\substack{\|\mathbf{w}\| \leq 1, \mathbf{w} \in \mathcal{R}_t^{\pi^*}(\lambda_t) \\ \|\mathbf{r}\| \leq 1, \mathbf{r} \in \mathcal{R}_t^{\pi^*}(\lambda_t)^\perp \\ \|\mathbf{v}\| \leq 1}} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \varepsilon \mathbf{w} + \frac{C_{\text{stab}}}{\log(1/\rho)} \sqrt{\lambda_t} \varepsilon \mathbf{r} + C_{\text{rem}} \varepsilon^2 \mathbf{v})\| \leq C_{\text{stab}} \varepsilon,$$

we are guaranteed $\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\| \leq \varepsilon$.

Proof of Proposition B.6. We prove this result by induction. Fix any $\varepsilon \in [0, 1]$. Define the quantity $C_\perp \triangleq \frac{1-\rho}{C_{\text{ISS}} \log(1/\rho)}$. Further define the shorthands $\mathcal{P}_t = \mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda_t)}$, and the relative orthogonal component $\mathcal{P}_t^\perp = \mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda_t)}^\perp$. Let us for each timestep t define the set

$$\mathcal{V}_t \triangleq \left\{ \varepsilon \mathbf{w} + \sqrt{\lambda_t} C_\perp \varepsilon \mathbf{r} + C_{\text{rem}} \varepsilon^2 \mathbf{v} : \|\mathbf{w}\|, \|\mathbf{r}\|, \|\mathbf{v}\| \leq 1, \mathbf{w} \in \mathcal{R}_t^{\pi^*}(\lambda_t), \mathbf{r} \in \mathcal{R}_t^{\pi^*}(\lambda_t)^\perp \right\}.$$

In addition to the statement of [Proposition B.6](#), we claim that $\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*} \in \mathcal{V}_t$ for each $t = 1, \dots, T$. Considering the base-case $T = 2$: since $\mathbf{x}_1^{\pi^*} = \mathbf{x}_1^{\hat{\pi}}$ by construction, and thus $\hat{\pi}(\mathbf{x}_1^{\hat{\pi}}) = \hat{\pi}(\mathbf{x}_1^{\pi^*})$, by assumption this satisfies $\|\hat{\pi}(\mathbf{x}_1^{\pi^*}) - \pi^*(\mathbf{x}_1^{\pi^*})\| \leq \frac{1-\rho}{C_{\text{ISS}}} \varepsilon$. By applying (C_{ISS}, ρ) -EISS, we have

$$\|\mathbf{x}_2^{\hat{\pi}} - \mathbf{x}_2^{\pi^*}\| \leq C_{\text{ISS}} \|\hat{\pi}(\mathbf{x}_1^{\pi^*}) - \pi^*(\mathbf{x}_1^{\pi^*})\| \leq C_{\text{ISS}} \frac{1-\rho}{C_{\text{ISS}}} \varepsilon \leq \varepsilon.$$

Furthermore, recalling the definitions in [Lemma B.1](#), we apply the C_{reg} -smoothness of the dynamics f and take a second-order Taylor expansion around $(\mathbf{x}, \mathbf{u}) = (\mathbf{x}_1^{\pi^*}, \pi^*(\mathbf{x}_1^{\pi^*}))$ to yield:

$$\mathbf{x}_2^{\hat{\pi}} - \mathbf{x}_2^{\pi^*} = \mathbf{B}_1(\hat{\pi} - \pi^*)(\mathbf{x}_1^{\hat{\pi}}) + \mathbf{r}_1^{\mathbf{x}}.$$

We observe this implies $\mathbf{B}_1(\hat{\pi} - \pi^*)(\mathbf{x}_1^{\hat{\pi}}) \in \mathcal{R}_1^{\pi^*}$, and [Lemma B.1](#) implies $\|\mathbf{B}_1\| \leq C_{\text{ISS}}$. On the other hand, since $\mathbf{W}_{1:2}^u = \mathbf{B}_1 \mathbf{B}_1^\top$, we know $\|\mathcal{P}_1^\perp \mathbf{B}_1\| \leq \sqrt{\lambda_1}$. Since $\mathbf{x}_1^{\hat{\pi}} = \mathbf{x}_1^{\pi^*}$, we have:

$$\begin{aligned} \|\mathcal{P}_1 \mathbf{B}_1(\hat{\pi} - \pi^*)(\mathbf{x}_1^{\hat{\pi}})\| &\leq C_{\text{ISS}} \|(\hat{\pi} - \pi^*)(\mathbf{x}_1^{\hat{\pi}})\| \leq C_{\text{ISS}} \frac{1-\rho}{C_{\text{ISS}}} \varepsilon \leq \varepsilon \\ \|\mathcal{P}_1^\perp \mathbf{B}_1(\hat{\pi} - \pi^*)(\mathbf{x}_1^{\hat{\pi}})\| &\leq \sqrt{\lambda_1} \|(\hat{\pi} - \pi^*)(\mathbf{x}_1^{\hat{\pi}})\| \leq \sqrt{\lambda_1} \frac{1-\rho}{C_{\text{ISS}}} \varepsilon \leq \sqrt{\lambda_1} C_\perp \varepsilon \\ \|\mathbf{r}_1^{\mathbf{x}}\| &\leq C_{\text{reg}} (\|\mathbf{x}_1^{\hat{\pi}} - \mathbf{x}_1^{\pi^*}\|^2 + \|\hat{\pi}(\mathbf{x}_1^{\hat{\pi}}) - \pi^*(\mathbf{x}_1^{\pi^*})\|^2) \leq C_{\text{reg}} \left(\frac{1-\rho}{C_{\text{ISS}}} \varepsilon \right)^2 \leq C_{\text{rem}} \varepsilon^2, \end{aligned}$$

which implies $\mathbf{x}_2^{\hat{\pi}} - \mathbf{x}_2^{\pi^*} \in \mathcal{V}_2$. This completes the base-case.

Now for $T > 2$, we assume the statement holds for $T-1$; in particular, we have $\max_{1 \leq t \leq T-1} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\| \leq \varepsilon$ and $\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*} \in \mathcal{V}_t$ for $t \in [T-1]$. Then, by (C_{ISS}, ρ) -EISS we have:

$$\|\mathbf{x}_T^{\hat{\pi}} - \mathbf{x}_T^{\pi^*}\| \leq C_{\text{ISS}} \sum_{t=1}^{T-1} \rho^{T-1-t} \|\hat{\pi}(\mathbf{x}_t^{\hat{\pi}}) - \pi^*(\mathbf{x}_t^{\pi^*})\|$$

$$\begin{aligned}
&\leq C_{\text{ISS}} \sum_{t=1}^{T-1} \rho^{T-1-t} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \Delta_{\mathbf{x}_t})\| \\
&\leq C_{\text{ISS}} \sum_{t=1}^{T-1} \rho^{T-1-t} \left(\frac{1-\rho}{C_{\text{ISS}}} \varepsilon \right), \quad (\text{Inductive hypothesis})
\end{aligned}$$

where $\Delta_{\mathbf{x}_t} \triangleq \mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}$ and the last line uses the induction hypothesis that each $\Delta_{\mathbf{x}_t} \in \mathcal{V}_t$, $t \in [T-1]$. This completes the first part of the induction step. It remains to show $\mathbf{x}_T^{\hat{\pi}} - \mathbf{x}_T^{\pi^*} \in \mathcal{V}_T$. From the definition of the linearizations Eq. (B.2), we may write:

$$\begin{aligned}
\mathbf{x}_T^{\hat{\pi}} - \mathbf{x}_T^{\pi^*} &= \sum_{s=1}^{T-1} \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s (\hat{\pi} - \pi^*)(\mathbf{x}_s^{\hat{\pi}}) + \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{r}_s^{\mathbf{x}} \\
&= \mathcal{P}_T \sum_{s=1}^{T-1} \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s (\hat{\pi} - \pi^*)(\mathbf{x}_s^{\hat{\pi}}) + \mathcal{P}_T^\perp \sum_{s=1}^{T-1} \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s (\hat{\pi} - \pi^*)(\mathbf{x}_s^{\hat{\pi}}) + \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{r}_s^{\mathbf{x}}
\end{aligned} \tag{B.8}$$

where $\mathbf{r}_s^{\mathbf{x}}$ are the second-order remainder terms from linearizing the dynamics around $(\mathbf{x}_s^{\pi^*}, \pi^*(\mathbf{x}_s^{\pi^*}))$ for $s \in [T-1]$. We first observe by definition $\sum_{s=1}^{T-1} \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s (\hat{\pi} - \pi^*)(\mathbf{x}_s^{\hat{\pi}}) \in \mathcal{R}_T^{\pi^*}$, i.e. the first term on the first line lies in the reachable subspace. Focusing on the first term of the second line, we may trivially bound:

$$\begin{aligned}
\left\| \mathcal{P}_T^\perp \sum_{s=1}^{T-1} \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s (\hat{\pi} - \pi^*)(\mathbf{x}_s^{\hat{\pi}}) \right\| &\leq \sum_{s=1}^{T-1} \|\mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s\| \|(\hat{\pi} - \pi^*)(\mathbf{x}_s^{\pi^*} + \Delta_{\mathbf{x}_s})\| \\
&\leq \sum_{s=1}^{T-1} C_{\text{ISS}} \rho^{T-1-s} \left(\frac{1-\rho}{C_{\text{ISS}}} \varepsilon \right) \leq \varepsilon,
\end{aligned}$$

where we used Lemma B.1 and the induction hypothesis for the last line. For the second term, we first observe that since $\mathbf{W}_{1:t}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}) \triangleq \sum_{s=1}^{t-1} \mathbf{A}_{s+1:t}^{\text{cl}} \mathbf{B}_s \mathbf{B}_s^\top \mathbf{A}_{s+1:t}^{\text{cl}^\top}$, we have $\|\mathcal{P}_T^\perp \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s\| \leq \|\mathcal{P}_T^\perp \mathbf{W}_{1:T}^{\mathbf{u}}(\mathbf{x}_1^{\pi^*}) \mathcal{P}_T^\perp\|^{1/2} \leq \sqrt{\lambda_T}$. Alternatively, we always have by Lemma B.1 $\|\mathcal{P}_T^\perp \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s\| \leq C_{\text{ISS}} \rho^{T-1-s}$. Therefore, picking any $k \in [T-1]$, we have:

$$\begin{aligned}
\left\| \mathcal{P}_T^\perp \sum_{s=1}^{T-1} \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s (\hat{\pi} - \pi^*)(\mathbf{x}_s^{\hat{\pi}}) \right\| &\leq \sum_{s=1}^{T-1} \|\mathcal{P}_T^\perp \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s\| \|(\hat{\pi} - \pi^*)(\mathbf{x}_s^{\pi^*} + \Delta_{\mathbf{x}_s})\| \\
&\leq \left(k\sqrt{\lambda_T} + \sum_{s=1}^{T-1-k} C_{\text{ISS}} \rho^{T-1-s} \right) \left(\frac{1-\rho}{C_{\text{ISS}}} \varepsilon \right) \\
&\leq \left(k\sqrt{\lambda_T} + \rho^{-k} \frac{C_{\text{ISS}}}{1-\rho} \right) \left(\frac{1-\rho}{C_{\text{ISS}}} \varepsilon \right).
\end{aligned}$$

Now, by solving for the optimal truncation point: $\min_{k \geq 1} k\sqrt{\lambda_T} + \rho^{-k} \frac{C_{\text{ISS}}}{1-\rho}$, we may upper bound the resulting value by:

$$\begin{aligned}
\min_{k \geq 1} k\sqrt{\lambda_T} + \rho^{-k} \frac{C_{\text{ISS}}}{1-\rho} &\leq \frac{\sqrt{\lambda_T}}{\log(1/\rho)} \left(1 + \log \left(\frac{\sqrt{\lambda_T}(1-\rho)}{C_{\text{ISS}} \log(1/\rho)} \right) \right) \\
&\leq \frac{\sqrt{\lambda_T}}{\log(1/\rho)},
\end{aligned}$$

where for the last line we observe that $\sqrt{\lambda_T} \leq \sqrt{\lambda_{\max}(\mathbf{W}_{1:T}^u)} \leq \frac{C_{\text{ISS}}}{1-\rho}$ by Lemma B.1, and thus $\log\left(\frac{\sqrt{\lambda_T}(1-\rho)}{C_{\text{ISS}}\log(1/\rho)}\right) \leq 0$. Therefore, we may plug this back in to yield:

$$\left\| \mathcal{P}_T^\perp \sum_{s=1}^{T-1} \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s (\hat{\pi} - \pi^*)(\mathbf{x}_s^{\hat{\pi}}) \right\| \leq \frac{\sqrt{\lambda_T}}{\log(1/\rho)} \left(\frac{1-\rho}{C_{\text{ISS}}} \varepsilon \right) = \sqrt{\lambda_T} C_\perp \varepsilon.$$

As for the last remainder term, we have:

$$\begin{aligned} \|\mathbf{A}_{s+1:T}^{\text{cl}}\| &\leq C_{\text{ISS}} \rho^{T-1-s} && \text{(Lemma B.1)} \\ \|\mathbf{r}_s^{\mathbf{x}}\| &\leq C_{\text{reg}} (\|\mathbf{x}_s^{\pi^*} - \mathbf{x}_s^{\hat{\pi}}\|^2 + \|\hat{\pi}(\mathbf{x}_s^{\hat{\pi}}) - \pi^*(\mathbf{x}_s^{\pi^*})\|^2) && \text{(Assumption 4.1)} \\ &\leq C_{\text{reg}} (\varepsilon^2 + 2\|(\hat{\pi} - \pi^*)(\mathbf{x}_s^{\pi^*} + \Delta_{\mathbf{x}_s})\|^2 + 2\|\pi^*(\mathbf{x}_s^{\pi^*} + \Delta_{\mathbf{x}_s}) - \pi^*(\mathbf{x}_s^{\pi^*})\|^2) \\ &\leq C_{\text{reg}} (1 + 2c_{\text{stab}}^2 + 2\|\nabla_{\mathbf{x}} \pi^*(\mathbf{x}_s^{\pi^*})^\top \Delta_{\mathbf{x}_s} + \mathbf{r}_s^\Delta\|^2) \varepsilon^2 \\ &\leq C_{\text{reg}} (1 + 2c_{\text{stab}}^2 + 2C_K + 2C_\pi^2 \varepsilon^2) \varepsilon^2 \\ &\leq 2C_{\text{reg}} (3 + 2C_K^2 + 2C_\pi^2) \varepsilon^2 \\ \left\| \sum_{s=1}^{T-1} \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{r}_s^{\mathbf{x}} \right\| &\leq \sum_{s=1}^{T-1} C_{\text{ISS}} \rho^{T-1-s} \|\mathbf{r}_s^{\mathbf{x}}\| \leq 2C_{\text{reg}} (3 + 2C_K^2 + 2C_\pi^2) \frac{C_{\text{ISS}}}{1-\rho} \varepsilon^2 \triangleq C_{\text{rem}} \varepsilon^2. \end{aligned}$$

Therefore, putting all the pieces back into Eq. (B.8), we have:

$$\begin{aligned} \mathbf{x}_T^{\hat{\pi}} - \mathbf{x}_T^{\pi^*} &= \mathcal{P}_T \sum_{s=1}^{T-1} \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s (\hat{\pi} - \pi^*)(\mathbf{x}_s^{\hat{\pi}}) + \mathcal{P}_T^\perp \sum_{s=1}^{T-1} \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{B}_s (\hat{\pi} - \pi^*)(\mathbf{x}_s^{\hat{\pi}}) + \mathbf{A}_{s+1:T}^{\text{cl}} \mathbf{r}_s^{\mathbf{x}} \\ &\leq \varepsilon \mathbf{w} + \sqrt{\lambda_T} C_\perp \mathbf{r} + C_{\text{rem}} \varepsilon^2 \mathbf{v}, \quad \|\mathbf{w}\|, \|\mathbf{r}\|, \|\mathbf{v}\| \leq 1, \quad \mathbf{w} \in \mathcal{R}_T^{\pi^*}(\lambda_T), \mathbf{r} \in \mathcal{R}_T^{\pi^*}(\lambda_T)^\perp \\ &\in \mathcal{V}_T. \end{aligned}$$

we have demonstrated $\mathbf{x}_T^{\hat{\pi}} - \mathbf{x}_T^{\pi^*} \in \mathcal{V}_T$, completing the induction step and thus the proof. \square

To review, we have established two key tools in Proposition B.5 and Proposition B.6, corresponding to Proposition 4.4 and Proposition 4.3 in the body, respectively. The first states that, fixing our attention to the component of the reachable subspace that is excitable above a threshold λ (to be determined in hindsight), we may bound the first-order, i.e. Jacobian error between $\hat{\pi}$ and π^* in terms of their error on the *mixture* distribution $\mathbb{P}_{\pi^*, \sigma_u, \alpha}$. The second states that, fixing an excitation level, as long as we ensure $\hat{\pi}$ matches π^* sufficiently well on the set \mathcal{V}_t for each t , which decomposes into the “excitable”, (linearly) reachable component in $\mathcal{R}^{\pi^*}(\lambda_t)$, the low-excitation (linearly) reachable component in $\mathcal{R}^{\pi^*}(\lambda_t)^\perp$, and a generic second-order term, the resulting closed-loop trajectories will remain close.

We are now ready to prove our main guarantee for noise injection.

B.3 Guarantees without Controllability: Proof of Theorem 2

We dedicate most of the effort into establishing the following result.

Proposition B.7. Let [Assumption 4.1](#) hold. Let C_r be defined as in [Corollary B.1](#) and C_{rem} as in [Proposition B.6](#). Let the noise-scale $\sigma_u > 0$ satisfy

$$\sigma_u \lesssim \min \left\{ \sqrt{\frac{\log(1/\rho)}{L_\pi d_u}} \frac{c_{\text{stab}}^3}{C_\pi}, \frac{\log(1/\rho)^2}{L_\pi^2 d_u} \frac{c_{\text{stab}}^4}{C_r}, c_{\text{stab}} \frac{\sqrt{1+4C_K^2}}{C_\pi} \right\}. \quad (\text{B.9})$$

Consider a candidate policy $\hat{\pi}$. Defining $C_{\text{traj}} \triangleq 2C_{\text{rem}}L_\pi + 6C_\pi \left(1 + \frac{1}{4} \frac{c_{\text{stab}}}{L_\pi} + C_{\text{rem}}^2\right)$, we have the following bound on the expected (clipped) trajectory error:

$$\begin{aligned} & \mathbb{E}_{\mathbf{x}_1 \sim \mathbb{P}_{\mathbf{x}_1^*}} \left[\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 \wedge 1 \right] \\ & \lesssim C_{\text{stab}}^2 \left(1 + C_{\text{traj}}^2 C_{\text{stab}}^2 + \frac{d_u L_\pi^2}{\log(1/\rho)^2 \sigma_u^2} \right) \mathbb{E}_{\mathbf{x}_1^*} \left[\max_{t \leq T-1} \|\hat{\pi}(\mathbf{x}_t^{\pi^*}) - \pi^*(\mathbf{x}_t^{\pi^*})\|^2 + \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^*} \|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2 \right] \\ & \leq C_{\text{stab}}^2 \left(1 + C_{\text{traj}}^2 C_{\text{stab}}^2 + \frac{d_u L_\pi^2}{\log(1/\rho)^2 \sigma_u^2} \right) \mathbf{J}_{\text{DEMO}, T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_u, \alpha}). \end{aligned}$$

Proof of Proposition B.7. Let us define the shorthands for the per-timestep trajectory and estimation errors:

$$\begin{aligned} r_t^{\text{traj}}(\hat{\pi}, \pi^*) & \triangleq \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 \wedge 1 \\ r_t^{\text{est}}(\hat{\pi}; \pi^*) & \triangleq \|\hat{\pi}(\mathbf{x}_t^{\pi^*}) - \pi^*(\mathbf{x}_t^{\pi^*})\|^2 \\ r_t^{\text{est}}(\hat{\pi}; \tilde{\pi}^*) & \triangleq \mathbb{E}_{\tilde{\mathbf{x}}_t | \mathbf{x}_1^*} [\|\hat{\pi}(\tilde{\mathbf{x}}_t) - \pi^*(\tilde{\mathbf{x}}_t)\|^2], \end{aligned}$$

As in [Proposition B.6](#), let us define a sequence $\{\lambda_t\}_{t=1}^{T-1}$, where each $\lambda_t \in [\lambda_{\min}^+(\mathbf{W}_{1:t}^u(\mathbf{x}_1^{\pi^*})), \lambda_{\max}(\mathbf{W}_{1:t}^u(\mathbf{x}_1^{\pi^*}))]$, as well as the truncated subspaces and projection matrices: $\mathcal{R}_t^{\pi^*}(\lambda_t)$, $\mathcal{P}_t = \mathcal{P}_{\mathcal{R}_t^{\pi^*}(\lambda_t)}$. By [Proposition B.5](#), noise injection certifies a norm bound on $\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_1^{\pi^*})$ restricted to $\mathcal{R}_t^{\pi^*}(\lambda_t)$, for each $t \geq 2$. Accordingly, we define the event:

$$\mathcal{E}_{\nabla \pi}(c) \triangleq \left\{ \max_{t \leq T-1} \|\mathcal{P}_t \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|_{\text{op}} \lesssim c \right\}.$$

We may decompose the desired quantity into:

$$\mathbb{E}_{\mathbf{x}_1^*} \left[\max_{t \leq T-1} r_t^{\text{traj}}(\hat{\pi}, \pi^*) \right] = \underbrace{\mathbb{E}_{\mathbf{x}_1^*} \left[\max_{t \leq T-1} r_t^{\text{traj}}(\hat{\pi}, \pi^*) \mathbf{1}\{\mathcal{E}_{\nabla \pi}(c_{\text{stab}})\} \right]}_{T_1} + \underbrace{\mathbb{E}_{\mathbf{x}_1^*} \left[\max_{t \leq T-1} r_t^{\text{traj}}(\hat{\pi}, \pi^*) \mathbf{1}\{\mathcal{E}_{\nabla \pi}(c_{\text{stab}})^c\} \right]}_{T_2}.$$

In addition to the requirements on σ_u in [Proposition B.5](#) for $\lambda = \lambda_t$, assume that σ_u satisfies across $t \geq 2$: $\sigma_u \lesssim \sqrt{\frac{\lambda_t}{d_u}} \frac{c_{\text{stab}}^3}{C_\pi}$, such that $\frac{d_u \sigma_u^2}{\lambda} C_\pi^2 C_{\text{stab}}^4 \lesssim c_{\text{stab}}^2$. Since $r_t^{\text{traj}}(\hat{\pi}, \pi^*) \leq 1$, we may then bound T_2 by:

$$\begin{aligned} T_2 & = \mathbb{E}_{\mathbf{x}_1^*} \left[\max_{t \leq T-1} r_t^{\text{traj}}(\hat{\pi}, \pi^*) \mathbf{1}\{\mathcal{E}_{\nabla \pi}(c_{\text{stab}})^c\} \right] \leq \mathbb{P} \left[\max_{t \leq T-1} \|\mathcal{P}_t \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_1^{\pi^*})\|_{\text{op}} \gtrsim c_{\text{stab}} \right] \\ & \leq \mathbb{P} \left[\max_{t \leq T-1} \frac{d_u}{\lambda_t \sigma_u^2} (r_t^{\text{est}}(\hat{\pi}; \pi^*) + r_t^{\text{est}}(\hat{\pi}; \tilde{\pi}^*)) \gtrsim c_{\text{stab}}^2 \right] \end{aligned}$$

$$\leq C_{\text{stab}}^2 \frac{d_u \max_t \lambda_t^{-1}}{\sigma_u^2} \mathbb{E}_{\mathbf{x}_1^{\pi^*}} \left[\max_{t \leq T-1} r_t^{\text{est}}(\hat{\pi}; \pi^*) + r_t^{\text{est}}(\hat{\pi}; \tilde{\pi}^*) \right],$$

where the second line arises from applying [Proposition B.5](#) and the noise-scale condition $\sigma_u \lesssim \sqrt{\frac{\lambda_t}{d_u} \frac{c_{\text{stab}}^3}{C_\pi}}$, and the last line comes from Markov's inequality. As for T_1 , we set the decomposition for a given $\tau \in (0, 1)$ to be determined later:

$$T_1 \leq \underbrace{\int_0^\tau \mathbb{P} \left[\max_{t \leq T-1} r_t^{\text{traj}}(\hat{\pi}, \pi^*) \mathbf{1}_{\{\mathcal{E}_{\nabla \pi}(c_{\text{stab}})\}} > \varepsilon \right] d\varepsilon}_{T_1^a} + \underbrace{\mathbb{P} \left[\max_{t \leq T-1} r_t^{\text{traj}}(\hat{\pi}, \pi^*) \mathbf{1}_{\{\mathcal{E}_{\nabla \pi}(c_{\text{stab}})\}} > \tau \right]}_{T_1^b}.$$

First, writing out the requirement of [Proposition B.6](#), casting $\varepsilon \mapsto \sqrt{\varepsilon}$, we have:

$$\begin{aligned} & \sup_{\substack{\|\mathbf{w}\| \leq 1, \mathbf{w} \in \mathcal{R}_t^{\pi^*}(\lambda_t) \\ \|\mathbf{r}\| \leq 1, \mathbf{r} \in \mathcal{R}_t^{\pi^*}(\lambda_t)^\perp \\ \|\mathbf{v}\| \leq 1}} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \sqrt{\varepsilon} \mathbf{w} + \frac{c_{\text{stab}}}{\log(1/\rho)} \sqrt{\lambda_t} \sqrt{\varepsilon} \mathbf{r} + C_{\text{rem}} \varepsilon \mathbf{v})\| \\ & \leq \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\| + \sup_{\substack{\|\mathbf{w}\| \leq 1, \mathbf{w} \in \mathcal{R}_t^{\pi^*}(\lambda_t) \\ \|\mathbf{r}\| \leq 1, \mathbf{r} \in \mathcal{R}_t^{\pi^*}(\lambda_t)^\perp \\ \|\mathbf{v}\| \leq 1}} \|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top (\sqrt{\varepsilon} \mathbf{w} + \frac{c_{\text{stab}}}{\log(1/\rho)} \sqrt{\lambda_t} \sqrt{\varepsilon} \mathbf{r} + C_{\text{rem}} \varepsilon \mathbf{v})\| \\ & \quad + 2C_\pi \|\sqrt{\varepsilon} \mathbf{w} + \frac{c_{\text{stab}}}{\log(1/\rho)} \sqrt{\lambda_t} \sqrt{\varepsilon} \mathbf{r} + C_{\text{rem}} \varepsilon \mathbf{v}\|^2 \\ & \leq \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\| + \sup_{\|\mathbf{w}\| \leq 1, \mathbf{w} \in \mathcal{R}_t^{\pi^*}(\lambda_t)} \|\nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})^\top \mathbf{w}\| \sqrt{\varepsilon} + 2L_\pi \frac{c_{\text{stab}}}{\log(1/\rho)} \sqrt{\lambda_t} \sqrt{\varepsilon} + 2C_{\text{rem}} L_\pi \varepsilon \\ & \quad + 6C_\pi \left(1 + \frac{c_{\text{stab}}}{\log(1/\rho)} \sqrt{\lambda_t} + C_{\text{rem}}^2 \right) \varepsilon \\ & \leq \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\| + \|\mathcal{D}_t \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|_{\text{op}} \sqrt{\varepsilon} \\ & \quad + 2L_\pi \frac{c_{\text{stab}}}{\log(1/\rho)} \sqrt{\lambda_t} \sqrt{\varepsilon} + \left(2C_{\text{rem}} L_\pi + 6C_\pi \left(1 + \frac{c_{\text{stab}}}{\log(1/\rho)} \sqrt{\lambda_t} + C_{\text{rem}}^2 \right) \right) \varepsilon. \end{aligned}$$

Let us interpret what this yields. On the last line, the first term is the on-expert error term $r_t^{\text{est}}(\hat{\pi}; \pi^*)$, the second term is controlled by [Proposition B.5](#), and the rest of the terms are the errors for which we do not guarantee control. To leverage [Proposition B.6](#), it suffices to have the last line bounded by $c_{\text{stab}} \sqrt{\varepsilon}$. Intuitively, the higher order error term scaling as ε automatically satisfies this for sufficiently small ε , which leaves the error term scaling as $\sqrt{\lambda_t} \varepsilon$. This is where we set the excitation levels $\{\lambda_t\}$ in hindsight. Observing the above, it suffices to set:

$$\lambda_t = \frac{1}{16} L_\pi^{-2} \log(1/\rho)^2, \quad t \geq 2.$$

In other words, for components of the controllability Gramian below this λ_t , the excitability is low enough such that we do not need to guarantee $\hat{\pi}, \pi^*$ match on them. For convenience, let us now define the quantity:

$$C_{\text{traj}} \triangleq 2C_{\text{rem}} L_\pi + 6C_\pi \left(1 + \frac{1}{4} \frac{c_{\text{stab}}}{L_\pi} + C_{\text{rem}}^2 \right).$$

Therefore, setting $\tau \approx C_{\text{traj}}^{-2} c_{\text{stab}}^2$, we may bound T_1^a by applying [Proposition B.6](#):

$$\begin{aligned}
T_1^a &= \int_0^\tau \mathbb{P} \left[\max_{t \leq T-1} r_t^{\text{traj}}(\hat{\pi}, \pi^*) \mathbf{1}_{\{\mathcal{E}_{\nabla\pi}(c_{\text{stab}})\}} > \varepsilon \right] d\varepsilon \\
&\leq \int_0^\tau \mathbb{P} \left[\max_{t \leq T-1} \sup_{\substack{\|\mathbf{w}\| \leq 1, \mathbf{w} \in \mathcal{R}_t^{\pi^*}(\lambda_t) \\ \|\mathbf{r}\| \leq 1, \mathbf{r} \in \mathcal{R}_t^{\pi^*}(\lambda_t)^\perp \\ \|\mathbf{v}\| \leq 1}} \|(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*} + \sqrt{\varepsilon} \mathbf{w} + \frac{1}{4} \frac{c_{\text{stab}}}{L_\pi} \sqrt{\varepsilon} \mathbf{r} + C_{\text{rem}} \varepsilon \mathbf{v})\|^2 \right. \\
&\quad \left. \cdot \mathbf{1}_{\{\mathcal{E}_{\nabla\pi}(c_{\text{stab}})\}} > c_{\text{stab}}^2 \varepsilon \right] d\varepsilon \\
&\leq \int_0^\tau \mathbb{P} \left[\max_{t \leq T-1} \left(r_t^{\text{est}}(\hat{\pi}; \pi^*) + \|\mathcal{D}_t \nabla_{\mathbf{x}}(\hat{\pi} - \pi^*)(\mathbf{x}_t^{\pi^*})\|_{\text{op}}^2 \varepsilon + C_{\text{traj}}^2 \varepsilon^2 \right) \mathbf{1}_{\{\mathcal{E}_{\nabla\pi}(c_{\text{stab}})\}} \gtrsim c_{\text{stab}}^2 \varepsilon \right] d\varepsilon \\
&= \int_0^\tau \mathbb{P} \left[\max_{t \leq T-1} r_t^{\text{est}}(\hat{\pi}; \pi^*) \gtrsim c_{\text{stab}}^2 \varepsilon \right] d\varepsilon \quad (\text{Def. of } \mathcal{E}_{\nabla\pi}(c_{\text{stab}}); C_{\text{traj}}^2 \varepsilon \lesssim c_{\text{stab}}^2 \text{ for } \varepsilon \leq \tau) \\
&\approx C_{\text{stab}}^2 \mathbb{E}_{\mathbf{x}_1^{\pi^*}} \left[\max_{t \leq T-1} r_t^{\text{est}}(\hat{\pi}; \pi^*) \right].
\end{aligned}$$

The bound on T_1^b follows similarly:

$$\begin{aligned}
T_1^b &\leq \mathbb{P} \left[\max_{t \leq T-1} r_t^{\text{est}}(\hat{\pi}; \pi^*) \gtrsim c_{\text{stab}}^2 \tau \right] \\
&\leq C_{\text{traj}}^2 C_{\text{stab}}^4 \mathbb{E}_{\mathbf{x}_1^{\pi^*}} \left[\max_{t \leq T-1} r_t^{\text{est}}(\hat{\pi}; \pi^*) \right]. \quad (\text{Markov's})
\end{aligned}$$

Putting everything together, we get the final bound:

$$\begin{aligned}
&\mathbb{E}_{\mathbf{x}_1^{\pi^*}} \left[\max_{t \leq T-1} r_t^{\text{traj}}(\hat{\pi}, \pi^*) \right] \leq T_1^a + T_1^b + T_2 \\
&\leq C_{\text{stab}}^2 \left((1 + C_{\text{traj}}^2 C_{\text{stab}}^2) \mathbb{E}_{\mathbf{x}_1^{\pi^*}} \left[\max_{t \leq T-1} r_t^{\text{est}}(\hat{\pi}; \pi^*) \right] + \frac{d_u \max_t \lambda_t^{-1}}{\sigma_u^2} \mathbb{E}_{\mathbf{x}_1^{\pi^*}} \left[\max_{t \leq T-1} r_t^{\text{est}}(\hat{\pi}; \pi^*) + r_t^{\text{est}}(\hat{\pi}; \tilde{\pi}^*) \right] \right) \\
&\approx C_{\text{stab}}^2 \left((1 + C_{\text{traj}}^2 C_{\text{stab}}^2) \mathbb{E}_{\mathbf{x}_1^{\pi^*}} \left[\max_{t \leq T-1} r_t^{\text{est}}(\hat{\pi}; \pi^*) \right] + \frac{d_u L_\pi^2}{\log(1/\rho)^2 \sigma_u^2} \mathbb{E}_{\mathbf{x}_1^{\pi^*}} \left[\max_{t \leq T-1} r_t^{\text{est}}(\hat{\pi}; \pi^*) + r_t^{\text{est}}(\hat{\pi}; \tilde{\pi}^*) \right] \right) \\
&\leq C_{\text{stab}}^2 \left(1 + C_{\text{traj}}^2 C_{\text{stab}}^2 + \frac{d_u L_\pi^2}{\log(1/\rho)^2 \sigma_u^2} \right) \mathbb{E}_{\mathbf{x}_1^{\pi^*}} \left[\max_{t \leq T-1} r_t^{\text{est}}(\hat{\pi}; \pi^*) + r_t^{\text{est}}(\hat{\pi}; \tilde{\pi}^*) \right],
\end{aligned}$$

which gives the desired result. \square

Therefore, by using the trivial bound $\mathbf{J}_{\text{TRAJ},2,T}^{\mathbf{x}}(\hat{\pi}) \leq T \mathbb{E}_{\mathbf{x}_1 \sim \mathbb{P}_{\mathbf{x}_1^{\pi^*}}} [\max_{1 \leq t \leq T} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\|^2 \wedge 1]$ and applying [Lemma A.1](#) to translate to $\mathbf{J}_{\text{TRAJ},2,T}(\hat{\pi})$, we get the final result.

Theorem 4 (Trajectory error bound; full ver. of [Theorem 2](#)). *Let [Assumption 4.1](#) hold. Let C_r be defined as in [Corollary B.1](#) and C_{rem} as in [Proposition B.6](#). Let the noise-scale $\sigma_u > 0$ satisfy*

$$\sigma_u \lesssim \min \left\{ \sqrt{\frac{\log(1/\rho)}{L_\pi d_u}} \frac{c_{\text{stab}}^3}{C_\pi}, \frac{\log(1/\rho)^2}{L_\pi^2 d_u} \frac{c_{\text{stab}}^4}{C_r}, c_{\text{stab}} \frac{\sqrt{1 + 4C_K^2}}{C_\pi} \right\}. \quad (\text{B.10})$$

Consider a candidate policy $\hat{\pi}$. Defining $C_{\text{traj}} \triangleq 2C_{\text{rem}}L_{\pi} + 6C_{\pi}\left(1 + \frac{1}{4}\frac{c_{\text{stab}}}{L_{\pi}} + C_{\text{rem}}^2\right)$, we may bound the trajectory error by the on-expert error on the mixture distribution $\mathbb{P}_{\pi^*, \sigma_u, 0.5}$ as:

$$\begin{aligned} \mathbf{J}_{\text{TRAJ}, 2, T}(\hat{\pi}) &\lesssim T(1 + L_{\pi}^2)C_{\text{stab}}^2 \left(1 + C_{\text{traj}}^2 C_{\text{stab}}^2 + \frac{d_u L_{\pi}^2}{\log(1/\rho)^2 \sigma_u^2}\right) \mathbf{J}_{\text{DEMO}, T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_u, \alpha}) \\ &= O_{\star}(T) \sigma_u^{-2} \mathbf{J}_{\text{DEMO}, T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_u, 0.5}). \end{aligned}$$

We conclude this section with a few technical remarks.

Remark B.1 (Horizon T dependence). We note the linear-in-horizon T dependence arises from a naive conversion between $\max_{1 \leq t \leq T}$ and $\sum_{t=1}^T$. We note that [Proposition B.7](#) can actually be interpreted as bounding $\mathbf{J}_{\text{TRAJ}, \infty, T}(\hat{\pi}) \leq O_{\star}(1) \mathbf{J}_{\text{DEMO}, \infty, T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_u, 0.5})$, for appropriately defined ∞ /"max"-norm, which does not exhibit any horizon dependence. We expect a more fine-grained analysis, e.g. leveraging [Lemma A.3](#), to similarly remove the T dependence from $\mathbf{J}_{\text{TRAJ}, p, T}$ and $\mathbf{J}_{\text{DEMO}, p, T}$, with the main technical barrier in extending [Proposition 4.3](#) ([Proposition B.6](#)).

Remark B.2 (Noise-scale σ_u dependence). We note that the final bound in [Theorem 4](#) has a σ_u^{-2} dependence. Firstly, we note that, by removing additive factors of σ_u (as in [Suboptimal Proposition 4.2](#) or [Proposition 4.1](#)), we do not need to trade-off σ_u with the on-expert error $\mathbf{J}_{\text{DEMO}, 2, T}(\hat{\pi}; \mathbb{P}_{\pi^*, \sigma_u, 0.5})$, and can in fact set σ_u as large as permissible up to the smoothness constraints, turning the dependence $O_{\star}(1)$. However, observing where σ_u arises in the proof of [Proposition B.7](#), it comes solely from applying Markov's inequality on the event $\mathbb{P}\left[\max_{t \leq T-1} \frac{d_u}{\lambda_t \sigma_u^2} (r_t^{\text{est}}(\hat{\pi}; \pi^*) + r_t^{\text{est}}(\hat{\pi}; \tilde{\pi}^*)) \gtrsim c_{\text{stab}}^2\right]$. We can envision instead applying a Chebyshev inequality. For example, if we square both sides, we raise the estimation error to quartic in $\|(\hat{\pi} - \pi^*)(\mathbf{x})\|$. If the estimation error satisfies moment-equivalence conditions, such as (4-2) hypercontractivity conditions that have appeared in prior learning-for-control literature [[Kakade et al., 2020](#), [Ziemann and Tu, 2022](#)], this pushes the σ_u dependence to an additive higher-order term. This crystallizes the intuition that the noise-level σ_u actually enters the trajectory error as a higher-order term (or equivalently, in the burn-in), explaining why huge differences in σ_u scale have similar effects on the final performance (see [Figure 1](#)). We avoid introducing these technical conditions in the body for clarity. Similarly, we note the proof of [Proposition B.7](#) also reveals that the on-expert error on the mixture distribution $\mathbb{P}_{\pi^*, \sigma_u, \alpha}$ enters only via the term depending on σ_u , and thus similarly the number of noised trajectories need not scale proportionally to n . This explains why the final performance of an imitator policy is often not sensitive to the exact proportion of noised trajectories α in the training data, as long as *some* trajectories are noised and some are clean; see [Figure 2](#).

B.4 Limitations of Prior Approaches

One may wonder what a control-oriented analysis as above buys compared to instantiating prior guarantees in the imitation learning literature. In particular, recent work in `LogLossBC` [[Foster et al., 2024](#)] reduces imitation learning to estimation in the Hellinger distance, which is achieved by regressing in the log-loss. However, as observed in [Simchowitz et al. \[2025\]](#), `LogLossBC` (and in the same vein earlier analyses [[Ross and Bagnell, 2010](#), [Ross et al., 2011](#)] that rely on the $\{0, 1\}$ loss) yields vacuous guarantees even for deterministic experts in continuous action spaces. Therefore, we consider fitting a noised expert and yield guarantees on the trajectory error of the resulting *noisy* rollouts. Contrast this with [Theorem 4](#), where the trajectories used in training may be executed noisily, but the trajectory error bound is measured on rolling out the *noiseless* expert and candidate policies. As a last caveat, we

note these works typically bound a cost suboptimality $J(\hat{\pi}) - J(\pi^*)$; this is generally a weaker notion than the trajectory error we consider, which via the formalism of integral probability metrics (IPMs) upper bounds the cost gap (see e.g. Sec 2.3 of Simchowitz et al. [2018]). We now introduce (stochastic) policies $\pi : \mathcal{X} \rightarrow \Delta(\mathcal{U})$, where:

$$\pi(\mathbf{x}) = \mathcal{N}(\pi(\mathbf{x}), \sigma_{\mathbf{u}}^2 d_u^{-1} \mathbf{I}_{d_u}), \pi \in \Pi. \quad (\text{B.11})$$

In other words, π encodes the deterministic policy π and a $\approx \sigma_{\mathbf{u}}$ -bounded noise-injection process [Definition 4.1](#), where we specify to scaled isotropic Gaussian noise for convenient evaluation of distributional distances.⁹ In particular, π^* denotes the *noisy expert policy*. A key step of `LOGLOSSBC` bounds the Hellinger error of a maximum likelihood estimator via a log-loss covering. Define an ε -log-loss-cover Π' of Π : for all $\pi \in \Pi$, there exists $\pi' \in \Pi'$ such that for all $\mathbf{x} \in \mathcal{X}$, $\mathbf{u} \in \mathcal{U}$, $\log(\mathbb{P}_{\pi}[\mathbf{u} | \mathbf{x}] / \mathbb{P}_{\pi'}[\mathbf{u} | \mathbf{x}]) \leq \varepsilon$. Denote $N_{\log}(\Pi, \varepsilon)$ as the smallest such cover. Then, the following guarantee on an MLE policy holds [Foster et al. \[2024, Prop. B.1\]](#).

Proposition B.8. *Given n trajectories of length T generated by the noised expert π^* , define the maximum likelihood policy:*

$$\hat{\pi} \in \arg \max_{\pi} \sum_{i=1}^n \sum_{t=1}^T \mathbb{P}_{\pi}[\tilde{\mathbf{u}}_t^{(i)} | \tilde{\mathbf{x}}_t^{(i)}].$$

Then, with probability at least $1 - \delta$, the resulting generalization error of $\hat{\pi}$ is bounded by

$$D_H(\hat{\pi}, \pi^*) \leq \inf_{\varepsilon > 0} \left\{ \frac{6 \log(2N_{\log}(\Pi, \varepsilon)/\delta)}{n} + 4\varepsilon \right\}.$$

Now, we observe for conditional-Gaussian policies (B.11) π, π' , the log-likelihood ratio is given by:

$$\log(\mathbb{P}_{\pi}[\mathbf{u} | \mathbf{x}] / \mathbb{P}_{\pi'}[\mathbf{u} | \mathbf{x}]) = \frac{d_u}{2\sigma_{\mathbf{u}}^2} (\|\pi'(\mathbf{x}) - \mathbf{u}\|^2 - \|\pi(\mathbf{x}) - \mathbf{u}\|^2).$$

Though the log-likelihood ratio is unbounded over the support $\mathbf{u} = \mathbb{R}^{d_u}$, we may truncate the domain, wherein the scaling is similar to $\text{KL}(\pi(\mathbf{x}) \parallel \pi'(\mathbf{x}))$, from which we have:

$$\text{KL}(\pi(\mathbf{x}) \parallel \pi'(\mathbf{x})) = \frac{d_u}{2\sigma_{\mathbf{u}}^2} \|\pi(\mathbf{x}) - \pi'(\mathbf{x})\|^2.$$

Notably, this implies an ε -cover in $\max_{\mathbf{x} \in \mathcal{X}} \text{KL}(\pi(\mathbf{x}) \parallel \pi'(\mathbf{x}))$ is equivalent to a $\sqrt{2\sigma_{\mathbf{u}}^2 d_u^{-1} \varepsilon}$ -cover of Π in $d(\pi, \pi') \triangleq \max_{\mathbf{x}} \|\pi(\mathbf{x}) - \pi'(\mathbf{x})\|_2$. For parametric classes with parameters in $\mathbb{R}^{d_{\theta}}$, $\log N_d(\Pi, \varepsilon) \approx d_{\theta} \log(1/\varepsilon)$, and thus converting between an ℓ^2 and KL cover only introduces additional logarithmic factors of $\sigma_{\mathbf{u}}$. However, for non-parametric classes such as those in the lower-bound constructions in [Theorem A \[Simchowitz et al., 2025\]](#), $\log N_d(\Pi, \varepsilon) \approx \text{poly}(1/\varepsilon)$, and thus converting to a KL cover worsens the dependence on ε and introduces additional *polynomial* factors of $\sigma_{\mathbf{u}}$ and d_u . Contrast this with [Suboptimal Proposition 4.2](#) or [Theorem 4](#), where the dependence is always $\sigma_{\mathbf{u}}^{-2}$, regardless of the statistical capacity of $\hat{\pi}, \pi^* \in \Pi$, since we are covering in ℓ^2 over the deterministic class, rather than in KL over the conditional-Gaussian class. In either case, we recall that this route of analysis ultimately only controls the rollout cost of *noised* policies. We now establish in the sequel, as insinuated by the upper bound in [Suboptimal Proposition 4.2](#), imitating purely on noised expert demonstrations yields an unavoidable bias scaling with $\sigma_{\mathbf{u}}$.

⁹This technically violates boundedness, but this is of minor concern by concentration of measure.

Suboptimality of only regressing on noise-injected trajectories To underscore the importance of imitating on *both* noise-injected and noiseless expert trajectories, we show via a simple example with maximally benign expert closed-loop dynamics that even perfect imitation on noise-injected trajectories necessarily incurs an additive factor in the trajectory error scaling with the smoothness of $\pi^* - \hat{\pi}$ and the noise-level σ^2 . Consider the system $\mathbf{x}_{t+1} = \mathbf{x}_t + \mathbf{u}_t$, expert policy $\pi^*(\mathbf{x}_t) = -\mathbf{x}_t$.

Proposition B.9 (Full ver. of Proposition 4.1). *Let the horizon $T = 3$ and $\mathbf{x}_1^{\pi^*}$ be fixed with $\|\mathbf{x}_1^{\pi^*}\| = 1$. Fixing any $\sigma_u \in (0, 1)$ and $C_\pi > 0$, let $\mathbf{z} \sim \mathcal{D}_z$ be any log-concave distribution with mean-zero and covariance satisfying $\Sigma_z \succeq \frac{1}{2d_u} \mathbf{I}_{d_u}$, and recall the corresponding noised expert states Definition 4.1. Then, there is a class of policies \mathcal{P} where any $\hat{\pi} \in \mathcal{P}$ satisfies: 1. $\frac{1}{n} \sum_{i=1}^n \|(\pi^* - \hat{\pi})(\sigma_u \mathbf{z}^{(i)})\| = 0$ with probability $\gtrsim 1 - n \exp(-\sqrt{d_u})$ where $\mathbf{z}^{(i)} \stackrel{\text{i.i.d.}}{\sim} \mathcal{D}_z$, 2. $\hat{\pi}(\mathbf{x}) = \pi^*(\mathbf{x})$ for all $\|\mathbf{x}\| \geq \sigma_u$, 3. $\hat{\pi}$ is C_π -smooth. However, the trajectory error induced by rolling out $\hat{\pi}$ is lower-bounded by:*

$$\mathbf{J}_{\text{TRAJ}, 2, T}(\hat{\pi}) \geq O(1) C_\pi^2 \sigma_u^4.$$

In other words, even when the candidate policy fits the expert *perfectly* on noise-injected expert trajectories, the trajectory error of the policies necessarily suffers a drift proportional to the smoothness budget C_π and noise-scale σ_u^2 , i.e. policies $\hat{\pi} \in \mathcal{P}$ and π^* are indistinguishable under purely noise-injected trajectories. On the other hand, a single un-noised trajectory from $\hat{\pi}$ and π^* can distinguish between the two policies perfectly.

Noting the expert closed-loop system here satisfies $C_{\text{reg}} = 0$, $C_K = 1$, $C_{\text{stab}} = 1$, we may compare to the key “expectation-to-uniform” step Lemma B.4 in establishing Suboptimal Proposition 4.2, where this lower bound matches the drift in the upper bound of Lemma B.4.

Proof of Proposition B.9. We first write out the noiseless expert’s trajectory:

$$\mathbf{x}_1^{\pi^*} = \mathbf{x}_1, \quad \mathbf{x}_2^{\pi^*} = \mathbf{x}_1^{\pi^*} - \mathbf{x}_1^{\pi^*} = \mathbf{0}, \quad \mathbf{x}_3^{\pi^*} = \mathbf{0}.$$

In other words, the expert reaches $\mathbf{0}$ in one timestep and stays there. Now consider the expert under the noising process $\tilde{\mathbf{u}} = \pi^*(\tilde{\mathbf{x}}) + \sigma_u \mathbf{z} = -\tilde{\mathbf{x}} + \sigma_u \mathbf{z}$, $\mathbf{z} \sim \mathcal{D}_z$: letting $\mathbf{z}_1, \mathbf{z}_2 \stackrel{\text{i.i.d.}}{\sim} \mathcal{D}_z$ be two i.i.d. draws of noise, we have

$$\tilde{\mathbf{x}}_1 = \mathbf{x}_1^{\pi^*}, \quad \tilde{\mathbf{x}}_2 = \tilde{\mathbf{x}}_1 + (-\tilde{\mathbf{x}}_1 + \sigma_u \mathbf{z}_1) = \sigma_u \mathbf{z}_1, \quad \tilde{\mathbf{x}}_3 = \tilde{\mathbf{x}}_2 + (-\tilde{\mathbf{x}}_2 + \sigma_u \mathbf{z}_2) = \sigma_u \mathbf{z}_2.$$

In other words, after timestep 1, since the expert policy always perfectly cancels out the previous state, the distribution of noised expert states is identical to the noise distribution $\sigma_u \mathbf{z}$. Therefore, the intuition for the lower bound is as follows: by concentration of measure, any “usual” distribution (e.g. log-concave, subgaussian) that has non-vanishing excitation, as captured by the second moment $\Sigma_z \succeq c \frac{1}{d_u} \mathbf{I}_{d_u}$, necessarily concentrates on the $O(1)\sigma_u$ -scaled unit sphere \mathbb{S}^{d_u} .¹⁰ Therefore, given n independent trajectories, i.e. n independent draws $\{(\mathbf{z}_1^{(i)}, \mathbf{z}_2^{(i)})\}$, with high probability we do not see any states $\tilde{\mathbf{x}}_1^{(i)}, \tilde{\mathbf{x}}_2^{(i)}$ within an $\approx \sigma_u$ radius of the origin. This is formalized in the following lemma [Paouris, 2006, Adamczak et al., 2014].

Lemma B.10 (Paouris’ Inequality [Paouris, 2006]). *Let \mathbf{z} be a log-concave random vector that with zero-mean and identity covariance supported on \mathbb{R}^d . Then, there exists a universal constant $c > 0$ such that for any $\gamma \geq 1$: $\mathbb{P}[\|\mathbf{z}\| \geq c\gamma\sqrt{d}] \leq \exp(-\gamma\sqrt{d})$.*

¹⁰We note that when \mathcal{D}_z is the uniform distribution on the unit sphere \mathbb{S}^{d_u} , then we may interchange the high-probability guarantee with expectation $\mathbb{E}[\|(\hat{\pi} - \pi^*)(\sigma_u \mathbf{z})\|^2] = 0$.

Therefore, re-scaling \mathbf{z} such that $\Sigma_{\mathbf{z}} \succeq \frac{1}{2d_u} \mathbf{I}_{d_u}$ and setting $\gamma = \frac{\sqrt{2}}{2c}$, this implies: $\mathbb{P}[\|\mathbf{z}\| \geq 1/2] \leq \exp(-\frac{\sqrt{2}}{2c} \sqrt{d_u}) \approx \exp(-\sqrt{d_u})$. Union bounding over $i = 1, \dots, n$, we have $\mathbb{P}[\|\mathbf{z}^{(i)}\| \geq 1/2, \forall i \in [n]] \gtrsim 1 - n \exp(-\sqrt{d_u})$.

Given that the noised expert states concentrate $\approx \sigma_{\mathbf{u}}$ away from the origin with overwhelming probability, we now task ourselves to constructing a family of candidate policies $\hat{\pi}$ that maximally deviate from the expert policy at the origin, given its smoothness budget C_{π} . This can be achieved, for example, by a straightforward bump function construction.

Lemma B.11 (Bump function existence, c.f. [Simchowitz et al. \[2025, Lemma A.15\]](#)). *For any $d \in \mathbb{N}$, we may construct a function $\text{bump}_d(\cdot) : \mathbb{R}^d \rightarrow \mathbb{R}$, $\text{bump}_d \in C^\infty$, such that the following hold:*

1. $\text{bump}_d(\mathbf{z}) = 1$ for all $\|\mathbf{z}\| \leq 1$.
2. $\text{bump}_d(\mathbf{z}) = 0$ for all $\|\mathbf{z}\| \geq 2$.
3. For each $p \geq 1$ and $\mathbf{z} \in \mathbb{R}^d$, $\|\nabla_p \text{bump}_d(\mathbf{z})\|_{\text{op}} \leq c_p$, where $c_p > 0$ is a constant depending on $p > 0$ but independent of dimension d .
4. $\nabla^p \text{bump}_d(\mathbf{z}) = \mathbf{0}$ for all $\|\mathbf{z}\| \geq 2$.

In other words, we may construct a function that always outputs 1 in the unit sphere, and 0 outside of the radius 2 sphere, and has bounded-norm derivatives in between. Before proceeding with the construction, we observe that $\pi^*(\mathbf{x}) = -\mathbf{x}$ is a linear function, and thus satisfies $\nabla^2 \pi^*(\mathbf{x}) = \mathbf{0}$ everywhere. For a given $\sigma_{\mathbf{u}} > 0$ and smoothness budget $C_{\pi} > 0$, it therefore suffices to determine $\Delta\pi = \hat{\pi} - \pi^*$ that satisfies the properties:

1. $\Delta\pi(\mathbf{x}) = \mathbf{0}$ for all $\|\mathbf{x}\| \geq \sigma_{\mathbf{u}}/2$.
2. $\|\nabla^2 \Delta\pi(\mathbf{x})\|_{\text{op}} \leq C_{\pi}$.

We construct $\Delta\pi$ as follows. Fix any $\mathbf{v} \in \mathbb{S}^{d_u}$, and let $\text{bump}_{d_x}(\cdot)$ be a function that satisfies the properties in [Lemma B.11](#). We propose:

$$\Delta\pi(\mathbf{x}) \triangleq L \text{bump}_{d_x}\left(\frac{\mathbf{x}}{\sigma_{\mathbf{u}}/4}\right) \mathbf{v}, \quad (\text{B.12})$$

where $L > 0$ is a constant to be determined later. We observe that by construction: $\Delta\pi = \mathbf{0}$ for all $\|\mathbf{x}\| \geq \sigma_{\mathbf{u}}/2$, $\|\Delta\pi(\mathbf{0})\| = L$, and $\|\nabla_{\mathbf{x}}^p \Delta\pi(\mathbf{x})\|_{\text{op}} = L \|\nabla_{\mathbf{x}}^p \text{bump}_{d_x}(4\mathbf{x}/\sigma_{\mathbf{u}})\|_{\text{op}} = L c_p \left(\frac{4}{\sigma_{\mathbf{u}}}\right)^p$. Therefore, to ensure $\Delta\pi$ is C_{π} -smooth, this informs choosing $L = \frac{\sigma_{\mathbf{u}}^2}{16c_2} C_{\pi}$. Therefore, the resulting policy $\hat{\pi} = \pi^* + \Delta\pi$ satisfies the following properties:

1. $\hat{\pi}$ is C_{π} -smooth.
2. $\|\hat{\pi}(\mathbf{0})\| = C_{\pi} \frac{\sigma_{\mathbf{u}}^2}{16c_2}$.
3. $\hat{\pi}(\mathbf{x}) = \pi^*(\mathbf{x})$ for all $\|\mathbf{x}\| \geq \sigma_{\mathbf{u}}/2$. In particular, by [Lemma B.10](#) that $\frac{1}{n} \sum_{i=1}^n (\hat{\pi} - \pi^*)(\sigma_{\mathbf{u}} \mathbf{z}^{(i)}) = \mathbf{0}$ with probability $\gtrsim 1 - n \exp(-\sqrt{d_u})$.

Now, we roll out $\hat{\pi}$ and π^* without noise injection. We have as aforementioned $\mathbf{x}_2^{\pi^*} = \mathbf{x}_3^{\pi^*} = \mathbf{0}$. On the other hand, since $\sigma_{\mathbf{u}} < 1$, we have $\hat{\pi}(\mathbf{x}_1^{\pi^*}) = \pi^*(\mathbf{x}_1^{\pi^*})$ and thus $\mathbf{x}_2^{\hat{\pi}} = \mathbf{x}_2^{\pi^*} = \mathbf{0}$. However, by our construction of $\hat{\pi}$, $\mathbf{x}_3^{\hat{\pi}} = \mathbf{x}_2^{\hat{\pi}} + \hat{\pi}(\mathbf{x}_2^{\hat{\pi}}) = \hat{\pi}(\mathbf{0}) = C_{\pi} \sigma_{\mathbf{u}}^2 (16c_2)^{-1}$, and thus:

$$\max_{t=1,2,3} \|\mathbf{x}_t^{\hat{\pi}} - \mathbf{x}_t^{\pi^*}\| = \|\mathbf{x}_3^{\hat{\pi}} - \mathbf{x}_3^{\pi^*}\| \geq O(1) C_{\pi} \sigma_{\mathbf{u}}^2.$$

After squaring both sides, we see the left-hand side is precisely $\mathbf{J}_{\text{TRAJ},2,T}^{\mathbf{x}}$, which is trivially upper bounded by $\mathbf{J}_{\text{TRAJ},2,T}$.

□

This lower bound establishes the unavoidable drift from noise-injection due to nonlinearity of the expert policy, thus highlighting the necessity of imitating on a dataset consisting of *both* noise-injected and clean expert trajectories; though, as discussed in the previous section, the exact proportion of each may not be important.

C Experiment Details

We describe the common specifications across experiments:

- **Model:** we use neural networks with two hidden layers of dimension 256 and GELU activations [Hendrycks and Gimpel, 2016]. We also additionally place batch-norm layers after the input and first hidden layers for the noise-injection experiments. For the action-chunking experiments, we remove the batchnorm layers as well as the layer biases to introduce a mild inductive bias of the model outputting $\mathbf{0}$ at the origin.
- **Optimizer and training:** we use the AdamW optimizer [Loshchilov, 2017] with a cosine decay learning rate schedule [Loshchilov and Hutter, 2016], with initial learning rate of 0.001 and other hyperparameters set as default. The models are trained for 4000 epochs with a batch size of 64. Evaluation statistics of each model are computed on an independent sample of 100 trajectories.

C.1 Action Chunking

For the action-chunking experiment, we consider a synthetic nonlinear system that is open-loop EISS, and closed-loop EISS under a deterministic expert, as constructed in Appendix E.1 and J of Simchowitz et al. [2025]. In particular, we first consider matrices:

$$\mathbf{A} = \begin{bmatrix} 1 + \mu & \frac{3}{2}\mu \\ -\frac{3}{2}\mu & 1 - 2\mu \end{bmatrix}, \quad \mathbf{K} = \begin{bmatrix} -(1 + \mu) & -\frac{3}{2}\mu \\ \frac{3}{2}\mu & 0 \end{bmatrix},$$

where we set $\mu = 1/8$. We then embed these matrices into a 6-dimensional state and input space:

$$\tilde{\mathbf{A}} = \begin{bmatrix} \mathbf{A} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} \end{bmatrix}, \quad \tilde{\mathbf{K}} = \begin{bmatrix} \mathbf{K} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} \end{bmatrix} \in \mathbb{R}^{6 \times 6}.$$

These matrices are respectively embedded into smooth nonlinear dynamics f and expert policy π^* as described in Construction E.1 [Simchowitz et al., 2025]. For the requisite smooth function g in the embedding, we generate a randomly initialized neural network with 1-hidden layer of dimension 16, with weights following the Xavier normal initialization [Glorot and Bengio, 2010] and biases sampled entrywise from $\text{Unif}([-1, 1])$; note that we only generate this network *once* to complete the problem instance. Having generated a “hard instance” indexed by $(\tilde{\mathbf{A}}, \tilde{\mathbf{K}}, g)$, the training data is comprised of 100 independent trajectories of length 64 rolled out under the expert policy π^* . For Figure 1, we train a behavior cloning agents; for each chunk length, the BC policy takes the form as described above, with the sole difference being the output dimension, which equates to $6 \times \text{chunk_len}$. Given the training recipe above, all BC policies across chunk lengths $\in \{1, 2, 4, 8, 16\}$ achieve training error of at most 10^{-6} , attaining near perfect imitation on the expert data.

Crucially, we note that, in contrast to our formal prescription in Intervention 1, we are not enforcing that the chunked BC policies accompany a simulated dynamics that it stabilizes, and purely treat the

policy as an ℓ -step action predictor. Beyond the soft inductive biases in ensuring the policies output $\mathbf{0}$ at the origin, we make no effort to enforce “simulated” stability, yet we still see the clear stabilization benefits of action-chunking in Figure 1.

C.2 Noise Injection

For the noise injection experiments depicted in Figure 1 (left) and Figure 2, we used the HalfCheetah-v5 environment through the Gymnasium library [Towers et al., 2024]. The expert policy is given by a Soft Actor-Critic [Haarnoja et al., 2018] RL policy pre-trained using the StableBaselines3 library [Raffin et al., 2021], downloaded from Huggingface [url]. When collecting expert demonstrations, we set `deterministic=True`. Given noise-scale $\sigma_{\mathbf{u}}$, we use scaled Gaussian noise $\mathcal{N}(\mathbf{0}, \frac{\sigma_{\mathbf{u}}^2}{d_{\mathbf{u}}} \mathbf{I}_{d_{\mathbf{u}}})$ as the noise-injection distribution. We train models for each problem parameter configuration, where the training data comprises of 40 *total* trajectories of $T = 300$ timesteps each. We then plot the resulting median per-timestep cumulative reward evaluated on 100 independent trajectories, normalized by the median expert cumulative reward at the problem horizon. For each figure specifically:

- **Figure 1 (right):** We sweep over noise-levels $\sigma_{\mathbf{u}} \in \{0.0, 0.001, 0.01, 0.1, 0.5, 1.0\}$, fixing the proportion of clean trajectories at 60%, equivalent to imitating over $\mathbb{P}_{\pi^*, \sigma_{\mathbf{u}}, 0.6}$ from Intervention 2. We note that noise-level 0.0 corresponds to vanilla behavior cloning.
- **Figure 2 (left):** We sweep over proportion of clean trajectories $\alpha \in \{0.0, 0.2, 0.4, 0.6, 0.8, 1.0\}$, holding noise-level $\sigma_{\mathbf{u}} = 1.0$ fixed. We note that $\alpha = 1.0$ corresponds to vanilla behavior cloning, and $\alpha = 0.0$ corresponds to pure noise-injection $\mathbb{P}_{\pi^*, \sigma_{\mathbf{u}}}$ (see Proposition 4.1).
- **Figure 2 (right):** We consider for $\sigma_{\mathbf{u}} \in \{0.5, 1.0\}$ the effect of recording clean versus noisy action labels. Recall that Intervention 2 prescribes *executing* expert actions noisily $\tilde{\mathbf{x}}_{t+1} = f(\tilde{\mathbf{x}}_t, \pi^*(\tilde{\mathbf{x}}_t) + \sigma_{\mathbf{u}} \mathbf{z}_t)$, but records the clean action label $\tilde{\mathbf{u}}_t = \pi^*(\tilde{\mathbf{x}}_t)$. On the other hand, the “RL-theoretic” approach, in order to achieve density, also requires recording the noisy label $\tilde{\mathbf{u}}_t = \pi^*(\tilde{\mathbf{x}}_t) + \sigma_{\mathbf{u}} \mathbf{z}_t$. We fix the proportion of clean trajectories to 0.0 for both set-ups for fair comparison.